

# Secure Serverless Computing using Dynamic Information Flow Control

**Kalev Alpernas**, Cormac Flanagan, Sadjad Fouladi, Leonid Ryzhyk,  
Mooly Sagiv, Thomas Schmitz, and Keith Winstein

 @kalevalp    kalevalp



vmware®



UC SANTA CRUZ

Stanford

# The Cloud is Growing Faster than Ever

Forbes

Billionaires Innovation Leadership Money Consumer Industry Lifestyle Fe

2,451 views | Oct 27, 2018, 06:00pm

## Stellar Intel Earnings Underscore Tech Sector Strength, Enormous Growth Potential In The Cloud



**Dave Altavilla** Contributor ⓘ

Consumer Tech

*I cover break-out tech in mobile, on desktop and in the data center.*

f

tw

in



# The Cloud is Growing Faster than Ever

Forbe

## Strong Cloud Growth Continues to Drive Microsoft Higher

Growth was strong across all segments, but Azure was still the highlight.

Danny Vena (TMFLifelsGood)  
Oct 28, 2018 at 10:37AM

Over the past several years, **Microsoft** (NASDAQ:MSFT) has accomplished an amazing feat, transforming from a shrink-wrapped business software provider to one of the biggest names in cloud computing. The company has been successful challenging the cloud leader, **Amazon** Web Services, and making meaningful headway, currently ranking in a strong second place.

As the global digital transformation gains steam, investors were watching closely for signs that Microsoft's cloud momentum would continue. When the company reported the financial results of its fiscal 2019 first quarter, which ended Sept. 30, 2018, Microsoft showed that it has much more growth in the tank, producing a record first quarter.



### AUTHOR



Danny Vena  
(TMFLifelsGood)

Daniel W. Vena, CPA, CGMA is long-term investor searching for intangibles that provide explosive growth opportunities in his investments. He served on active duty with the US Army and has a Bachelors degree in accounting.

 Follow @dannyyvena

335 followers

### ARTICLE INFO

Oct 28, 2018 at 10:37AM

Technology and Telecom

### STOCKS

# The Cloud is Growing Faster than Ever

Forbes

Str  
Hig  
Growth  
Danny Ver  
Oct 28, 2  
Over the  
transform  
cloud co  
Services  
As the gl  
Microsof  
of its fisc  
more gro

f  
tw  
in

Forbes

Billionaires Innovation Leadership Money Consumer Industry Lifestyle Feat

18,495 views | Jan 7, 2018, 07:36pm

## 83% Of Enterprise Workloads Will Be In The Cloud By 2020



**Louis Columbus** Contributor ⓘ

### TWEET THIS



Digitally transforming enterprises (63%) is the leading factor driving greater public cloud engagement or adoption today.



66% of IT professionals say security is their most significant concern in adopting an enterprise cloud computing strategy.

f



crosoft

i-term investor searching for growth opportunities in his duty with the US Army and has a

followers

VI  
n

# The Cloud is Growing Faster than Ever

Forbe

Str For

CLOUD

Back to Home

## Is security now a driver of cloud adoption?

JAMES NUNNS EDITOR  
23RD OCTOBER 2017

+ INCREASE / DECREASE 1

Growth

Danny Ver  
Oct 28, 2017

Over the  
transfor  
cloud co  
Services

As the gl  
Microsof  
of its fisc  
more gro

f

tw

in





# The Internet is a Leaky Place

## Facebook just had its worst hack ever – and it could get worse

By [Donie O'Sullivan, CNN Business](#)

Updated 9:22 AM ET, Thu October 4, 2018





**New York (CNN)** – On Sunday, September 16, engineers at Facebook detected some unusual activity on the social media platform's networks. It was an attack, the [biggest security breach in Facebook's history](#). And it would take the company 11 more days to stop it.

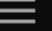
Now, almost a week since the public was first told of the attack, we still barely know anything about what happened.


We don't know who the hackers were, or what they were looking for. We don't know whether they were targeting particular people in certain countries. We don't know how long they had access to users' information. And we don't know what, if anything, they took.

# The Internet is a Leaky Place



Markets Tech Media Success Perspectives Video

U.S. Edition + 

  
get

By Donle O  
Updated 9:

New York C  
activity o  
in Facebo  
Now, alm  
about wh  
We don't  
they were  
access to

## Cathay Pacific got hacked, compromising the data of millions of passengers





By Jethro Mullen, CNN Business

Updated 12:22 AM ET, Thu October 25, 2018

**Hong Kong (CNN Business)** — One of Asia's top airlines has discovered a data breach in which the personal information of more than 9 million passengers may have been stolen.

Cathay Pacific ([CPCAY](#)) [said late Wednesday](#) that a wide range of data — including passengers' names, dates of birth, phone numbers, email addresses and passport numbers — was exposed in a hack of its information systems earlier this year.

"We are very sorry for any concern this data security event may cause our passengers," CEO Rupert Hogg [said in a statement](#). The Hong Kong-based carrier is in the process of contacting affected people, he added.



# The Internet is a Leaky Place

CNN BU CNN

Forbes

Billionaires Innovation Leadership Money Consumer Industry Lifestyle Featured BrandVoice Lists

Q

7,148 views | Oct 26, 2018, 07:30am

## Amazon Web Services Customers Can Hack AWS Cloud And Steal Data, Says Oracle CTO Larry Ellison



**Bob Evans** Contributor ⓘ  
Enterprise & Cloud



(Note: After an award-winning career in the media business covering the tech industry, Bob Evans was VP of Strategic Communications at SAP in 2011, and Chief Communications Officer at Oracle from 2012 to 2016. He now runs his own firm, Evans Strategic Communications LLC.)

CLOUD WARS -- Oracle founder Larry

f

t

in



# The Internet is a Leaky Place

CNN BU CNN I For

HealthcareITNews

GLOBAL EDITION TOPICS SIGN UP MAIN MENU

Face Ca  
get mi

By Donle O By Jeth

Updated 9: Updated

New York ( Hong K  
activity of person  
in Facebo Cathay  
Now, alm passer  
about wh — was  
We don't "We ar  
they were Rupert  
access to affecte

## The real victim in health data breaches? Patients' medical identities

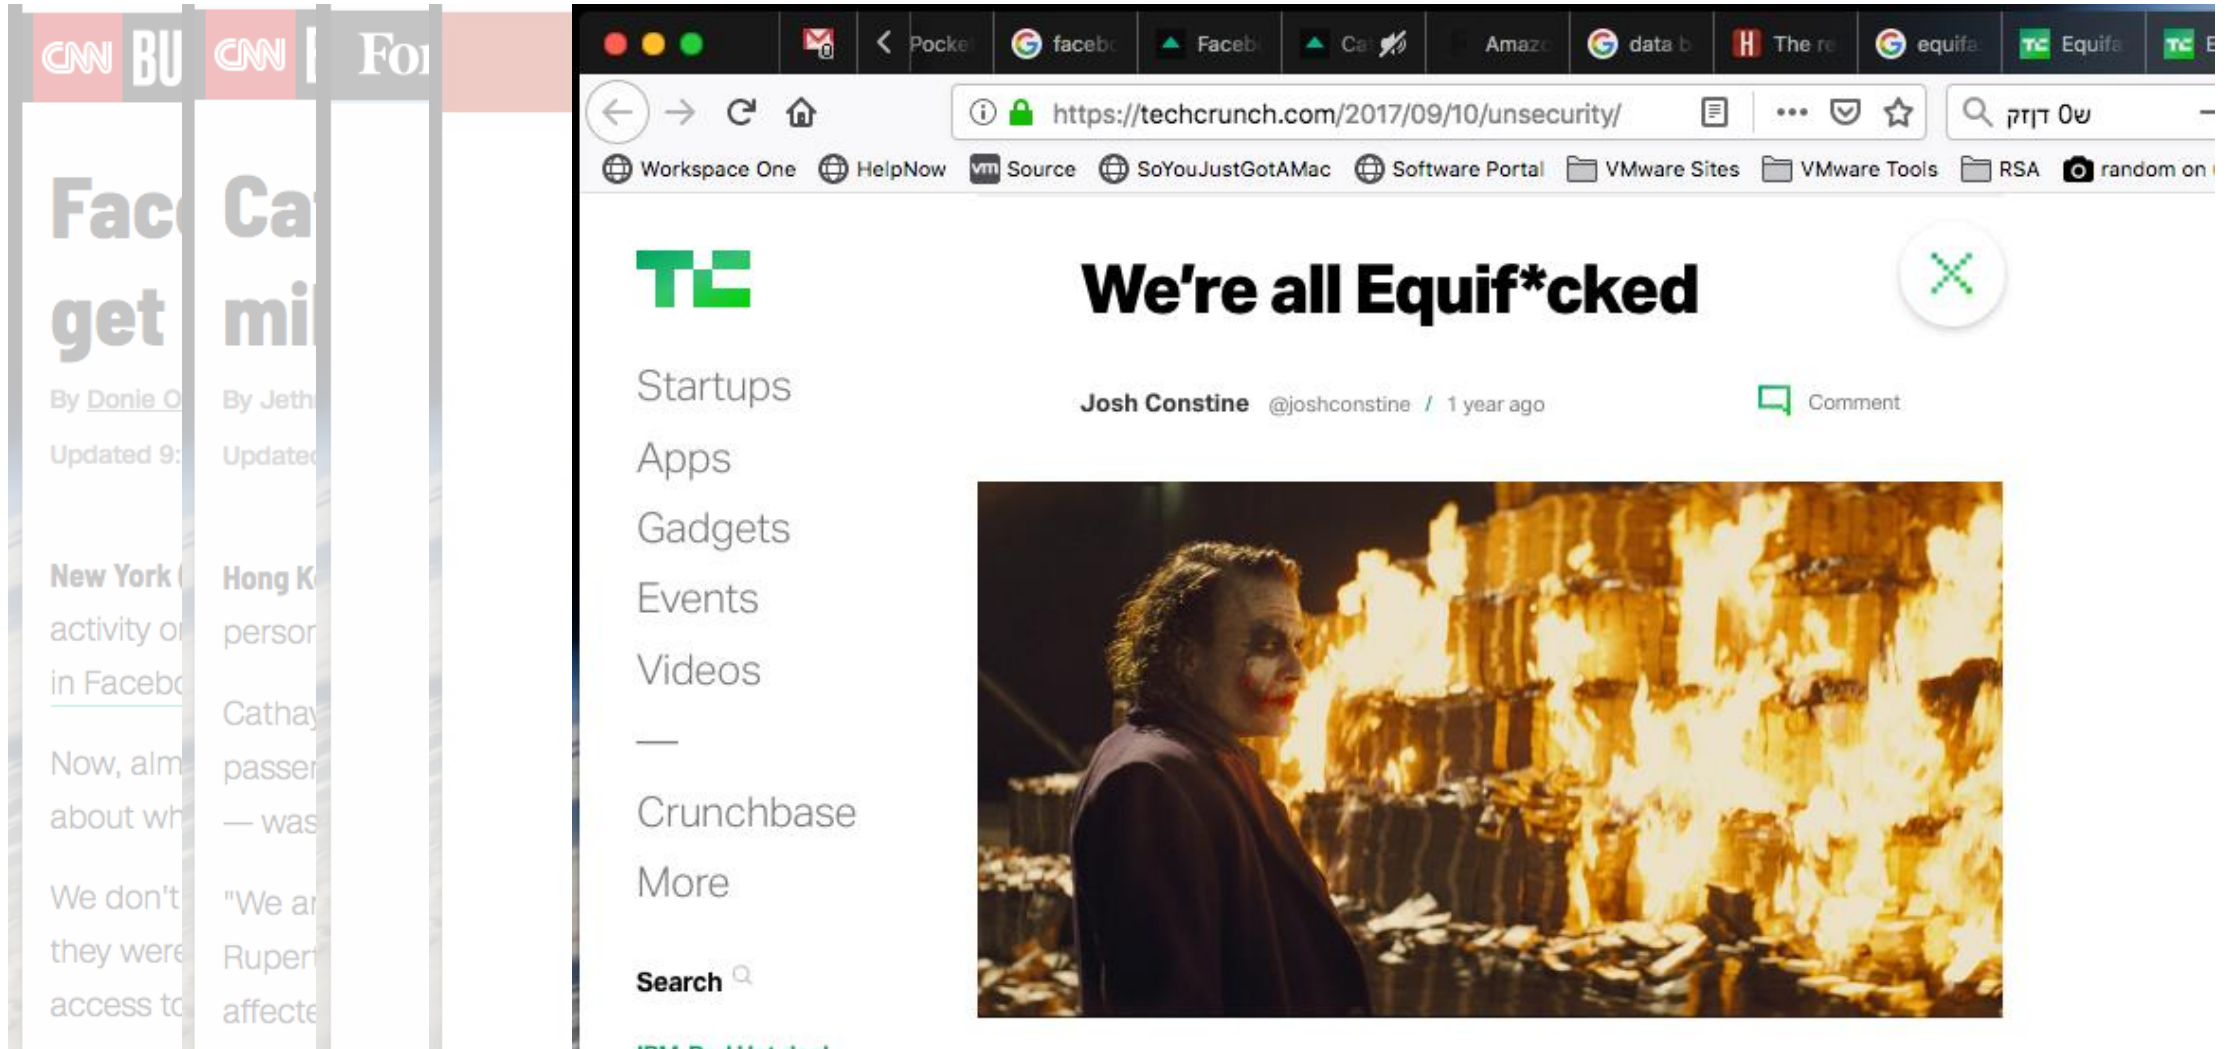
By [Jessica Davis](#) | October 29, 2018 | 05:03 PM

[f](#) [t](#) [in](#) [✉](#)

Hackers can perform synthetic identify theft long after a cybersecurity event by piecing together data to conduct medical and insurance fraud.

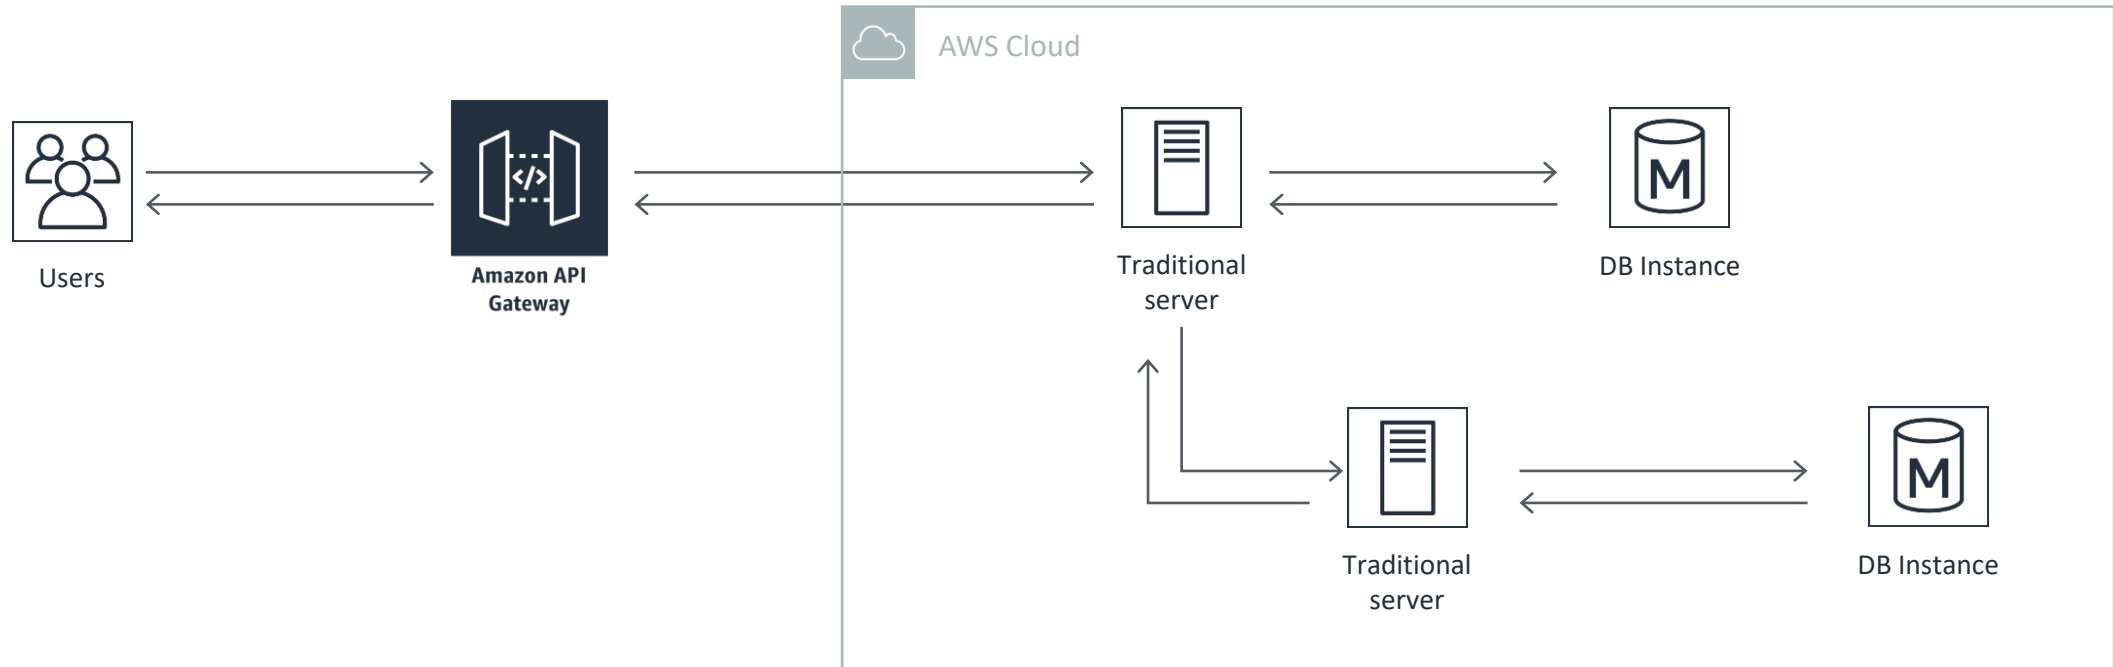


# The Internet is a Leaky Place



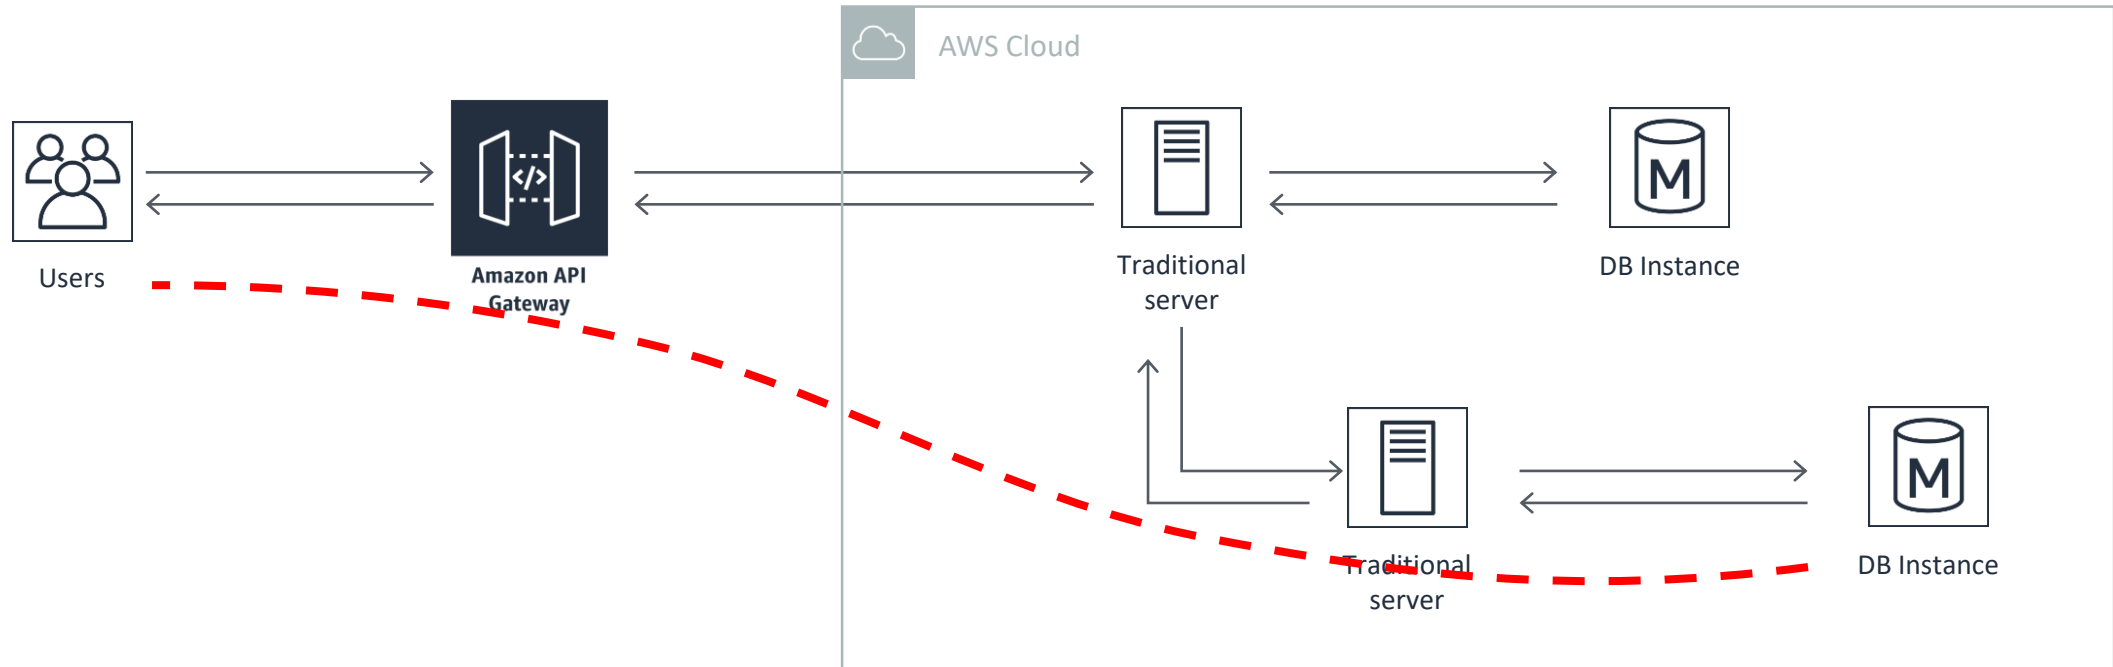
# Cloud Applications

Traditional (monolithic) three tiered application



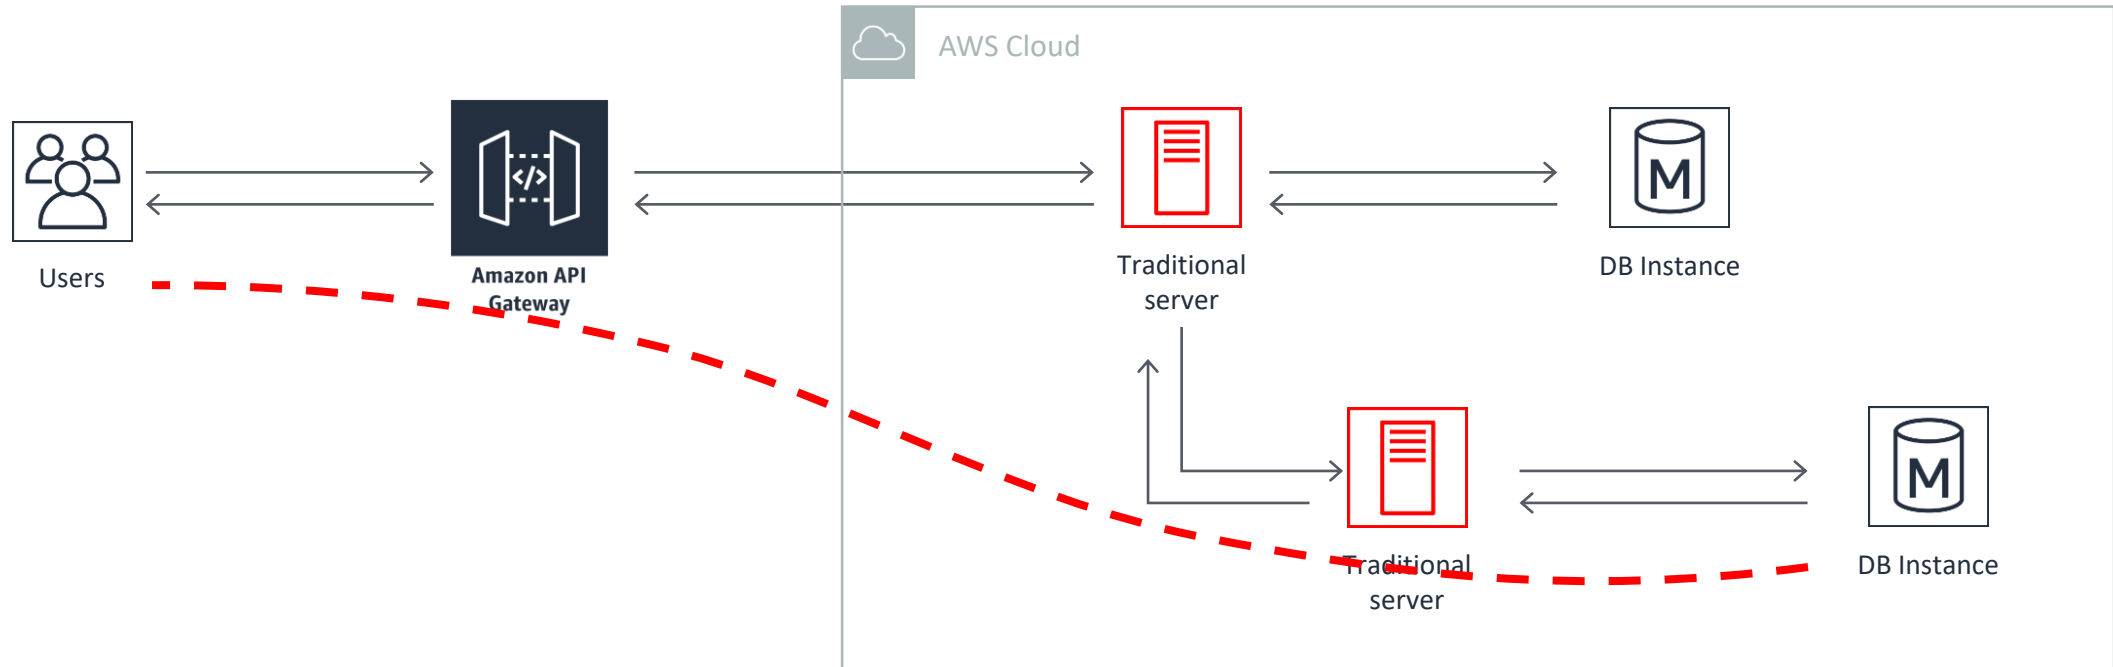
# Cloud Applications

Goal: keep data secure – Information Flow Control (IFC)  
*Non-Interference*



# Cloud Applications

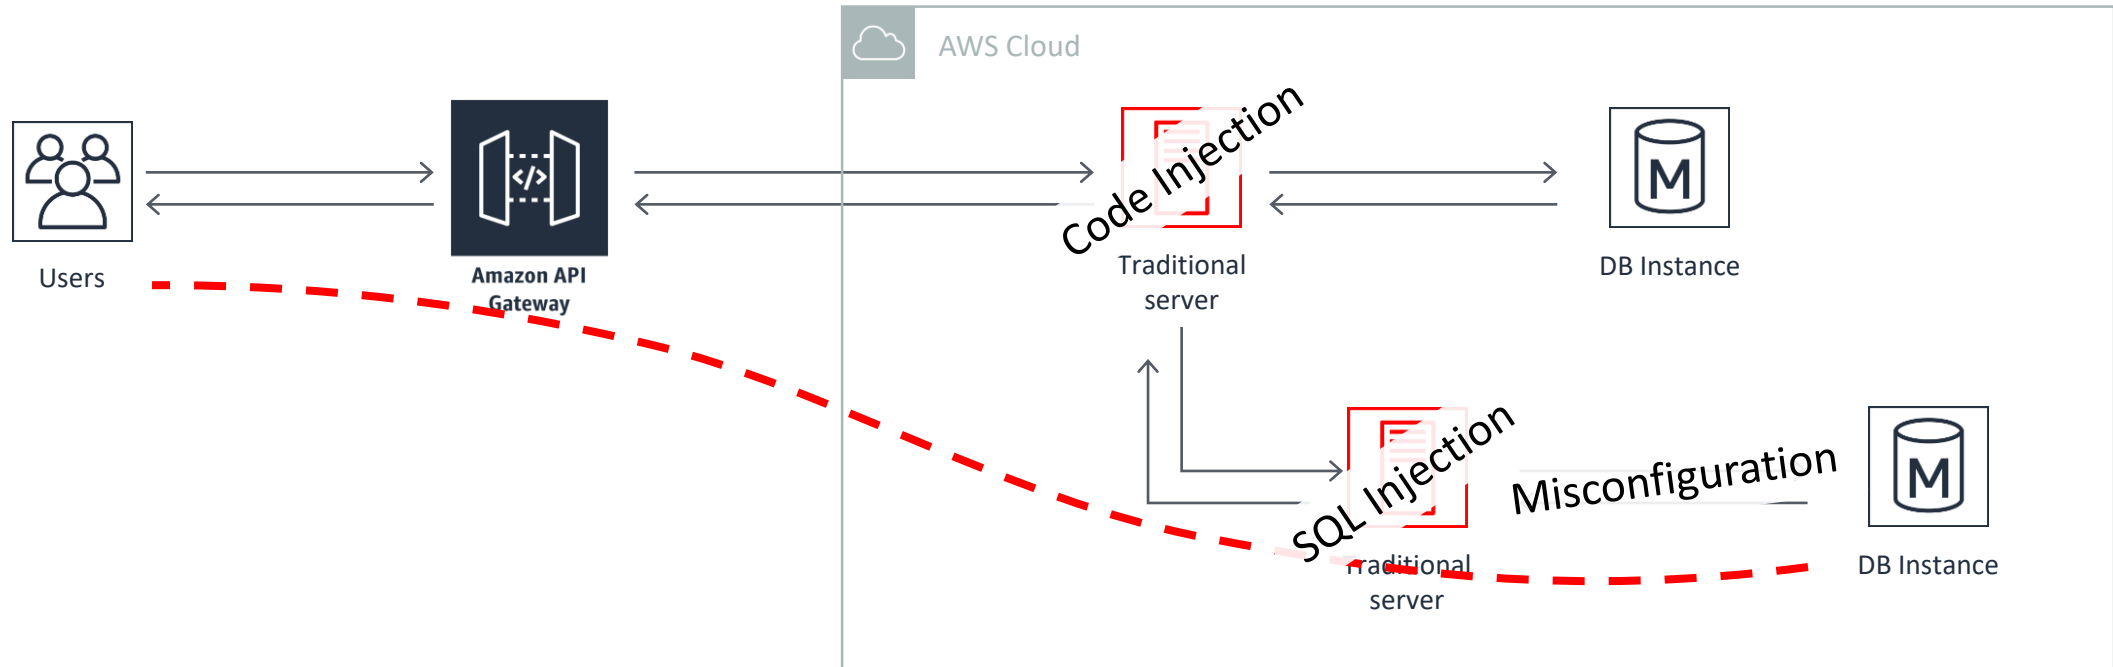
Servers are vulnerable





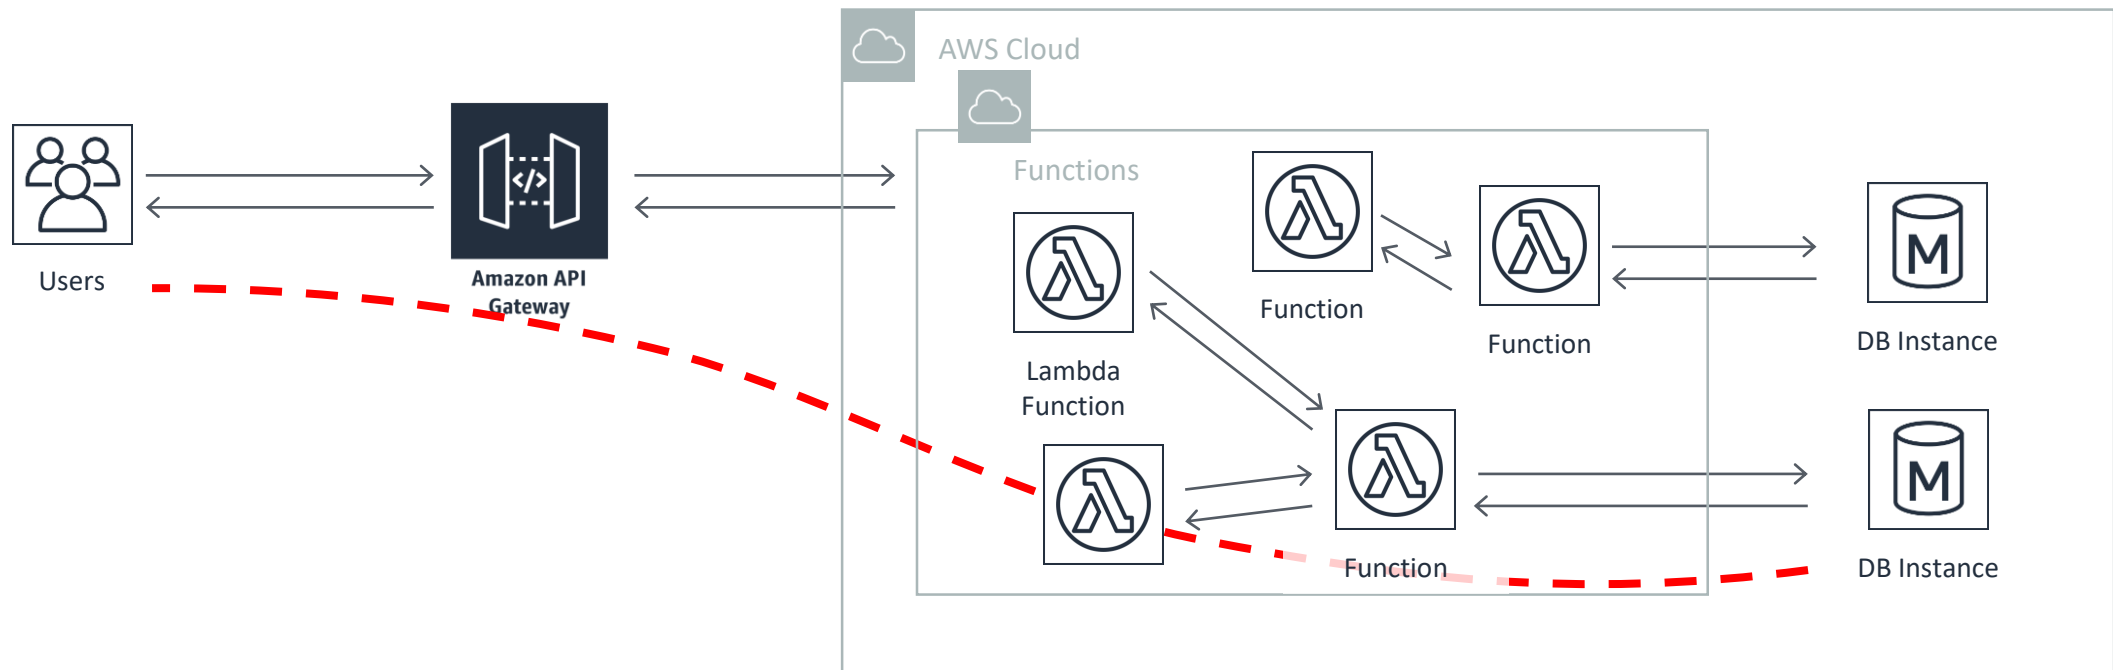
# Cloud Applications

Servers are vulnerable



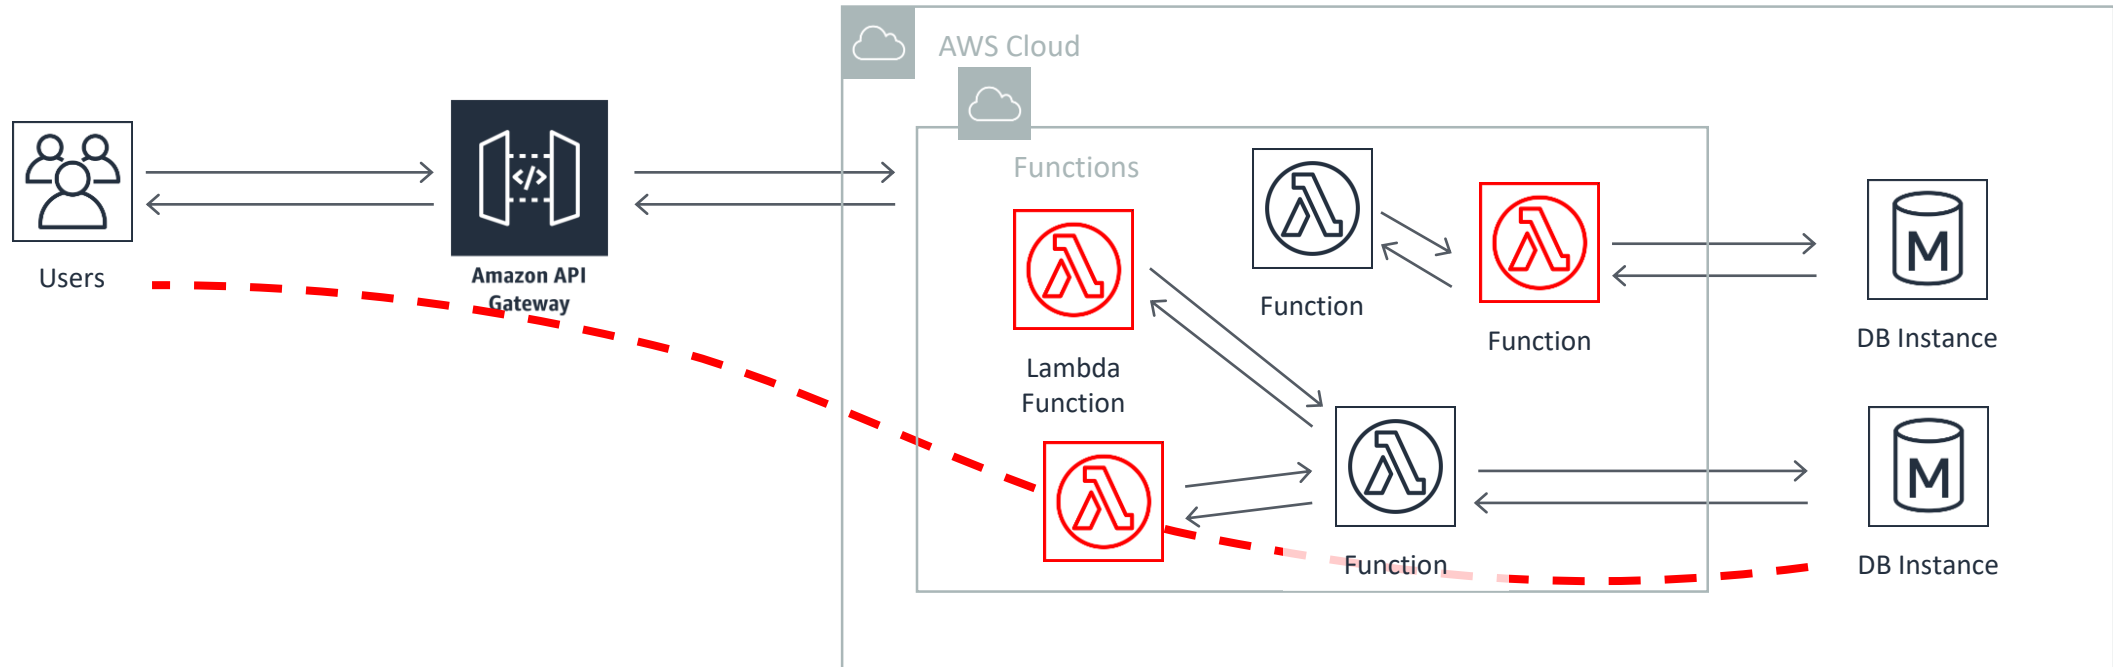
# Cloud Applications

What happens when we replace them with serverless functions?



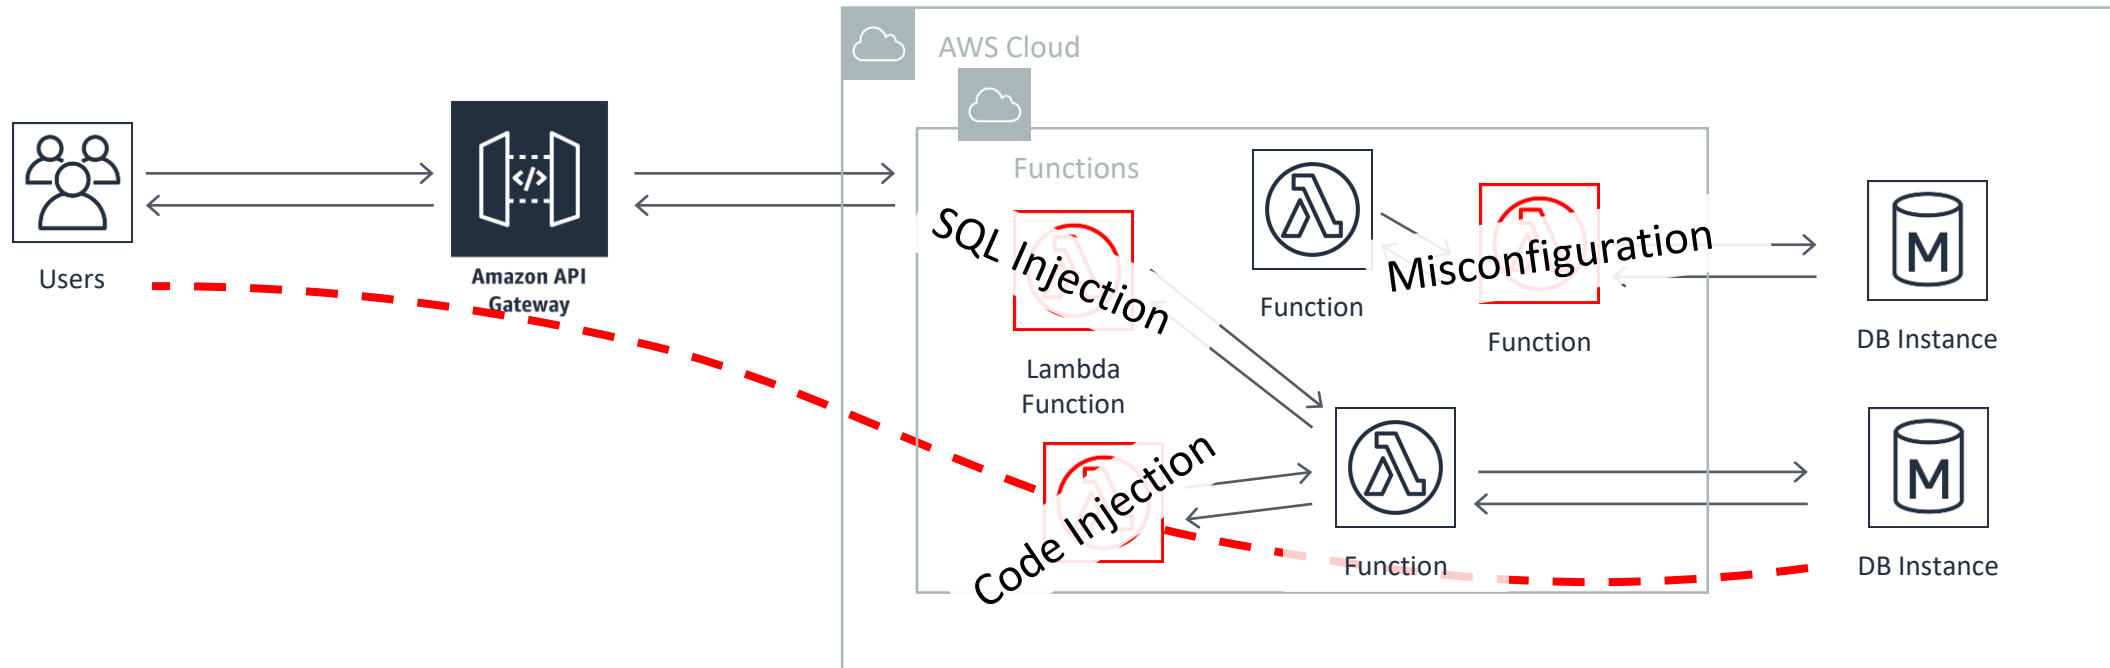
# Cloud Applications

Functions are vulnerable



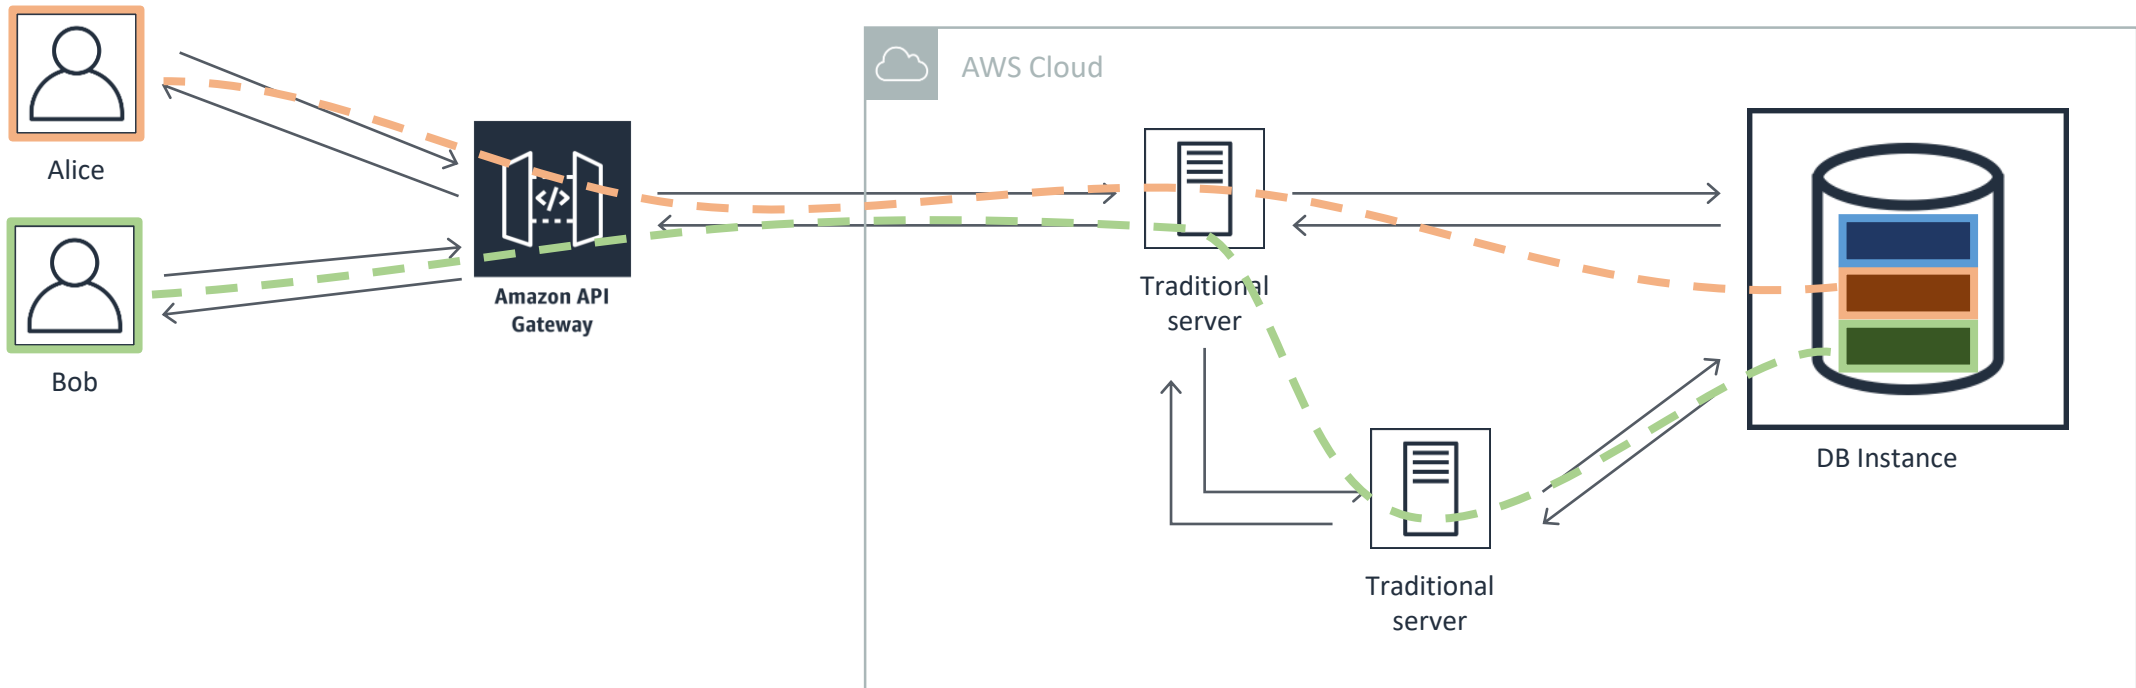
# Cloud Applications

Polyglot, distributed environment makes things even harder



# Monolithic Applications

Long running processes, serving multiple users, managing complex local state





# Serverless Applications

Every execution serves a single request, no local state



Alice



Bob

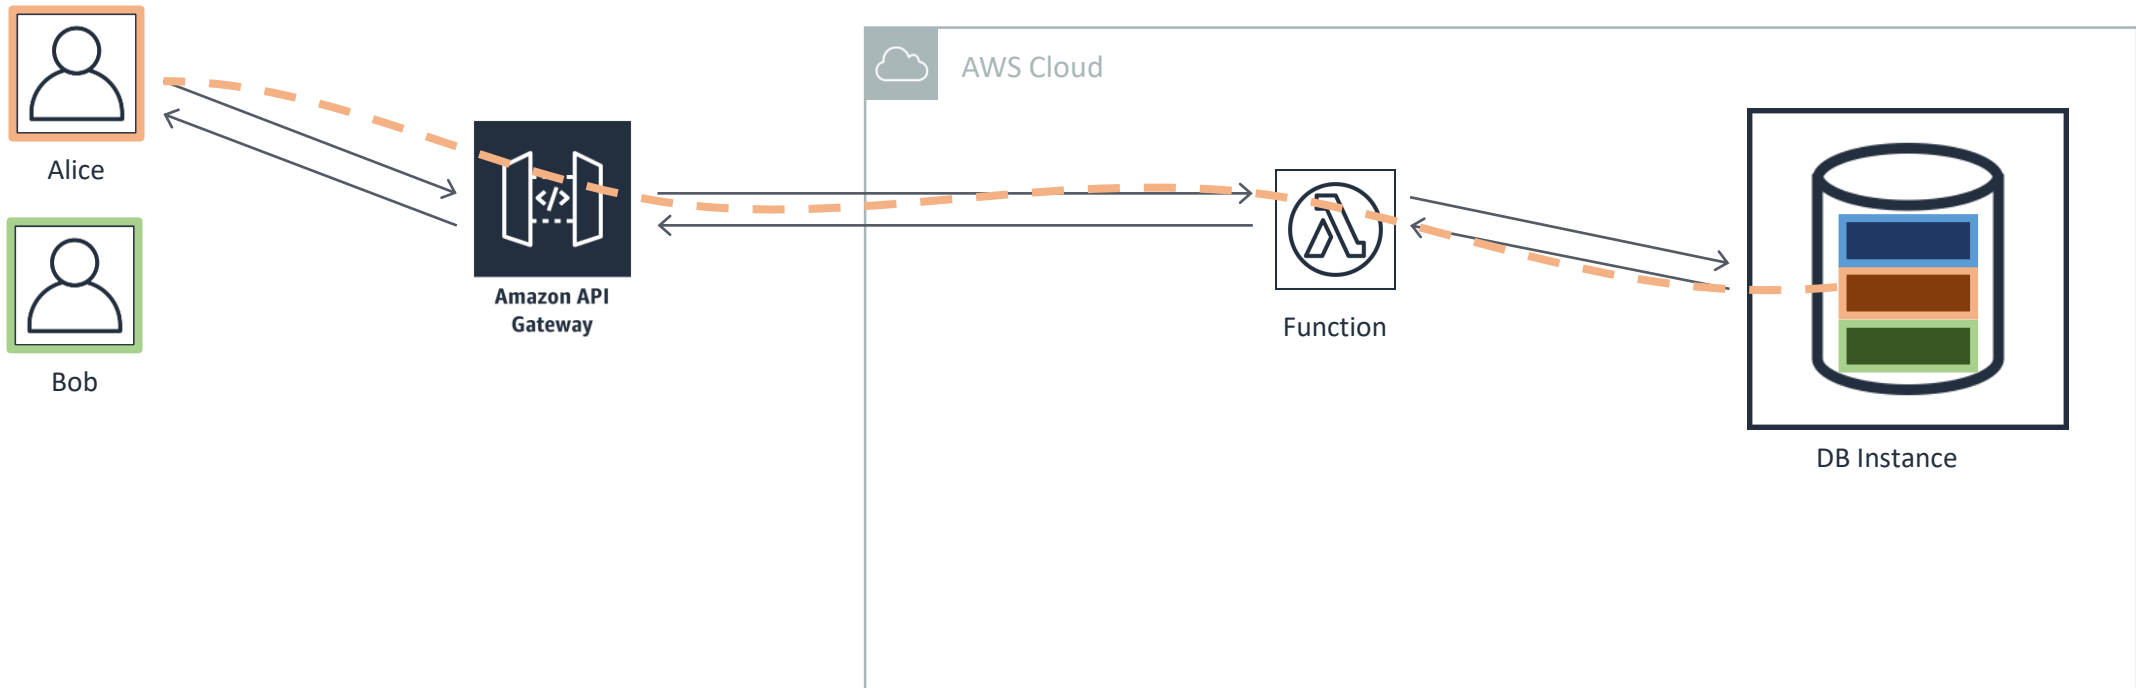


Amazon API  
Gateway



# Serverless Applications

Every execution serves a single request, no local state



# Serverless Applications

Every execution serves a single request, no local state



Alice



Bob

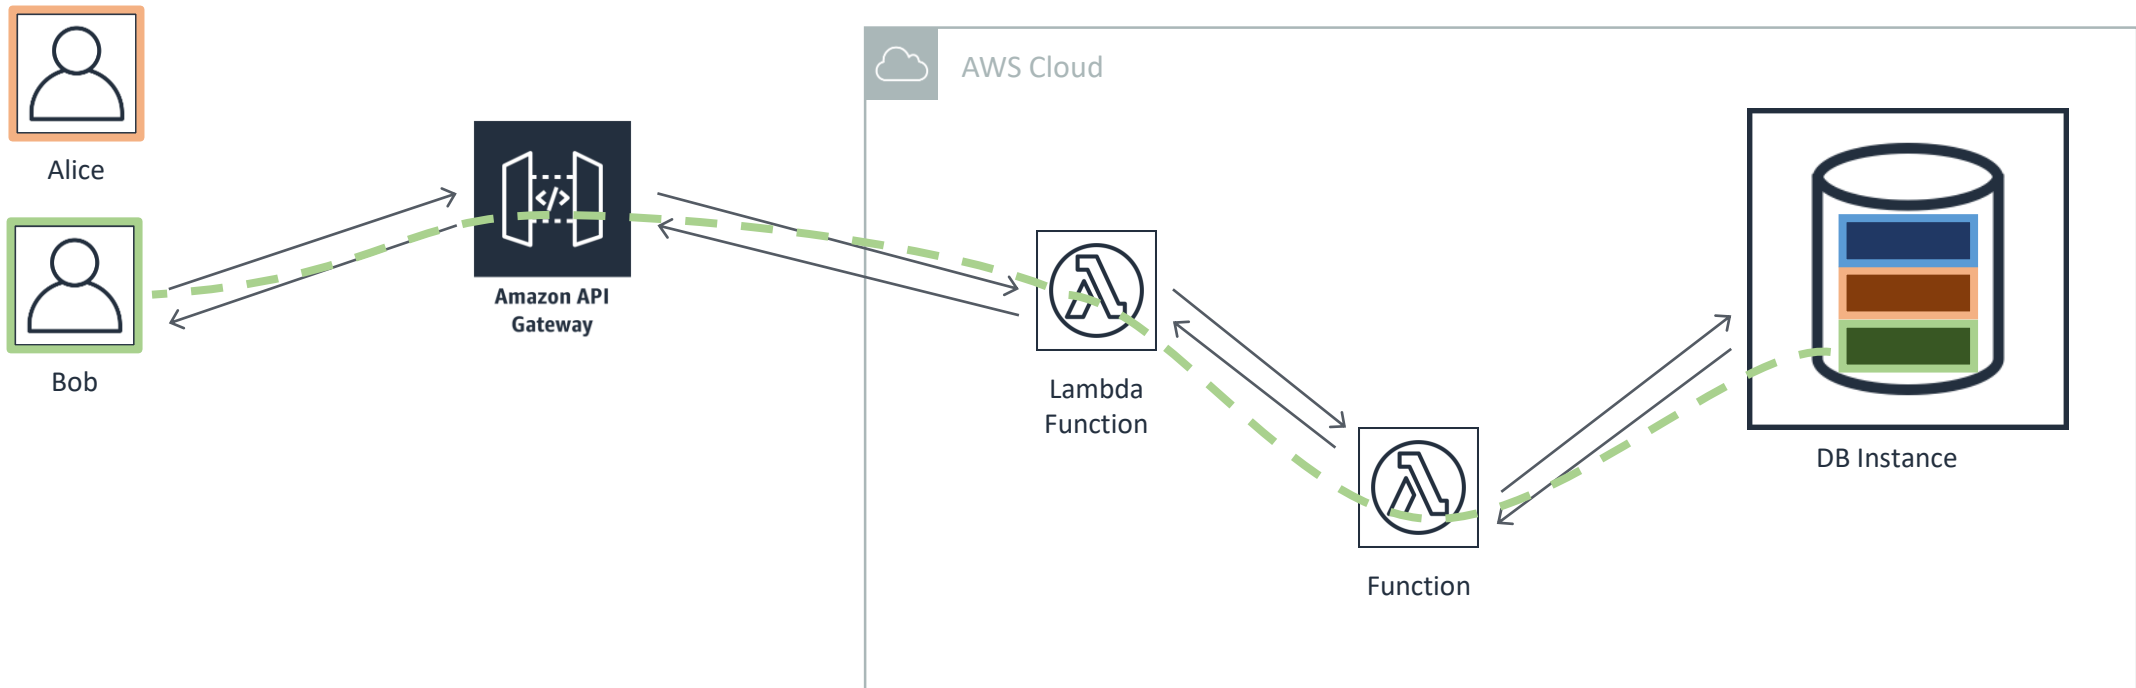


Amazon API  
Gateway



# Serverless Applications

Every execution serves a single request, no local state



# Serverless Applications

Every execution serves a single request, no local state



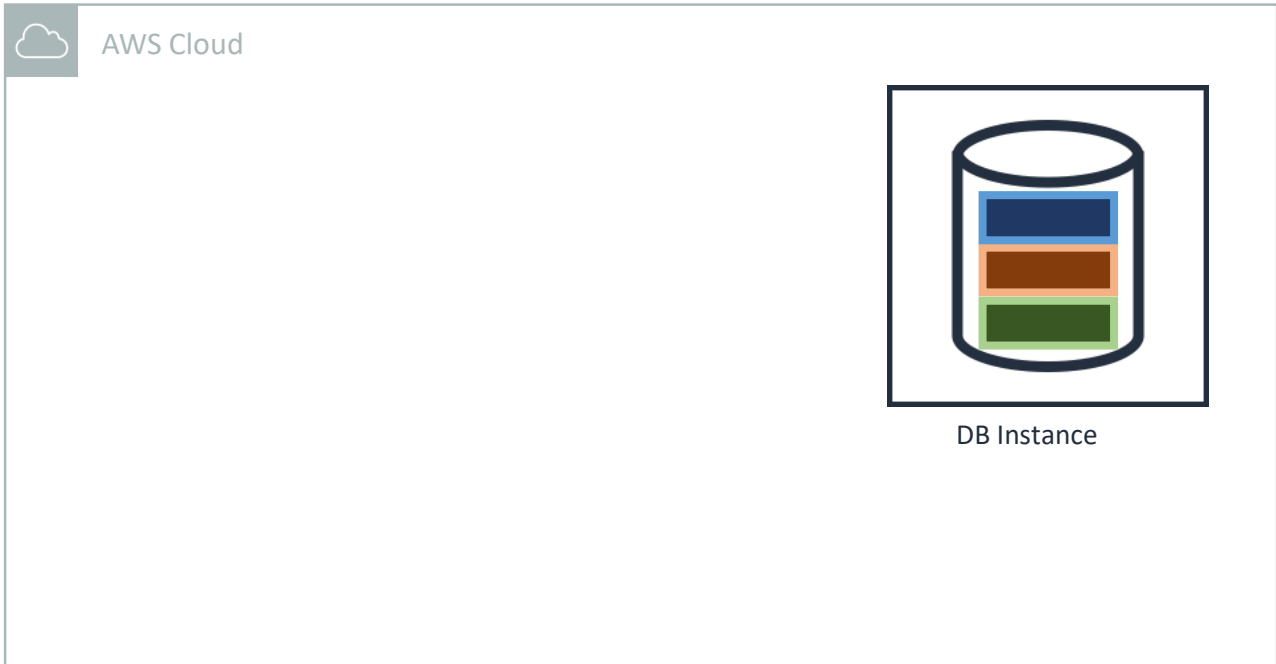
Alice



Bob



Amazon API  
Gateway





# Underneath Everything, there is Hope

Serverless computing is our opportunity to implement **efficient IFC**

Perform IFC tracking at **function boundaries**



# Serverless+IFC

How do we exploit Serverless to design an efficient IFC system?



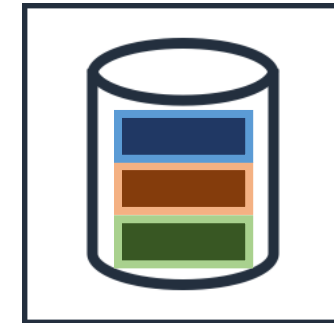
Alice



Bob



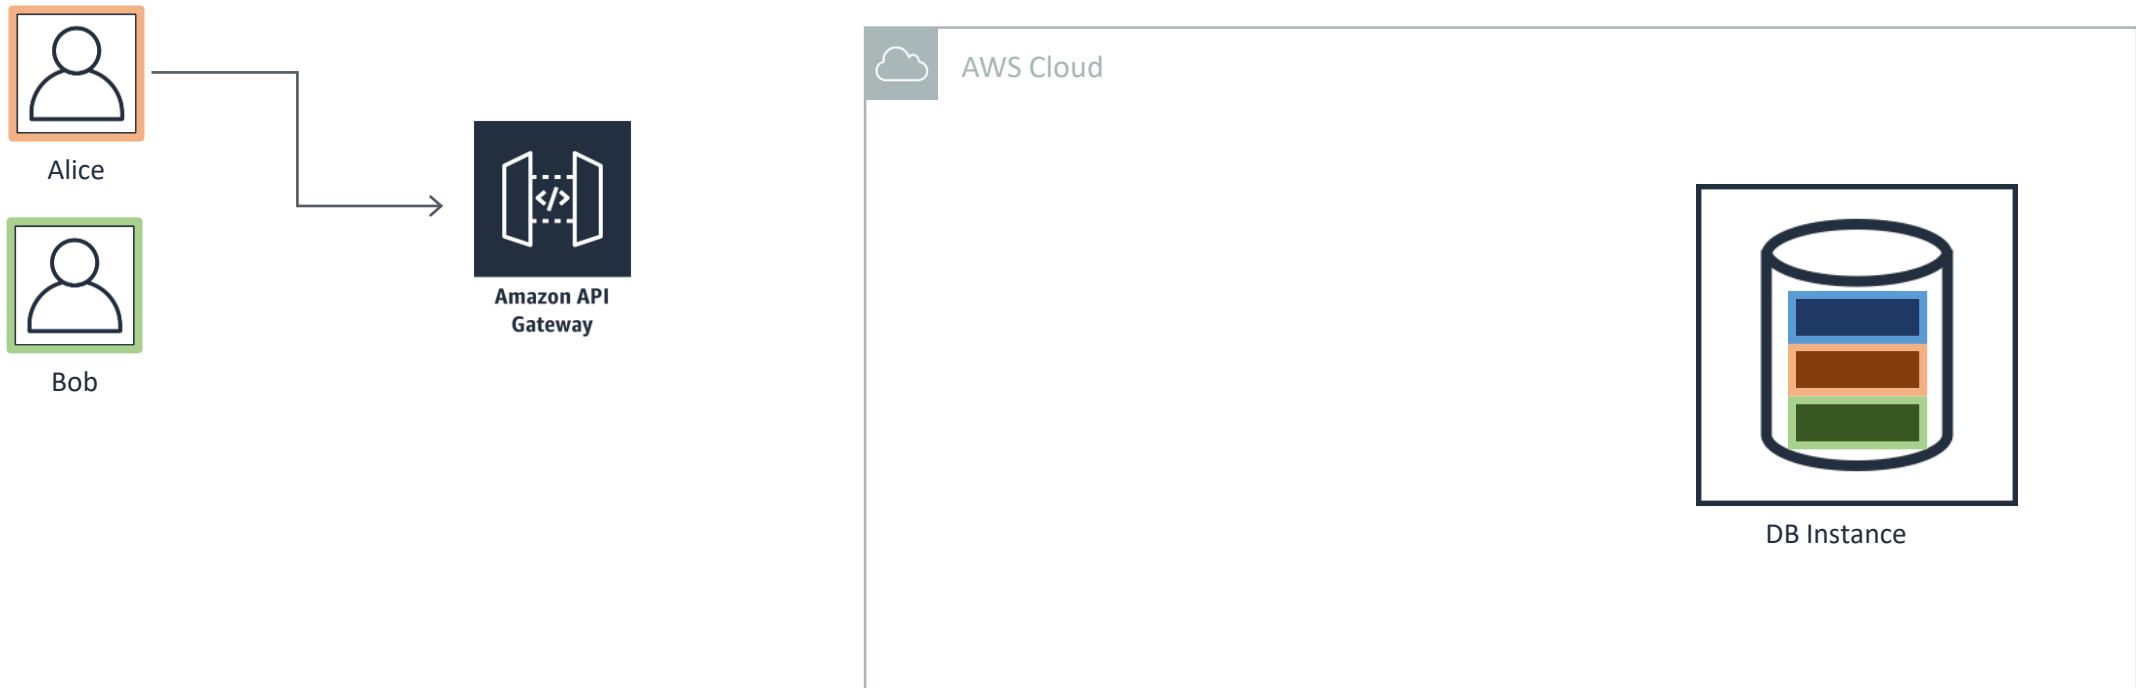
Amazon API  
Gateway



DB Instance

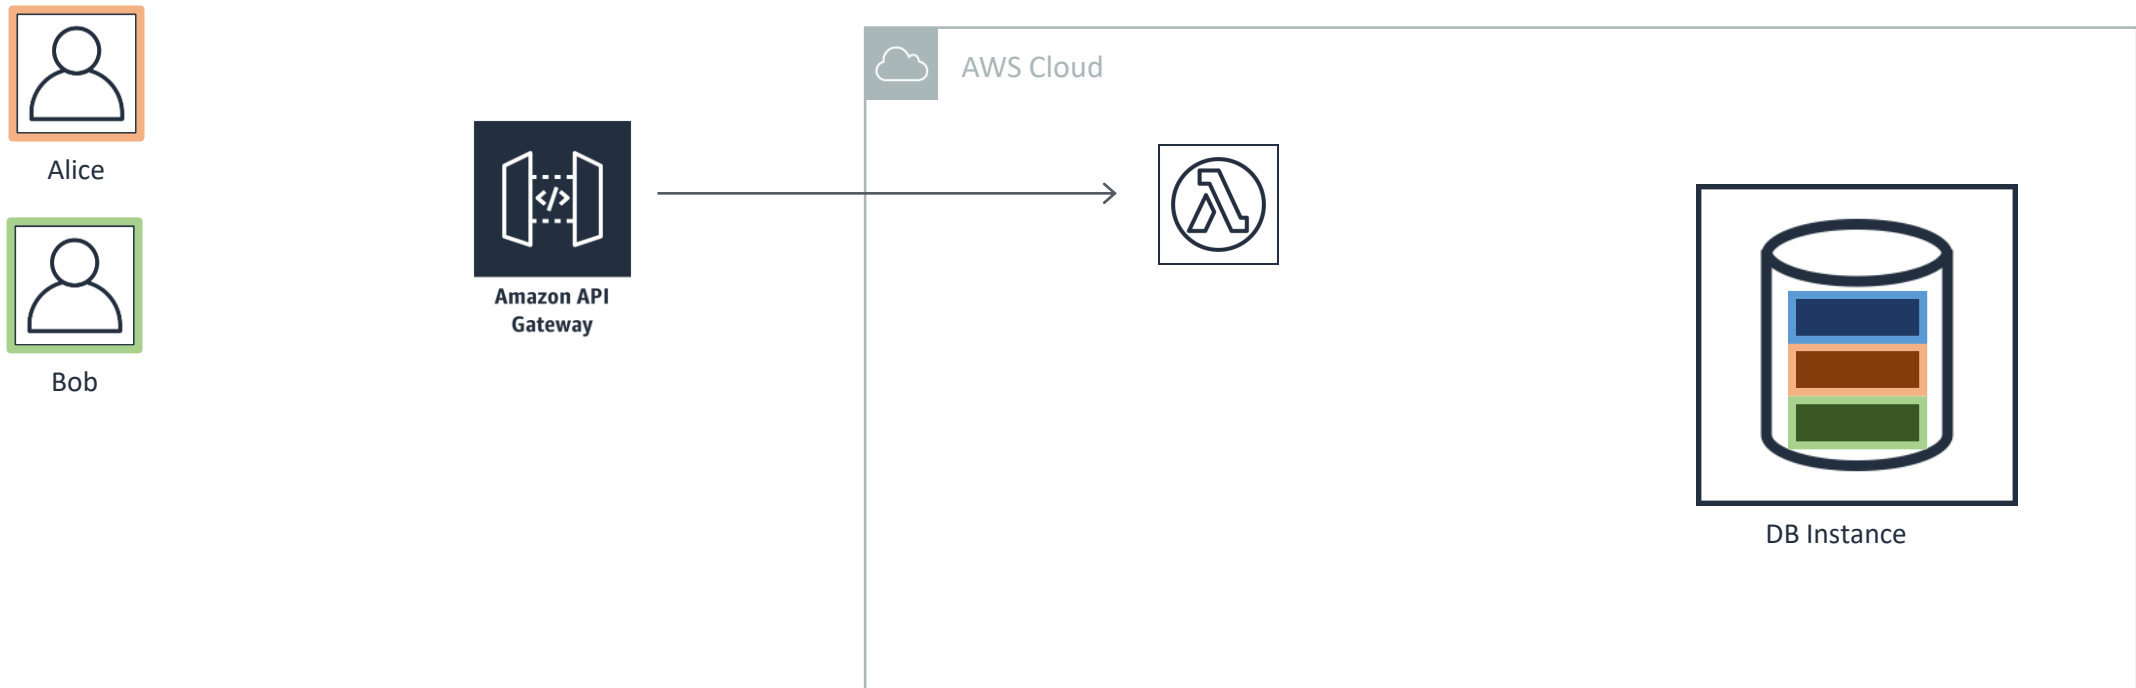
# Serverless+IFC

Alice makes a request



# Serverless+IFC

The Serverless framework spawns a function to serve the request



# Serverless+IFC

Functions may spawn additional functions



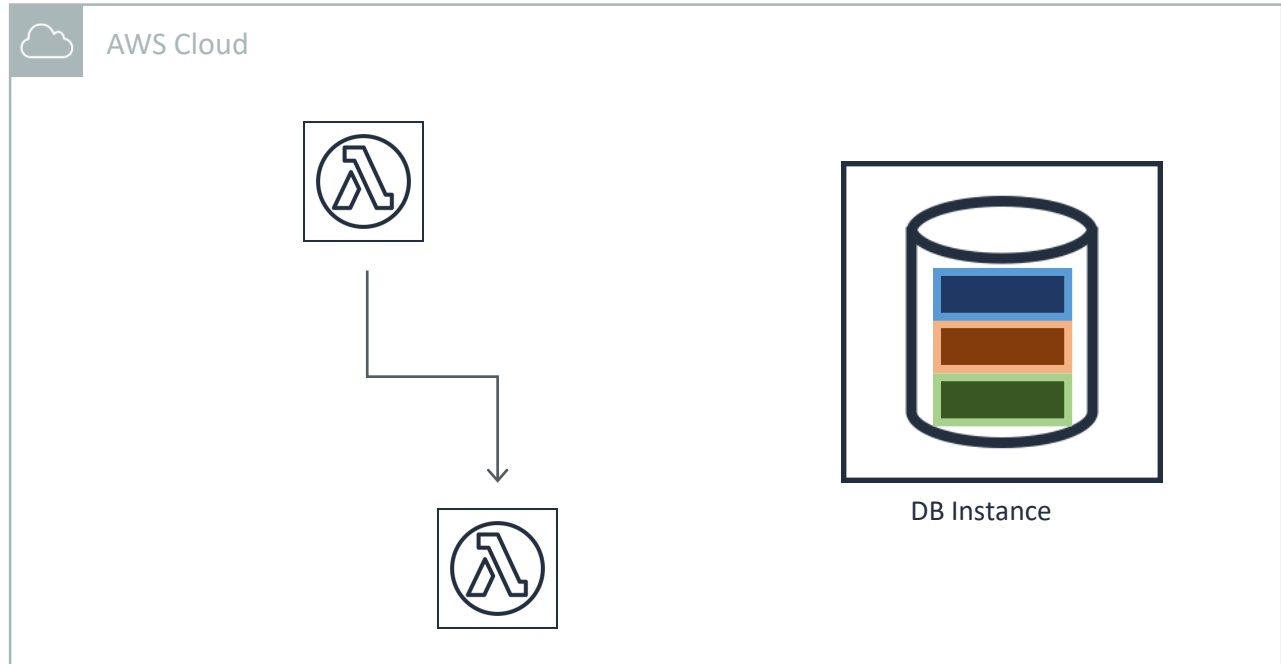
Alice



Bob



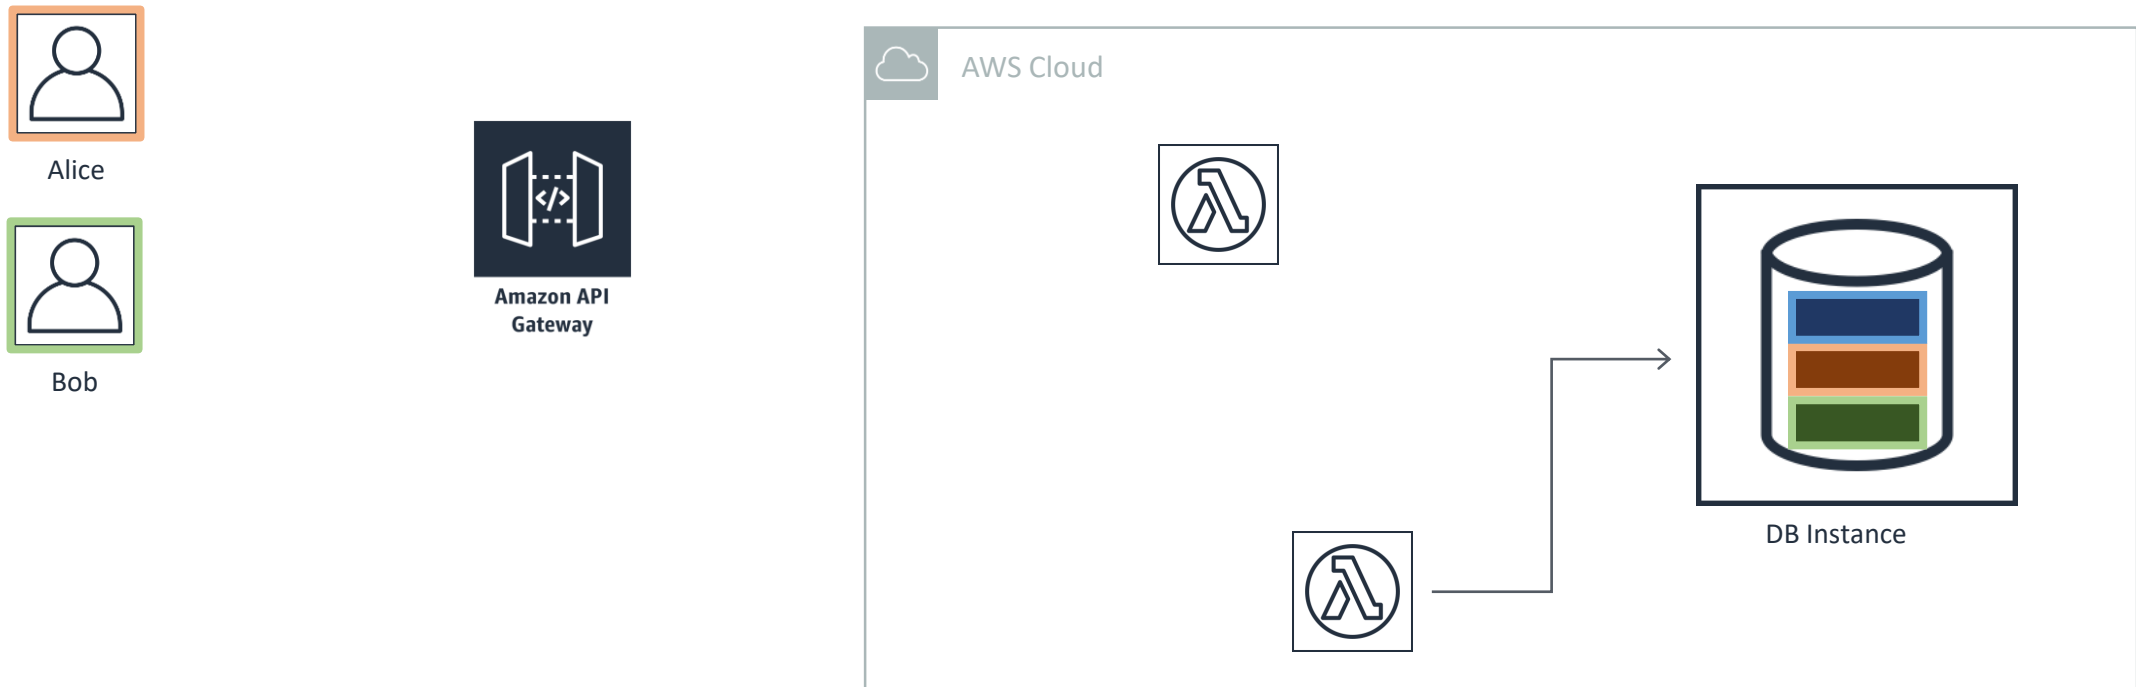
Amazon API Gateway





# Serverless+IFC

Functions interact with DB



# Serverless+IFC

Function reads **labeled** data and is marked with the **label**



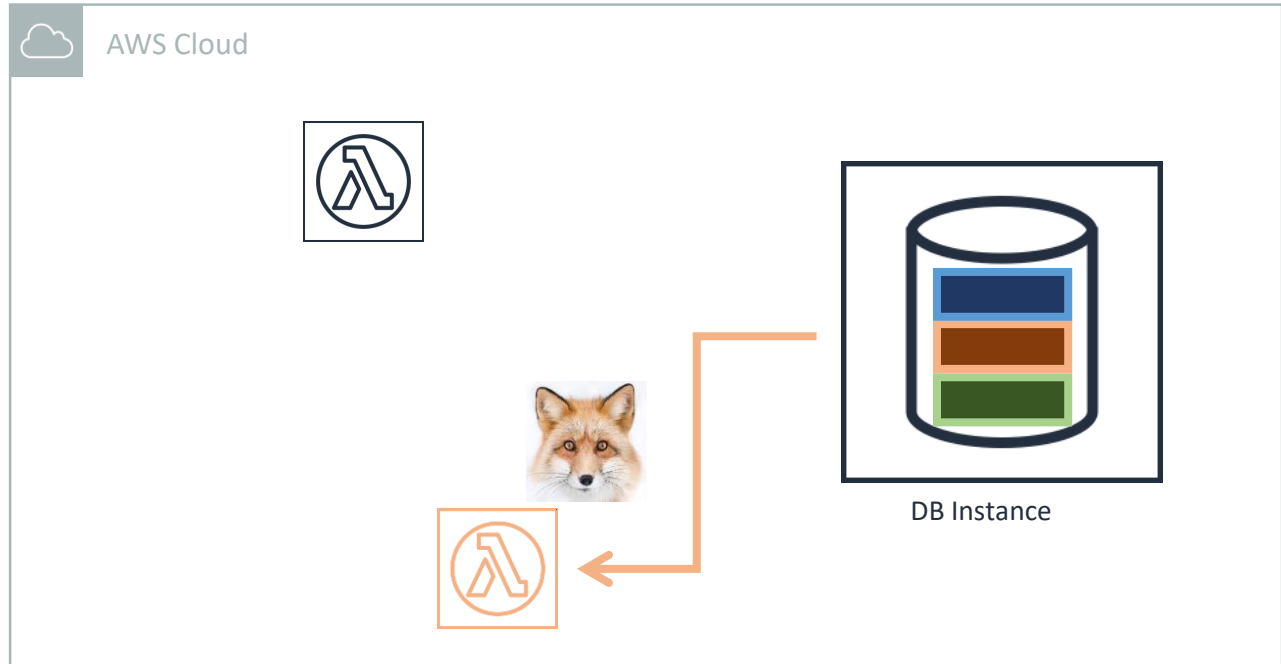
Alice



Bob



Amazon API Gateway



# Serverless+IFC

**Label** is propagated across function calls



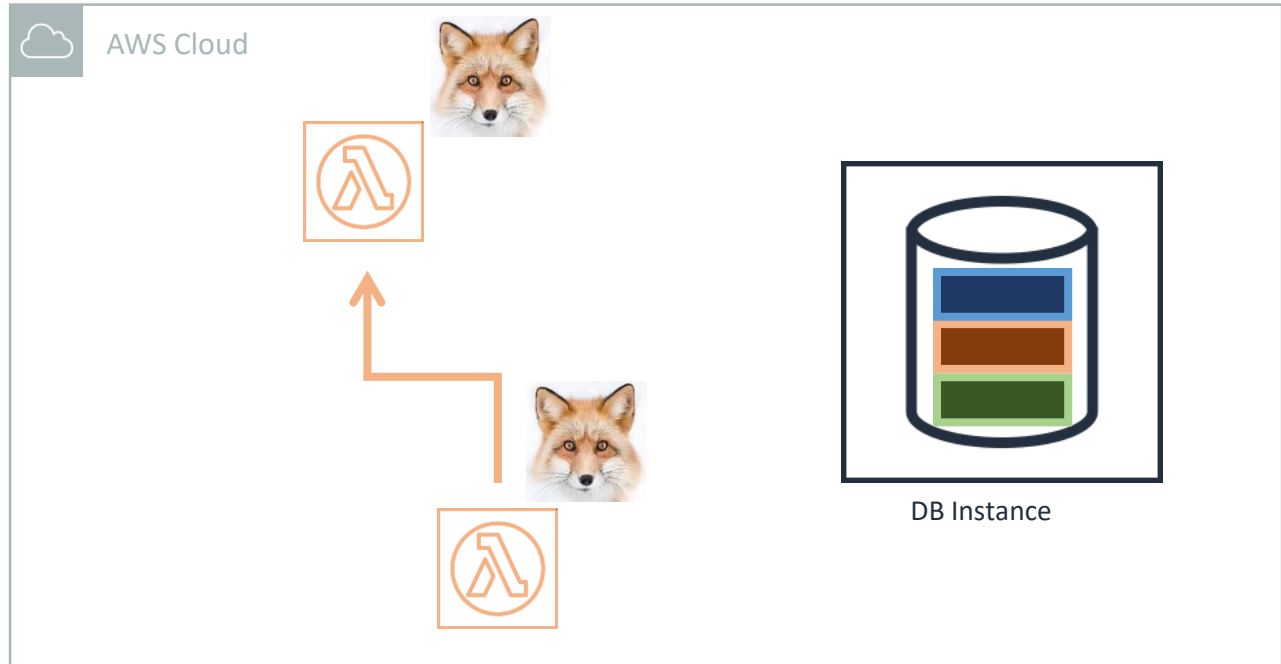
Alice



Bob

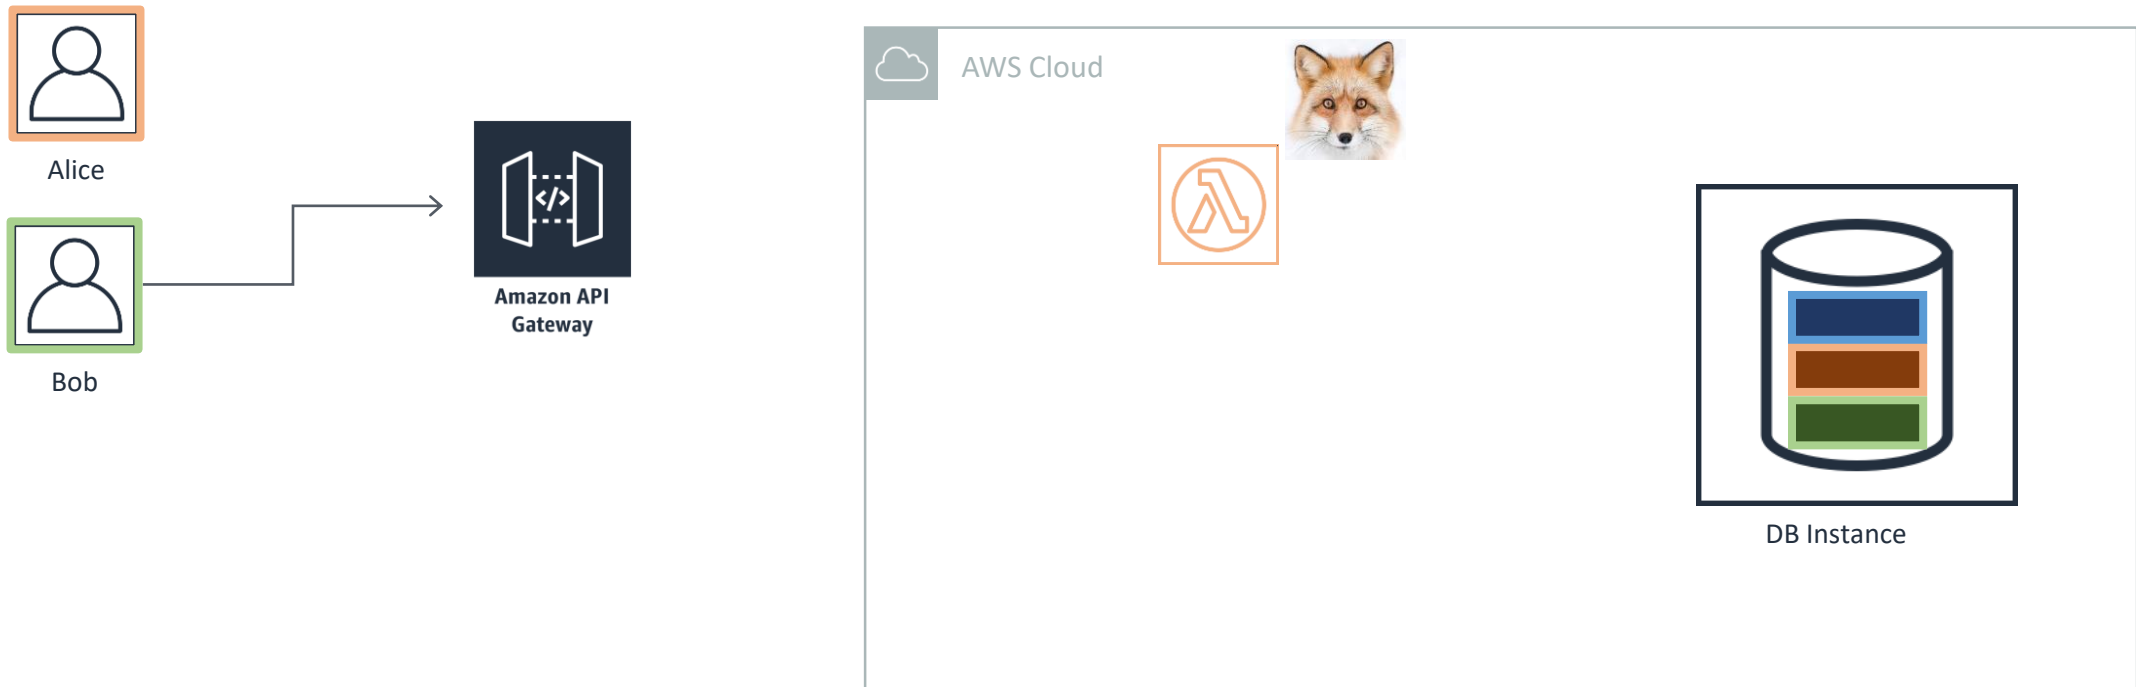


Amazon API Gateway



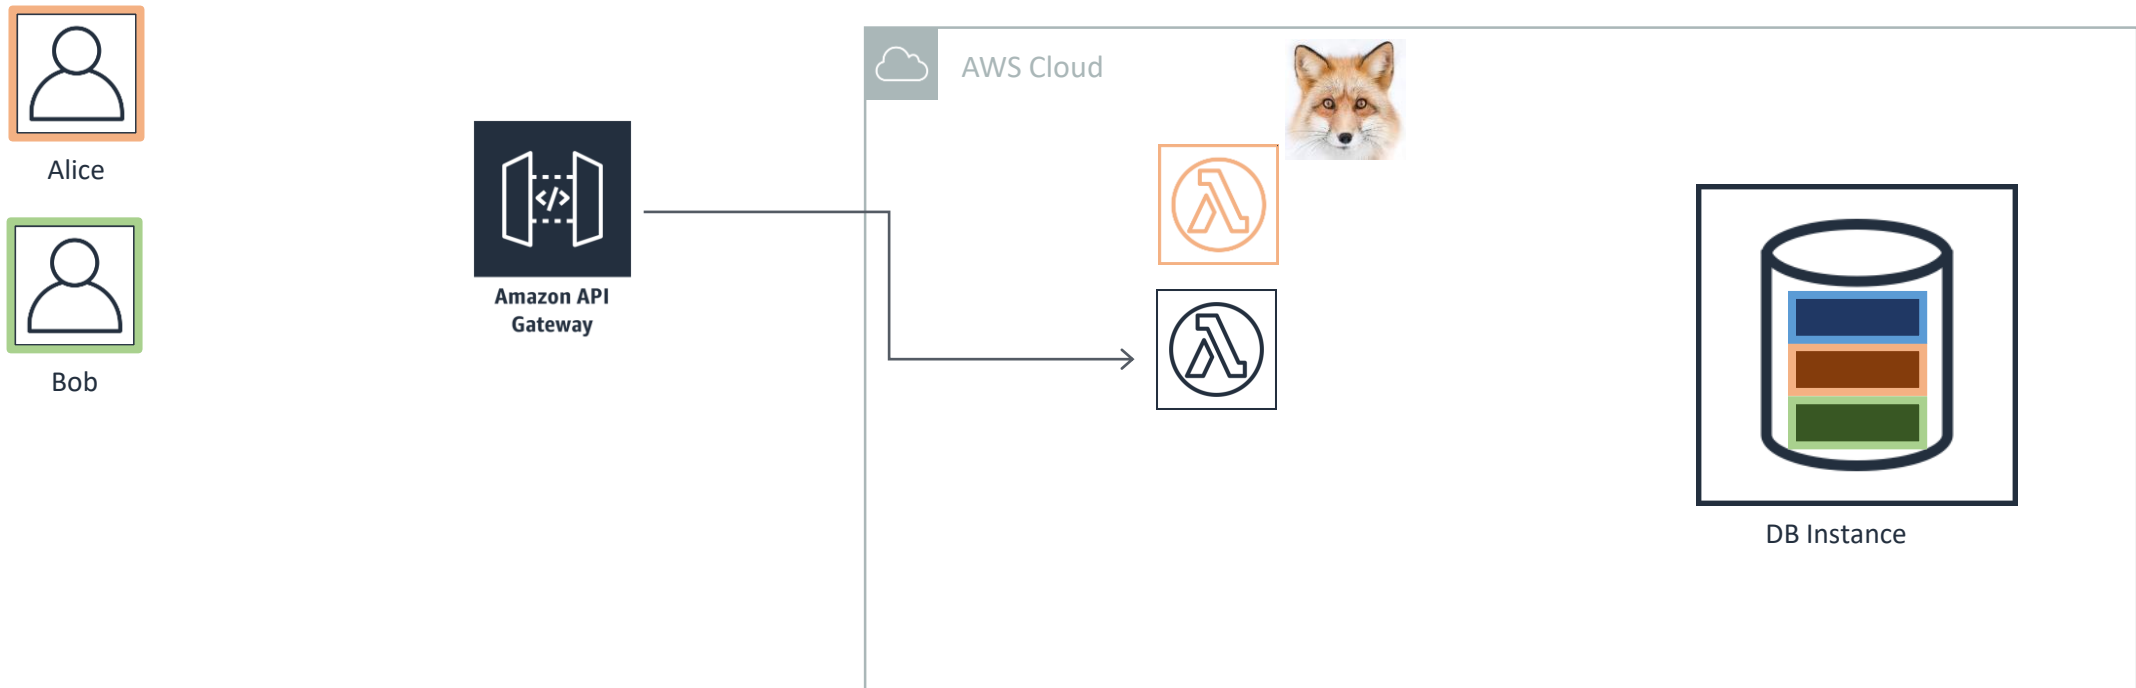
# Serverless+IFC

Other requests may occur in parallel



# Serverless+IFC

Parallel requests spawn new functions, that are independent from one another



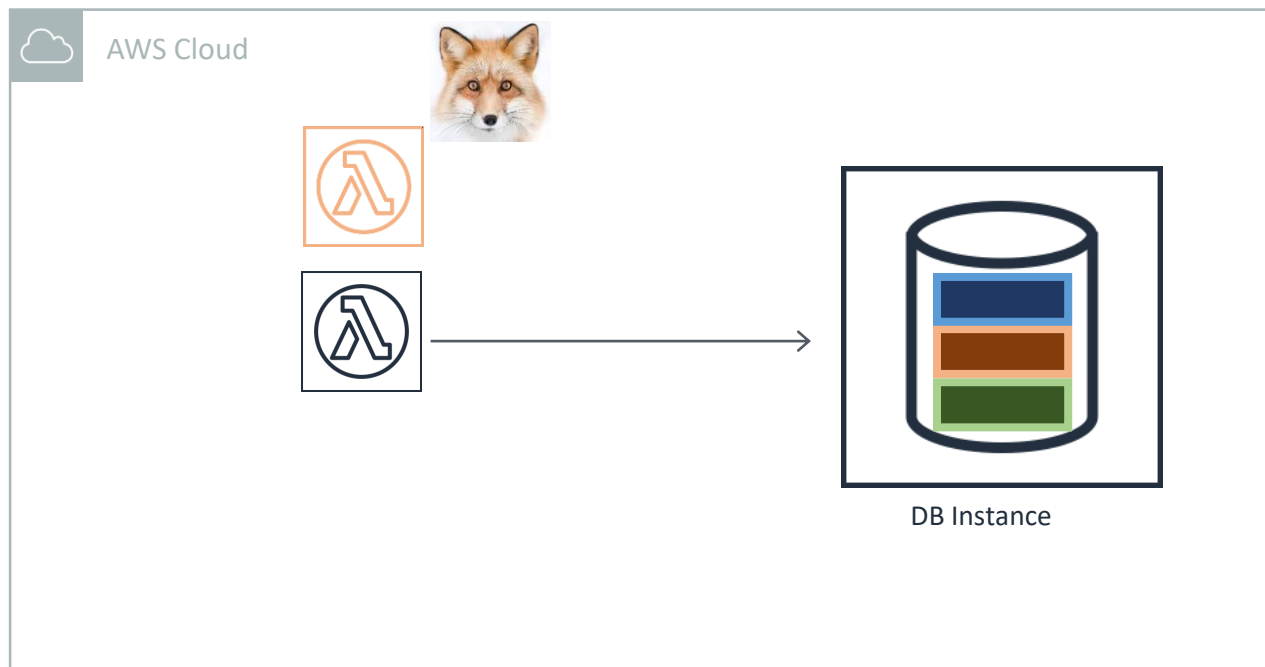
# Serverless+IFC



Alice



Bob



# Serverless+IFC

Function reads **labeled** data and is marked with the **label**



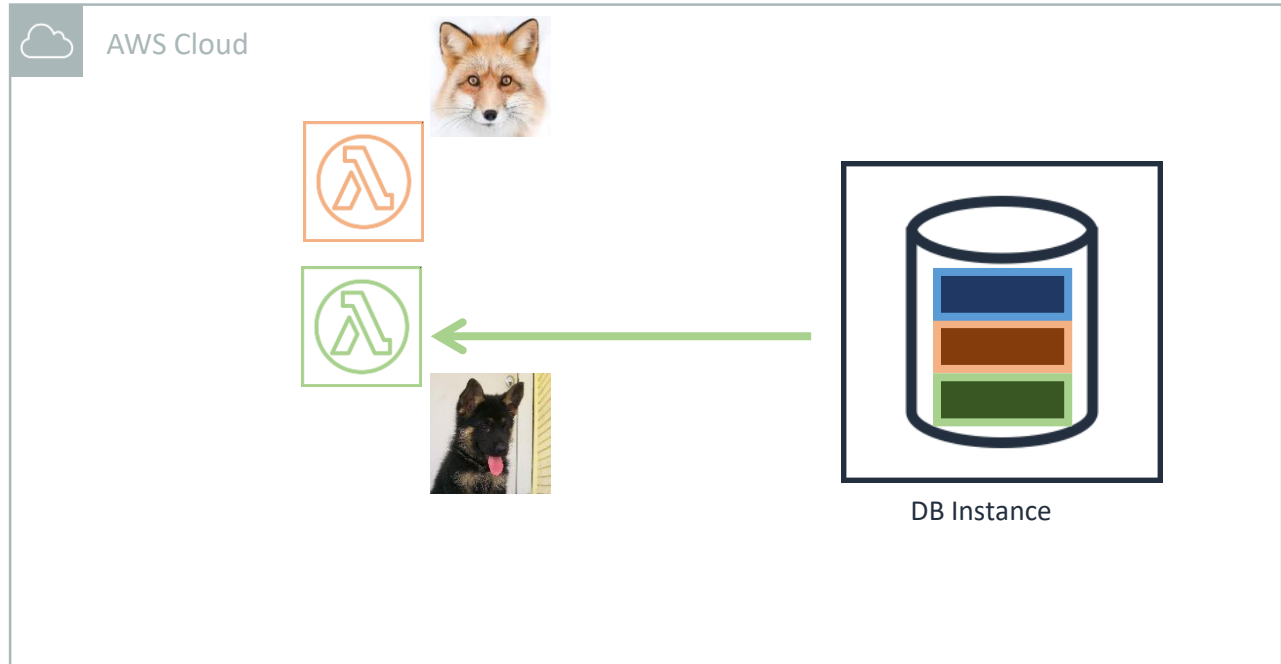
Alice



Bob

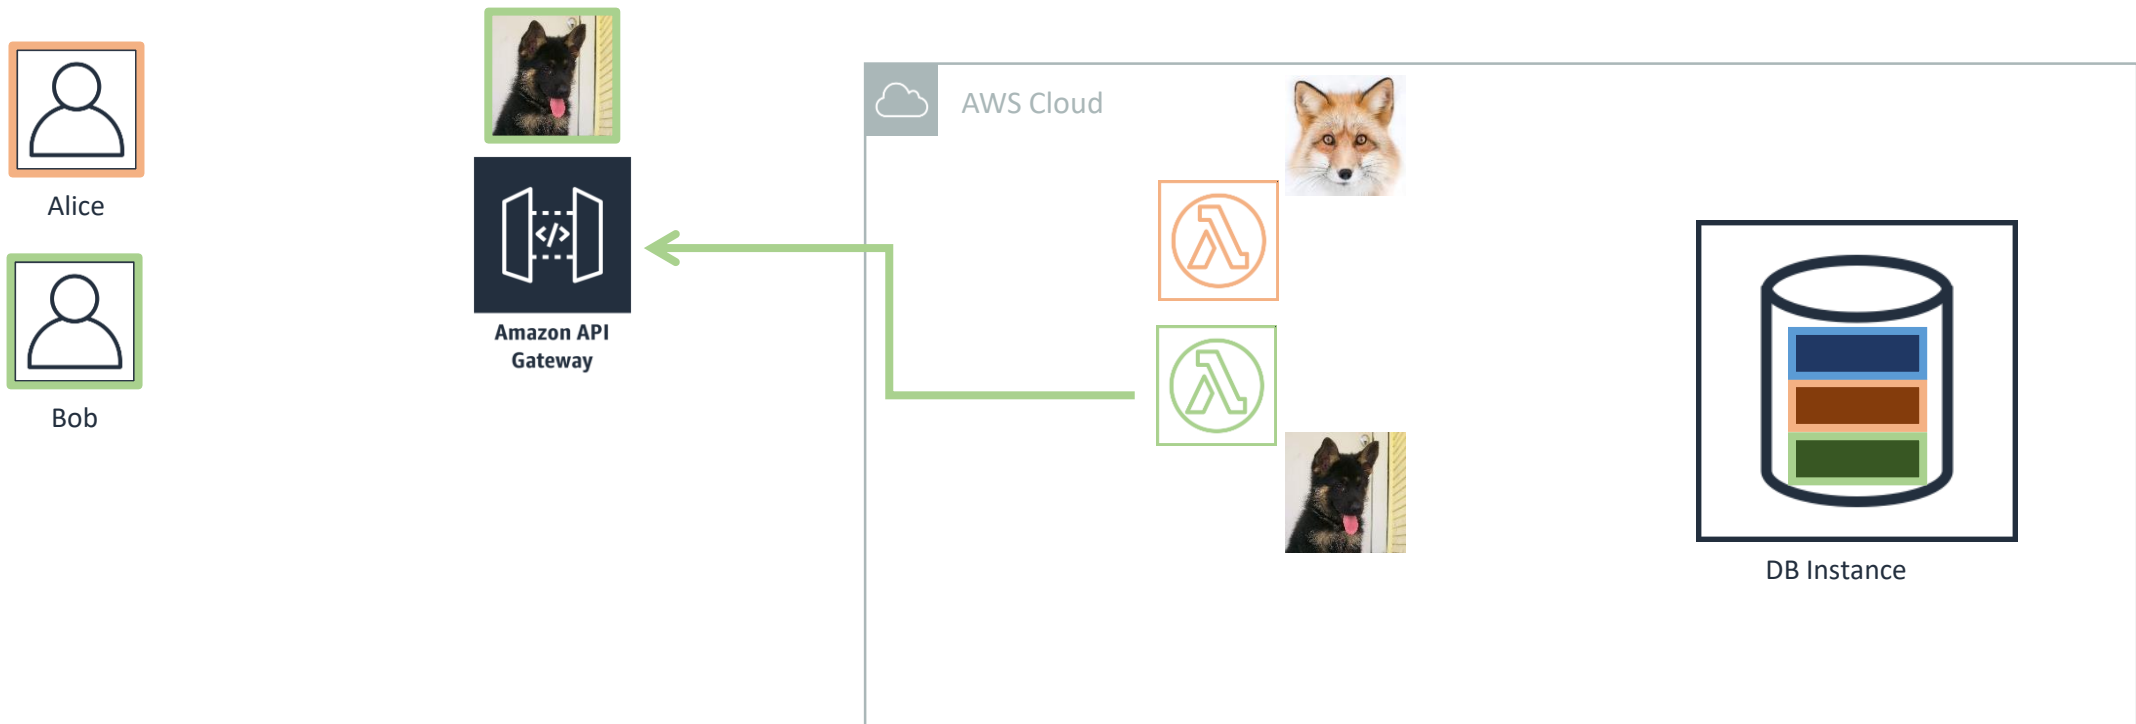


Amazon API  
Gateway



# Serverless+IFC

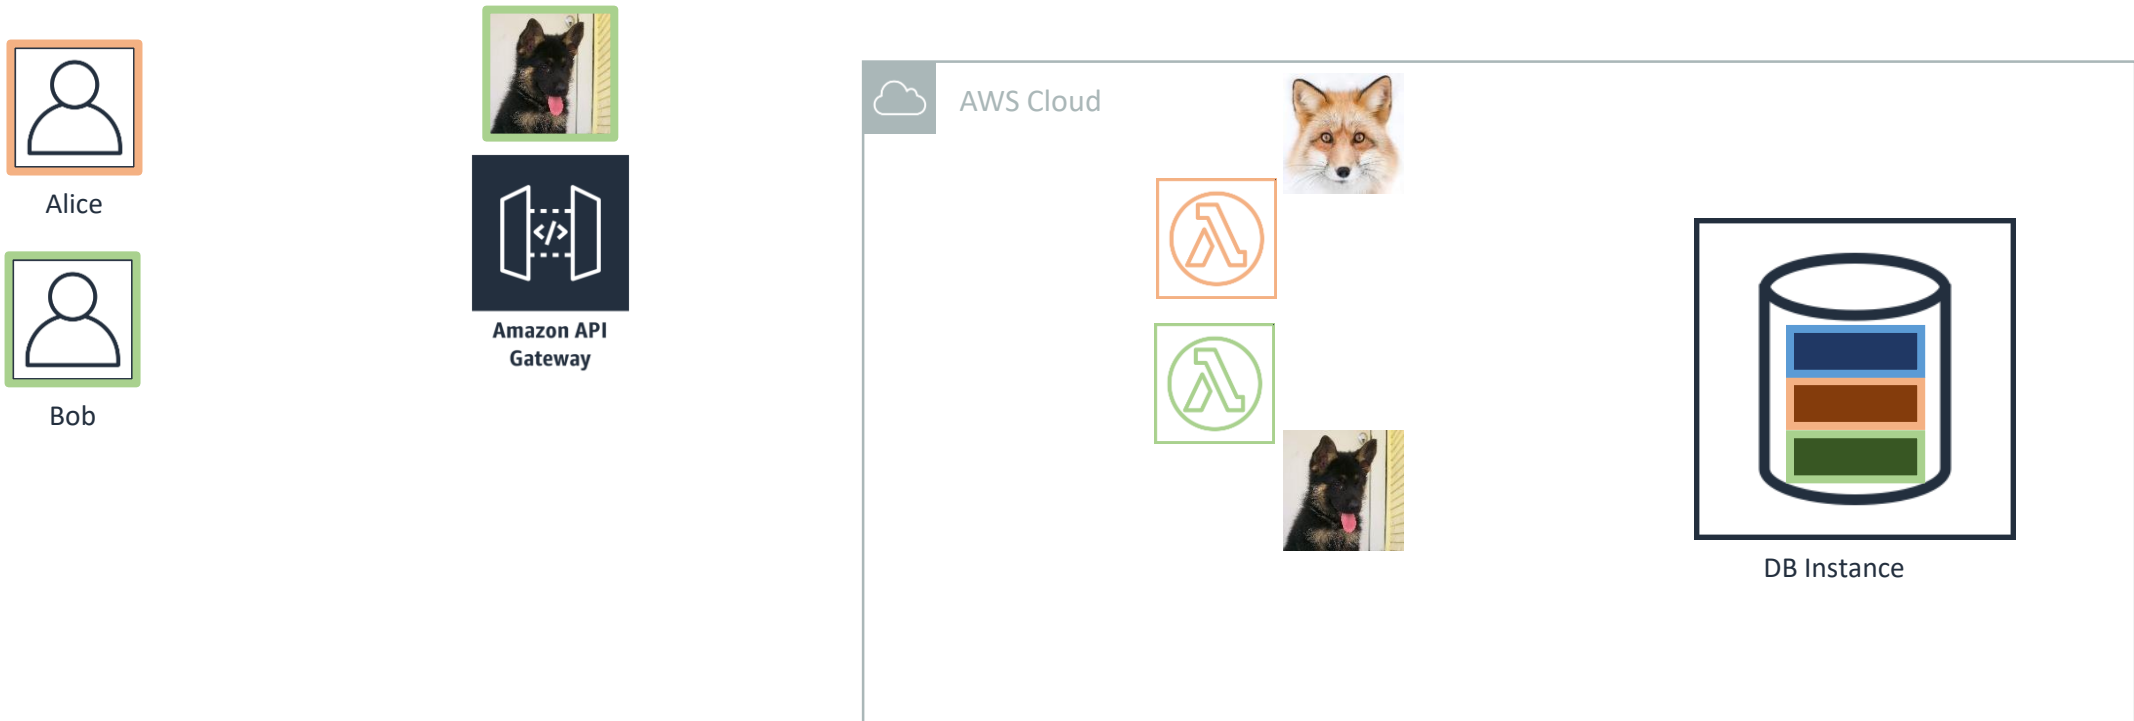
Function returns result





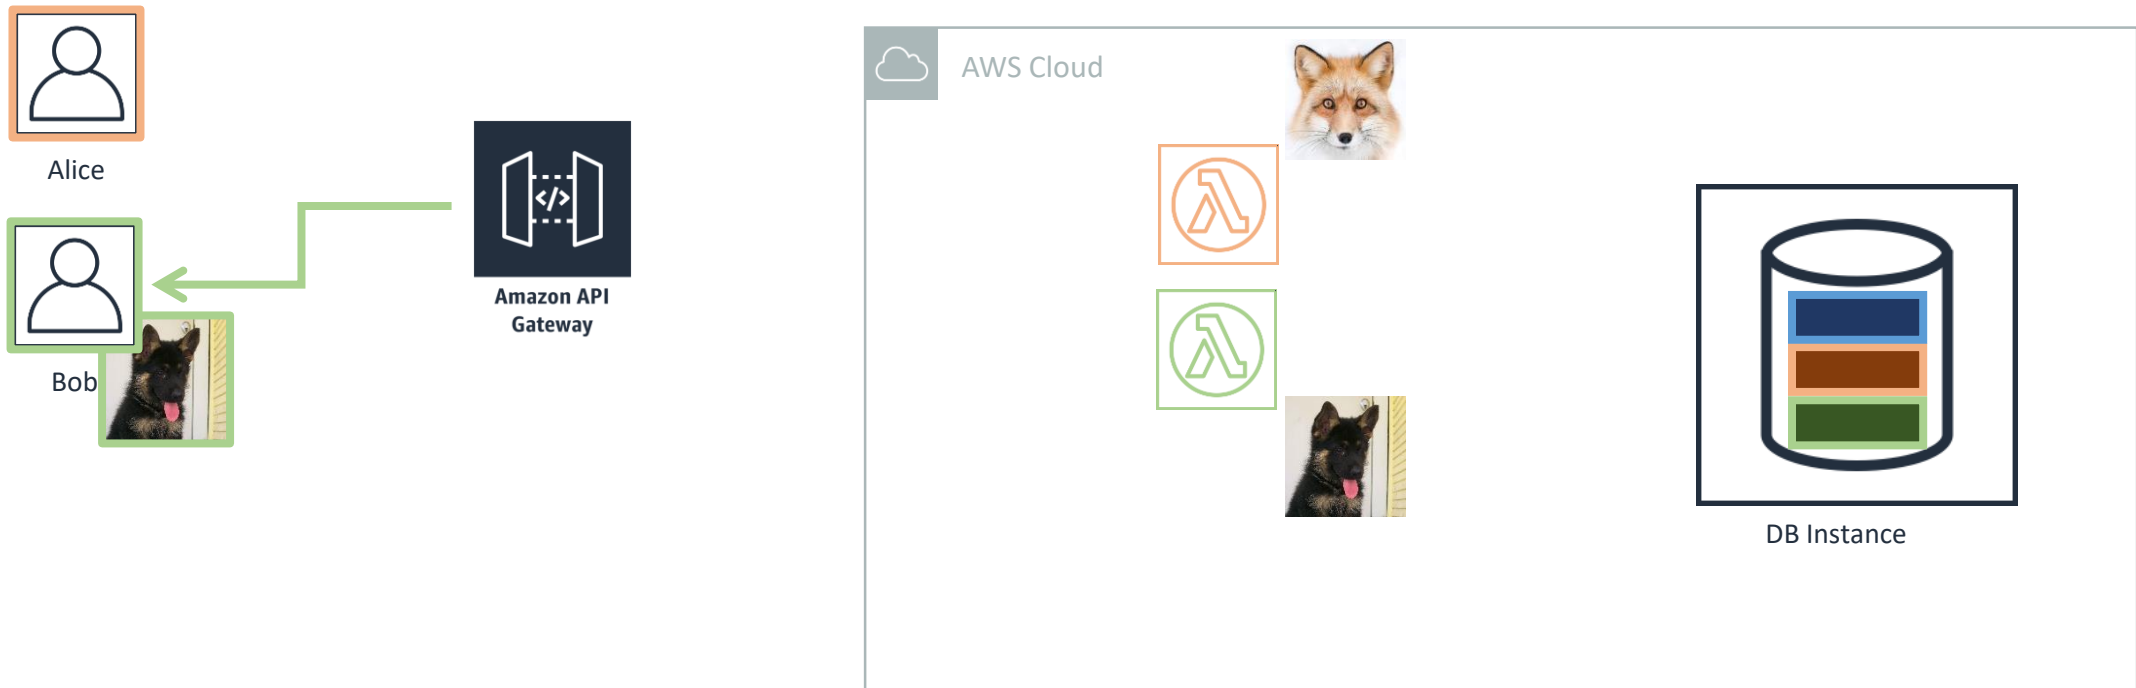
# Serverless+IFC

Before the result is sent to the user, we can check that the security policy is not violated

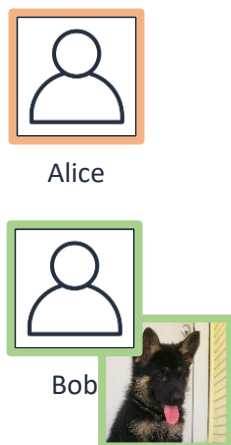


# Serverless+IFC

And send the response if it does not violate the security policy

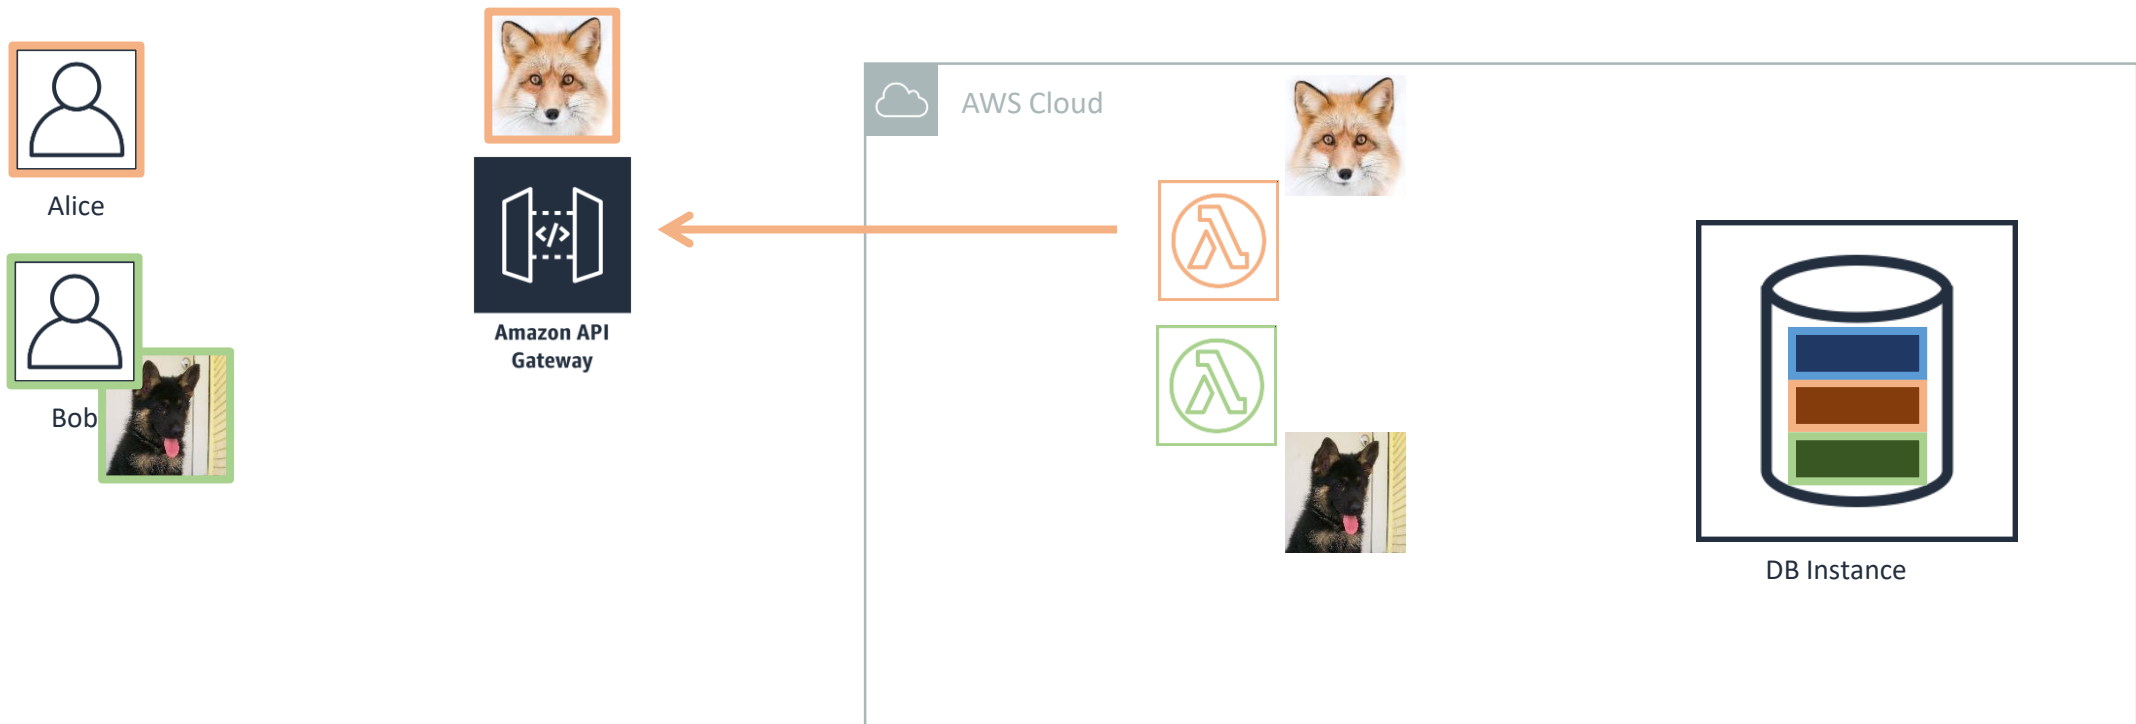


# Serverless+IFC



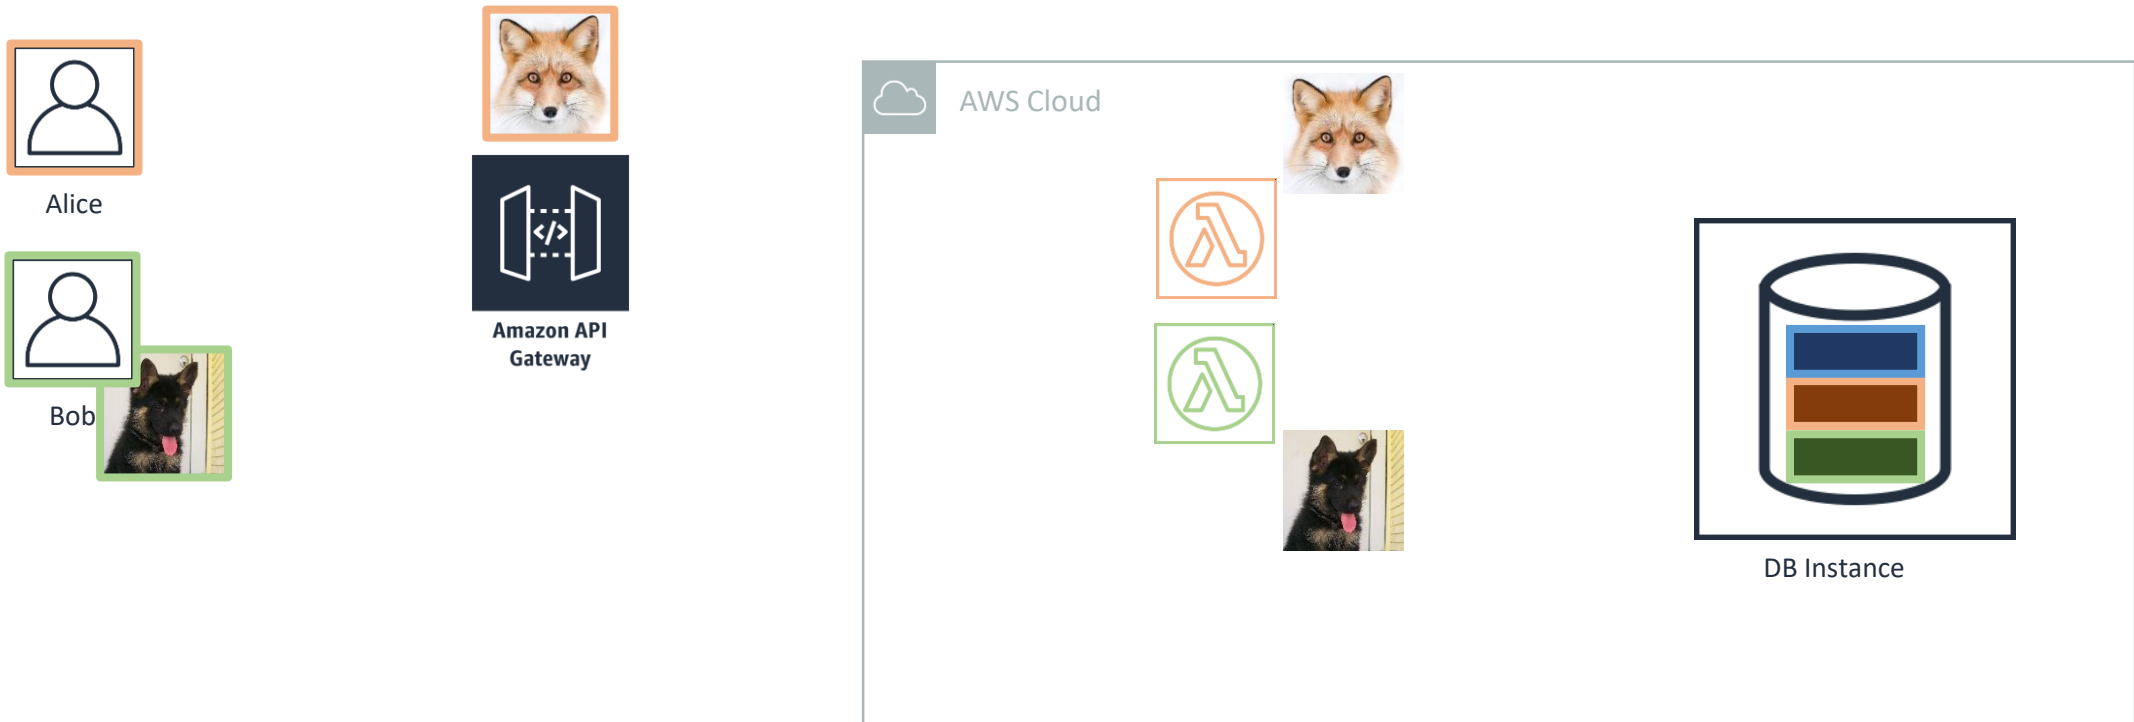
# Serverless+IFC

Function returns result



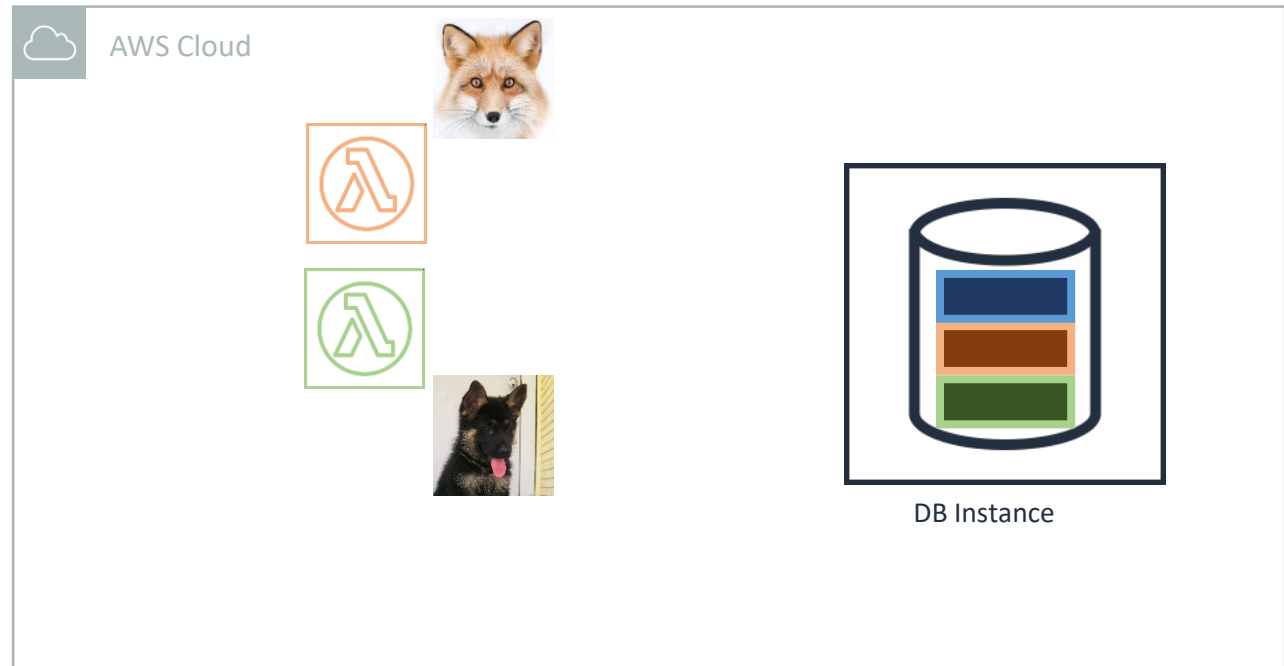
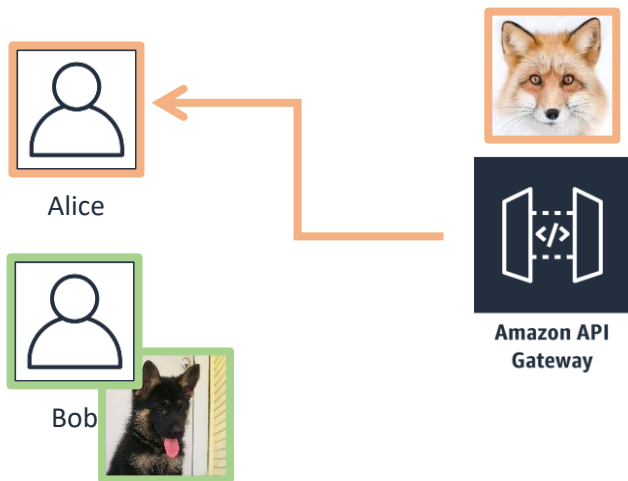
# Serverless+IFC

Before the result is sent to the user, we can check that the security policy is not violated



# Serverless+IFC

And send the response if it does not violate the security policy

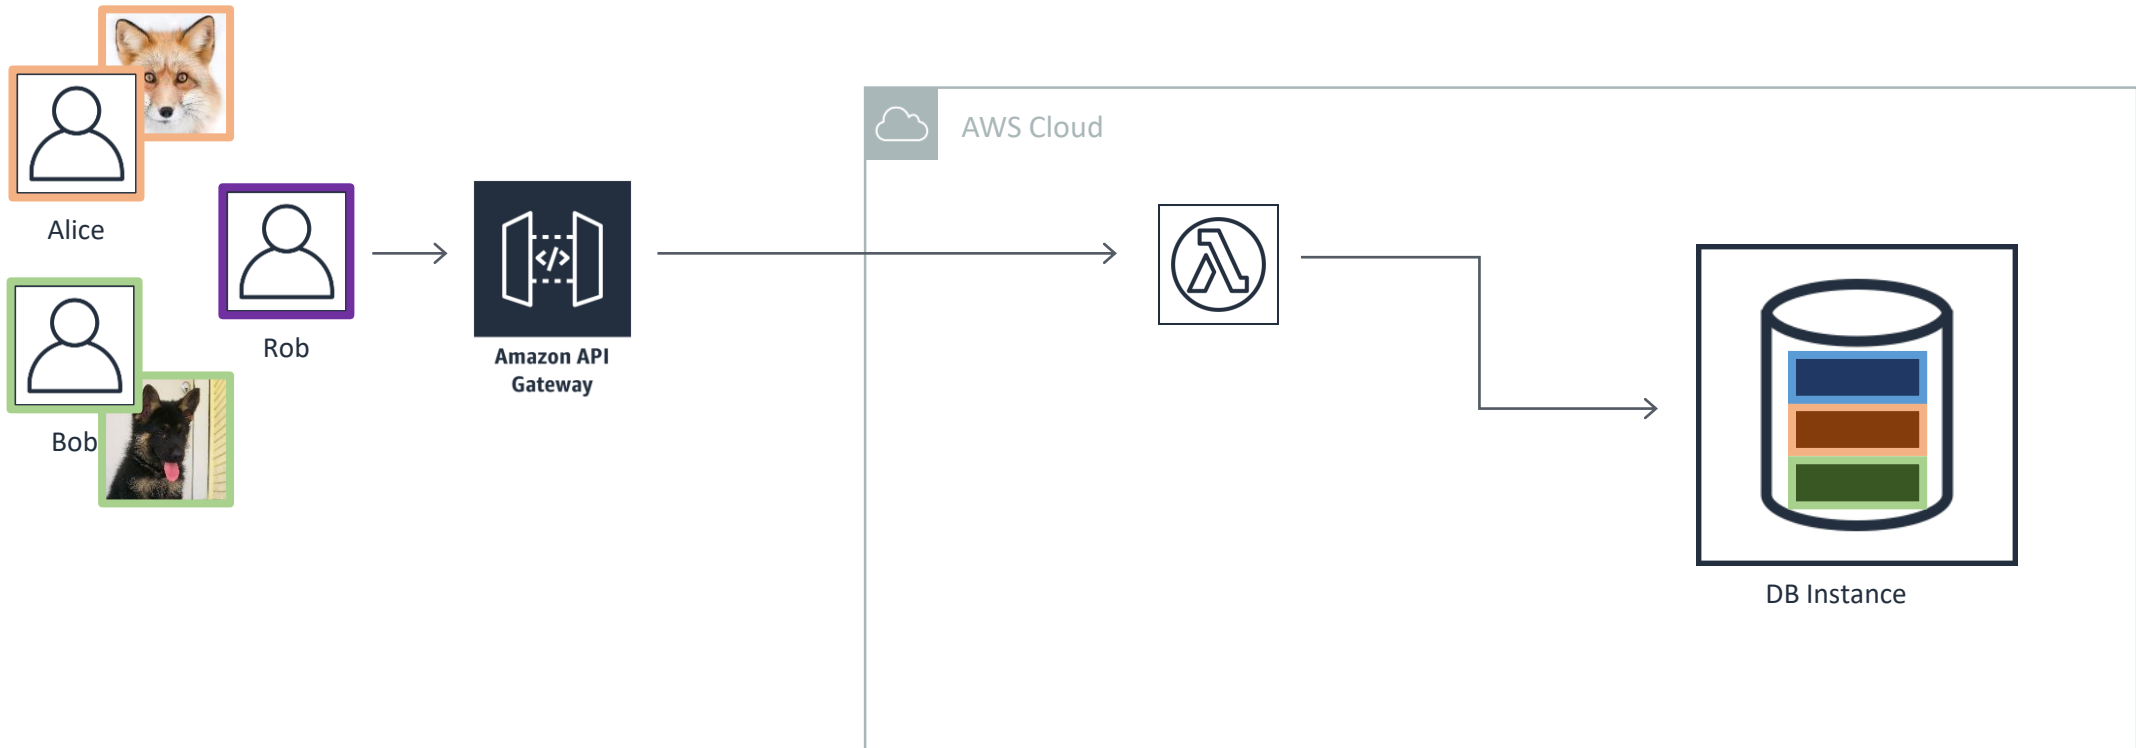


# Serverless+IFC



# Serverless+IFC

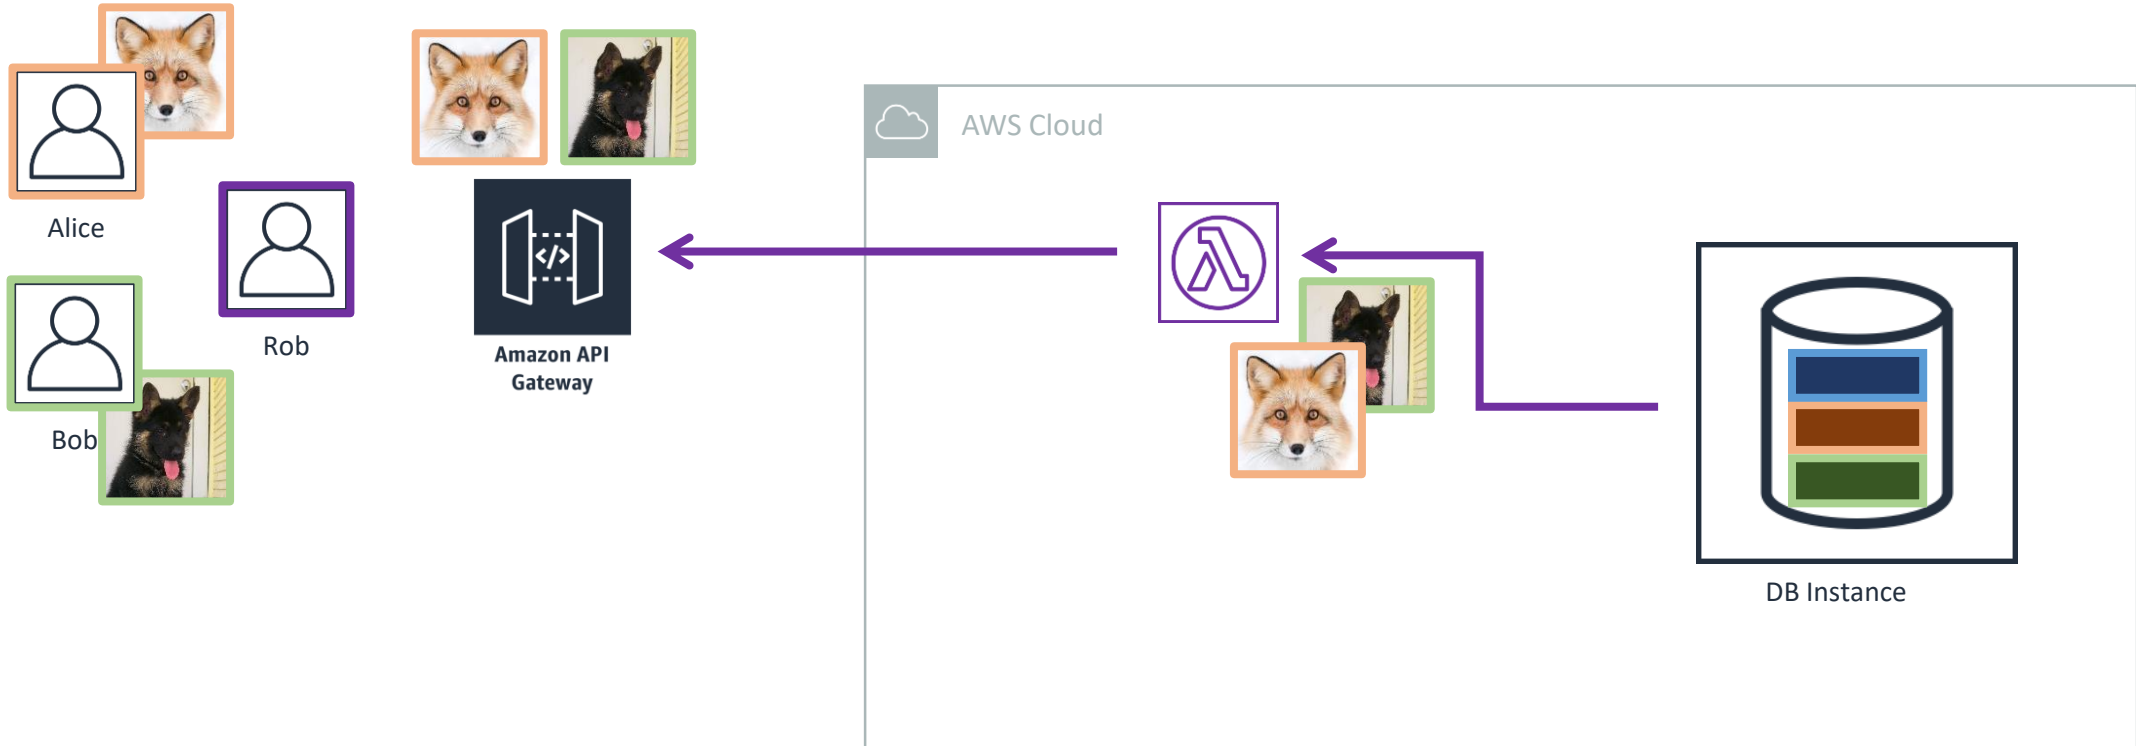
Security policy might have a hierarchy – making shared dbs unavoidable





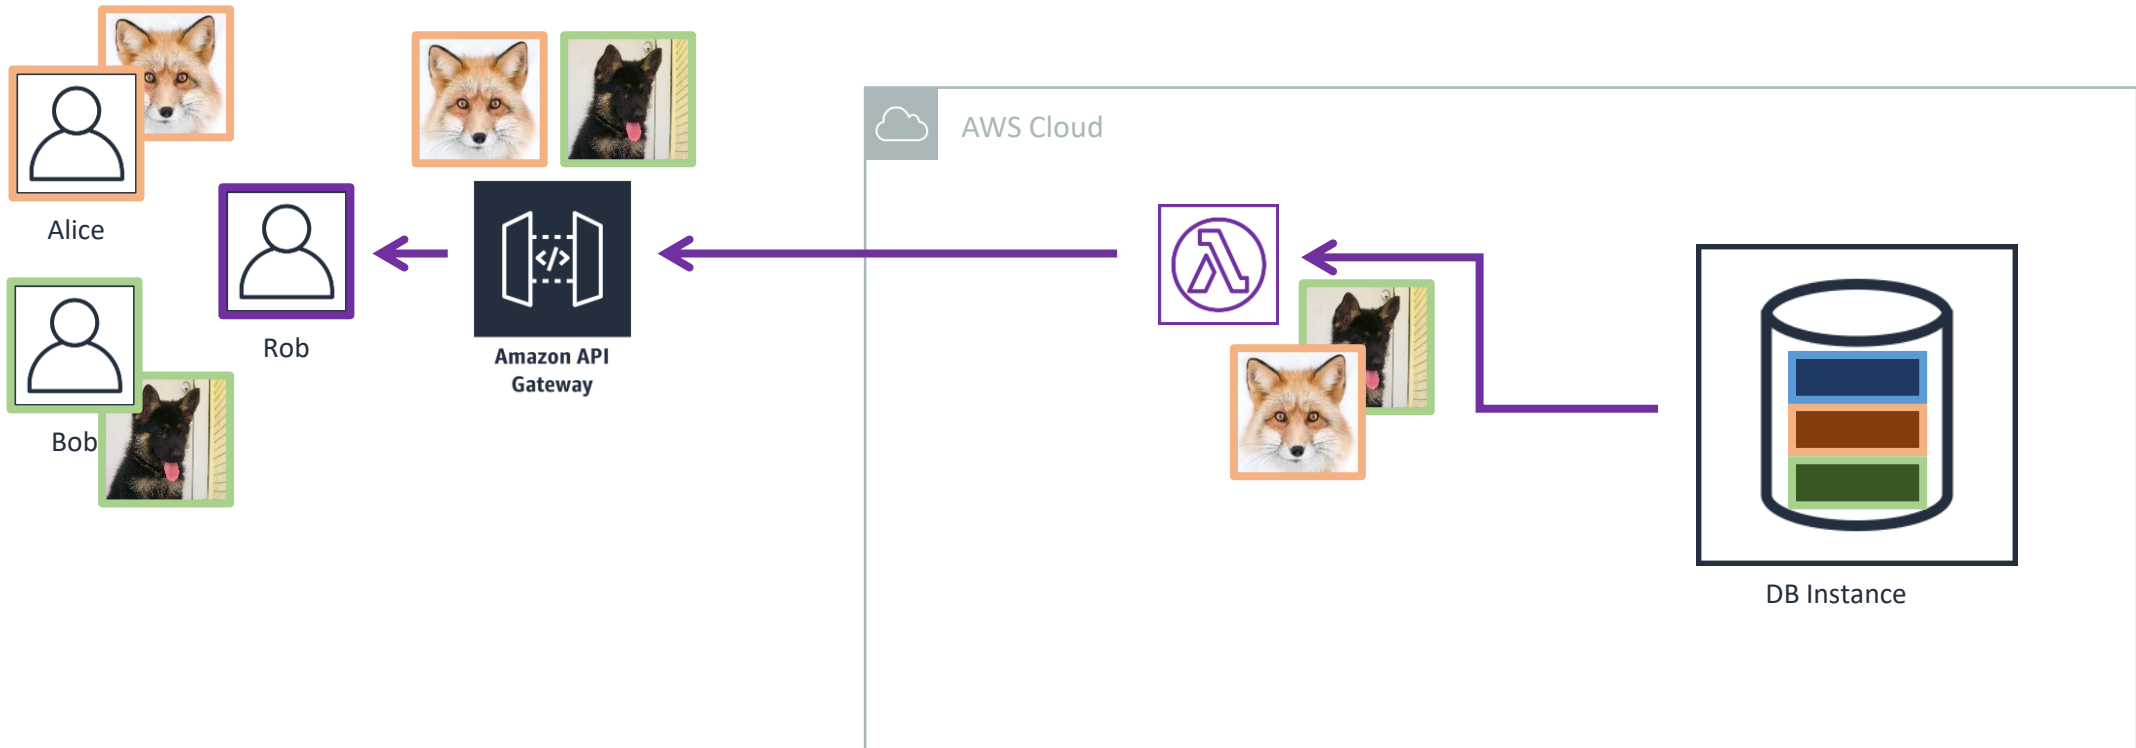
# Serverless+IFC

Security policy might have a hierarchy – making shared dbs unavoidable



# Serverless+IFC

Security policy might have a hierarchy – making shared dbs unavoidable



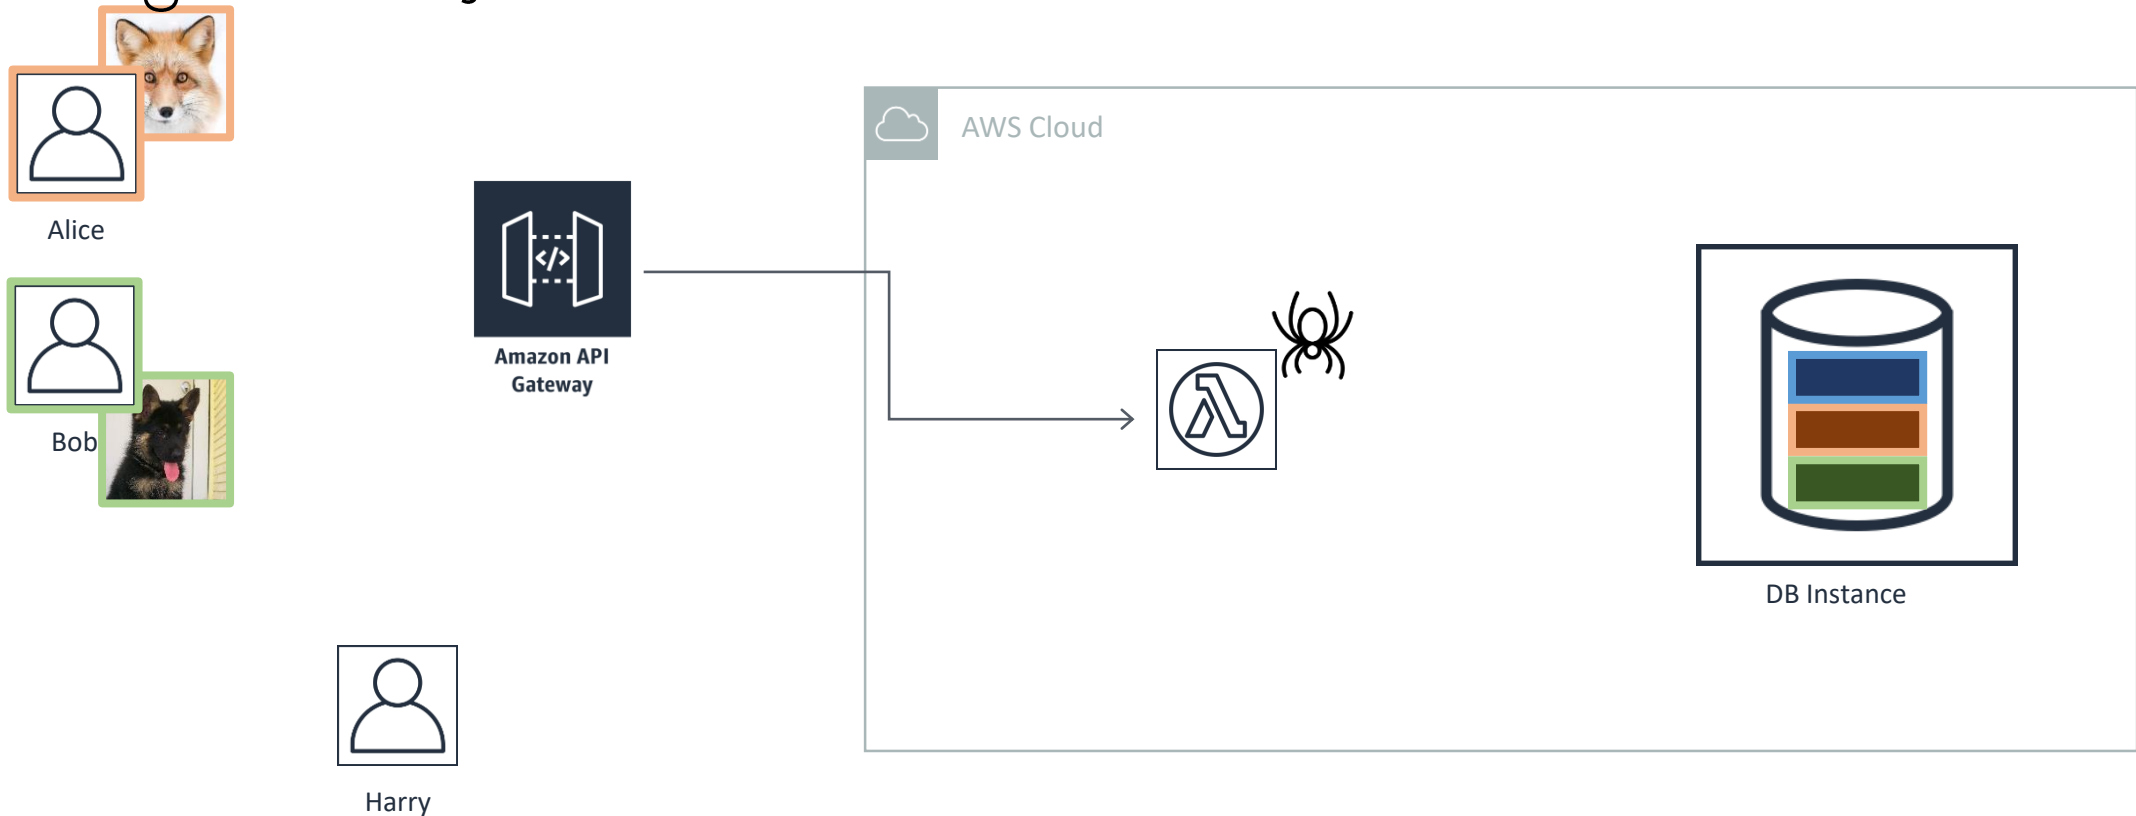
# Serverless+IFC

What happens when a hacker tries to steal data?



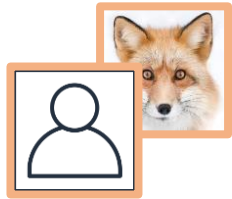
# Serverless+IFC

The hacker might exploit some vulnerability in the application (e.g. code injection)

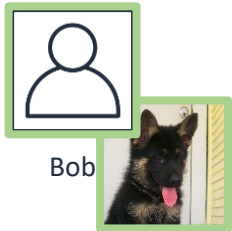


# Serverless+IFC

Trapeze guarantees security even if the underlying function code is completely compromised



Alice



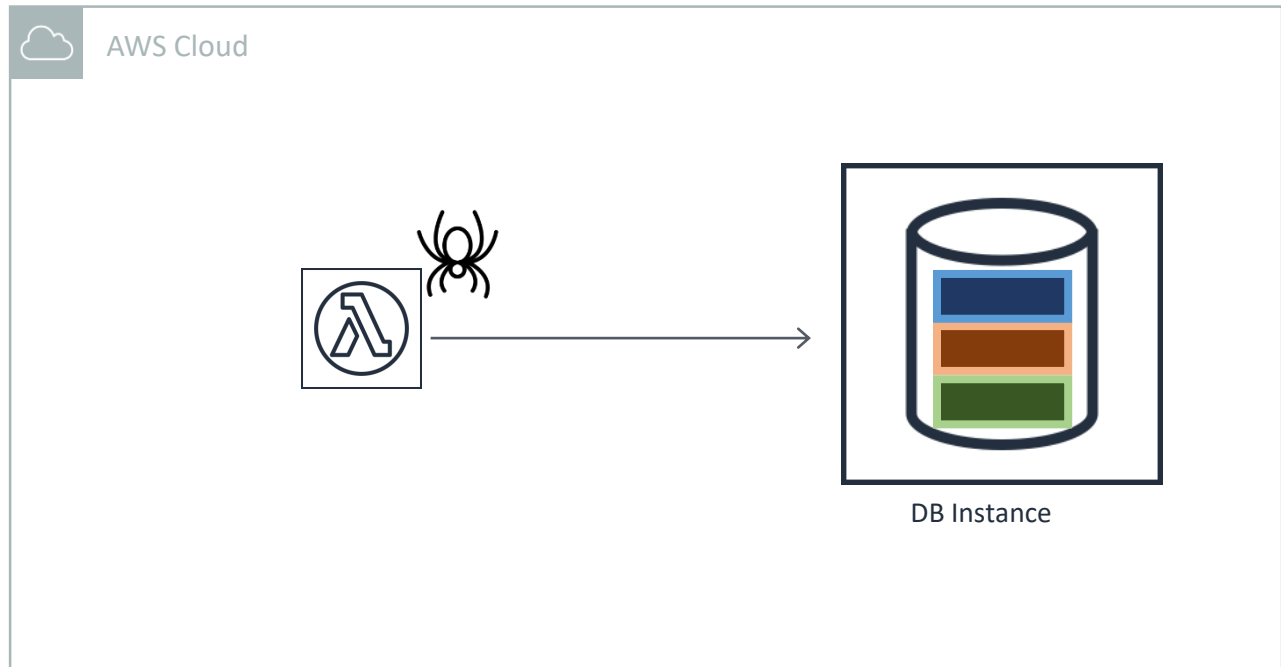
Bob



Harry

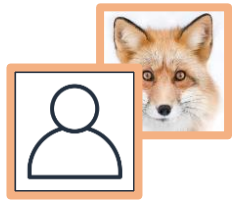


Amazon API Gateway

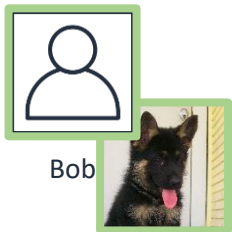


# Serverless+IFC

The malicious code may read a secret, and is marked with a **label**



Alice



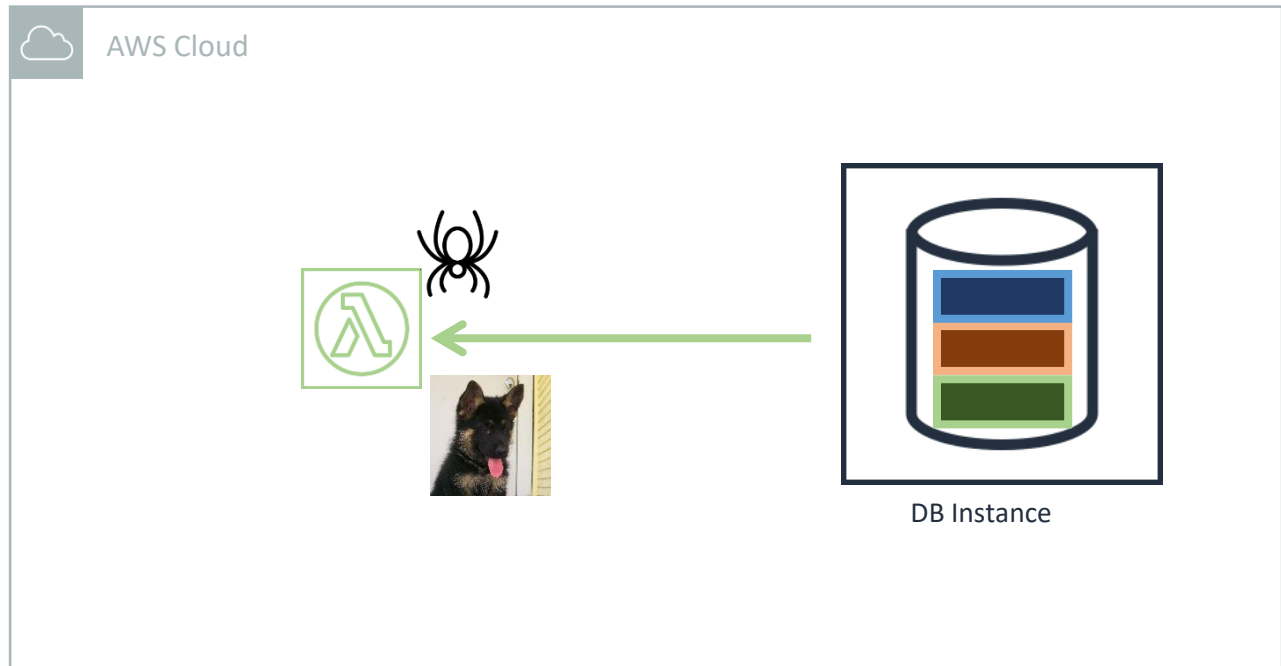
Bob



Harry

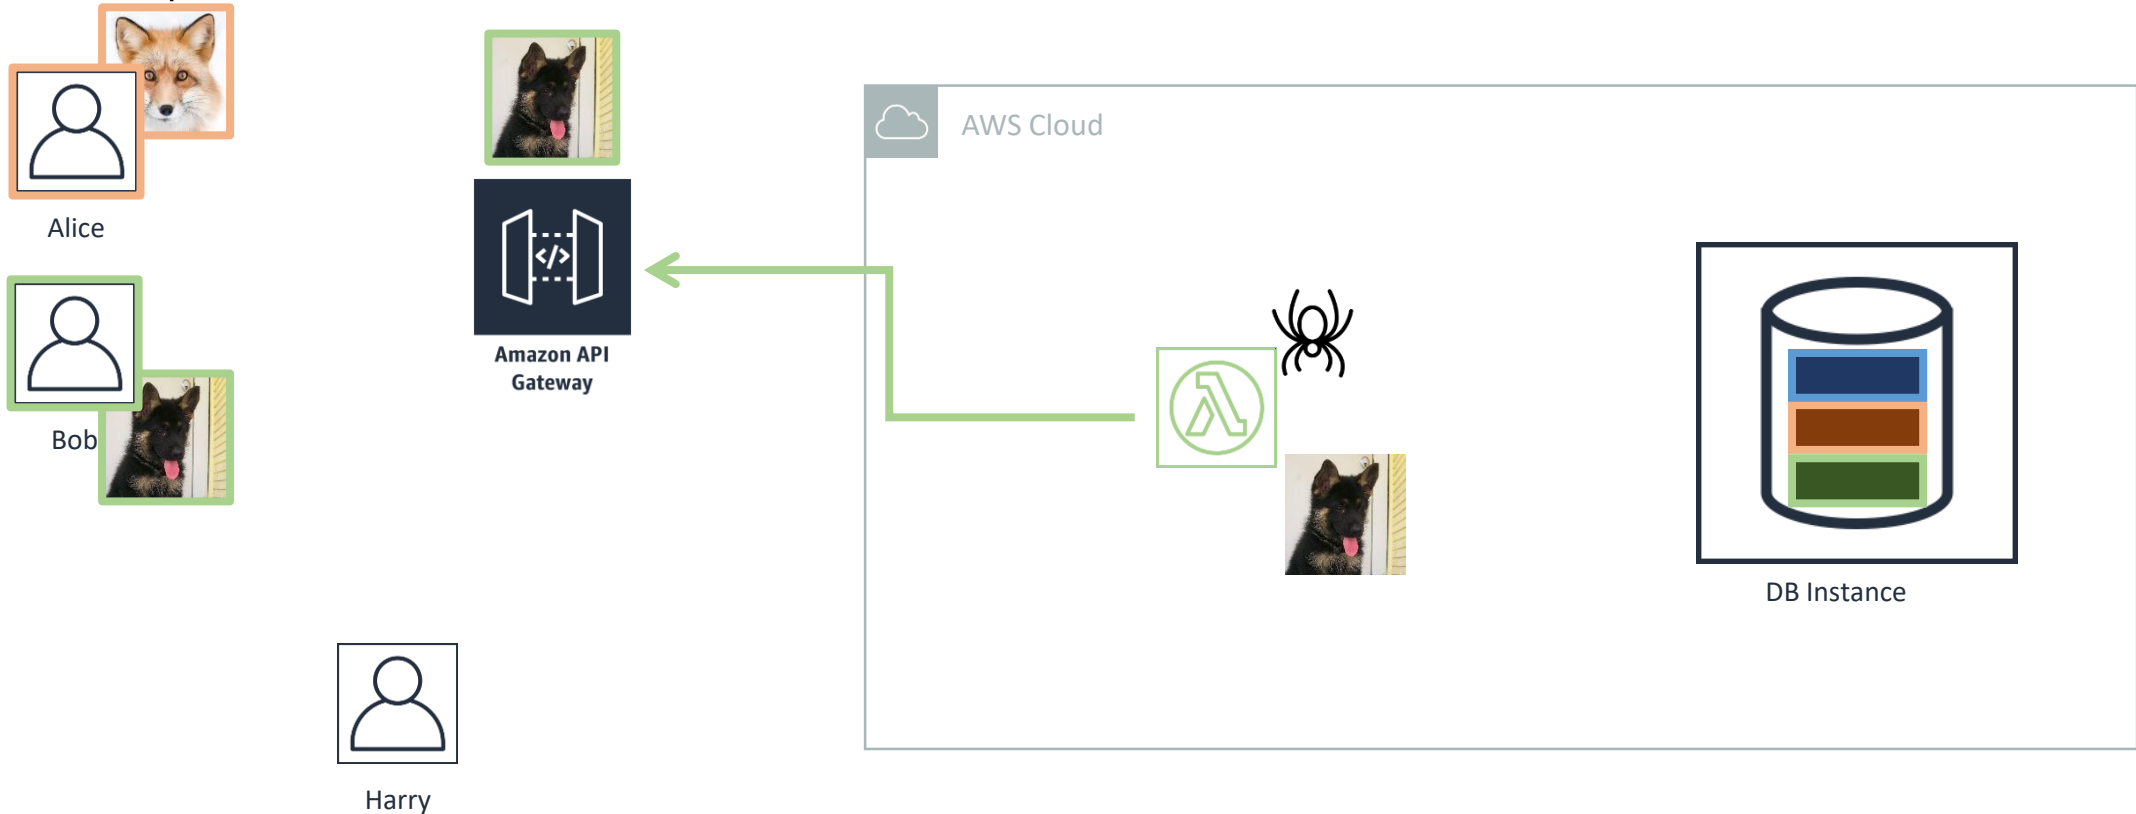


Amazon API Gateway



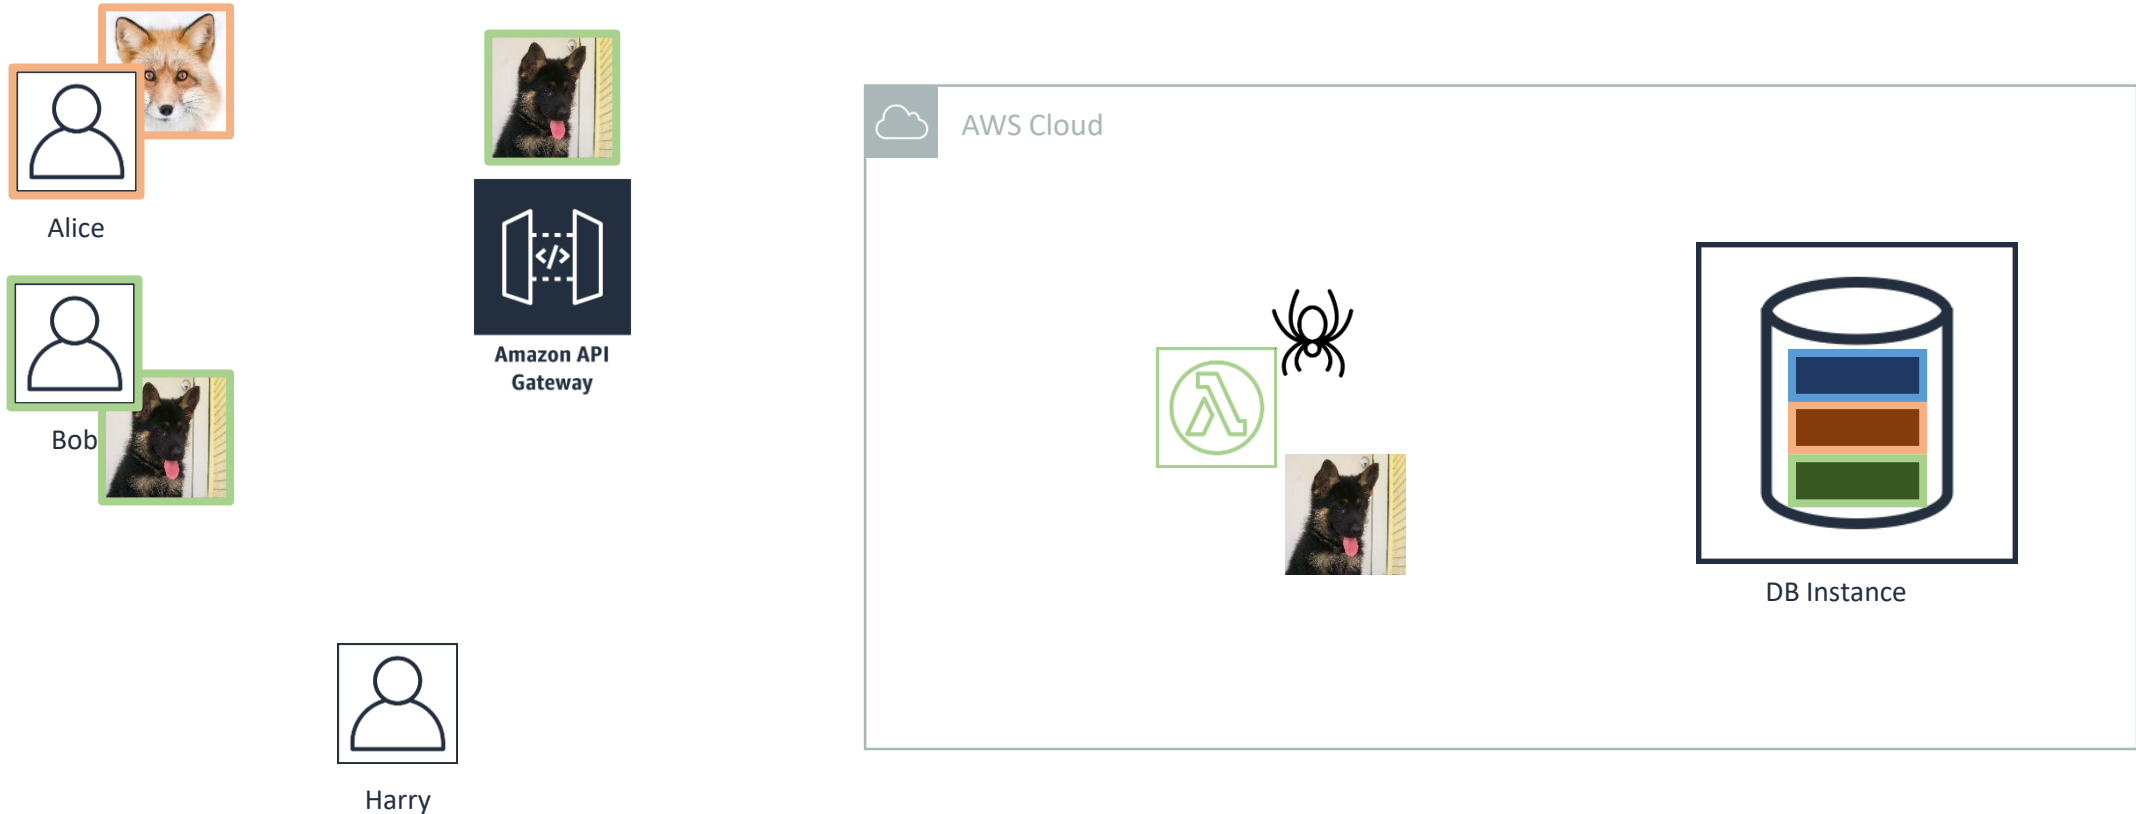
# Serverless+IFC

The malicious function then tries to send the secret in a response



# Serverless+IFC

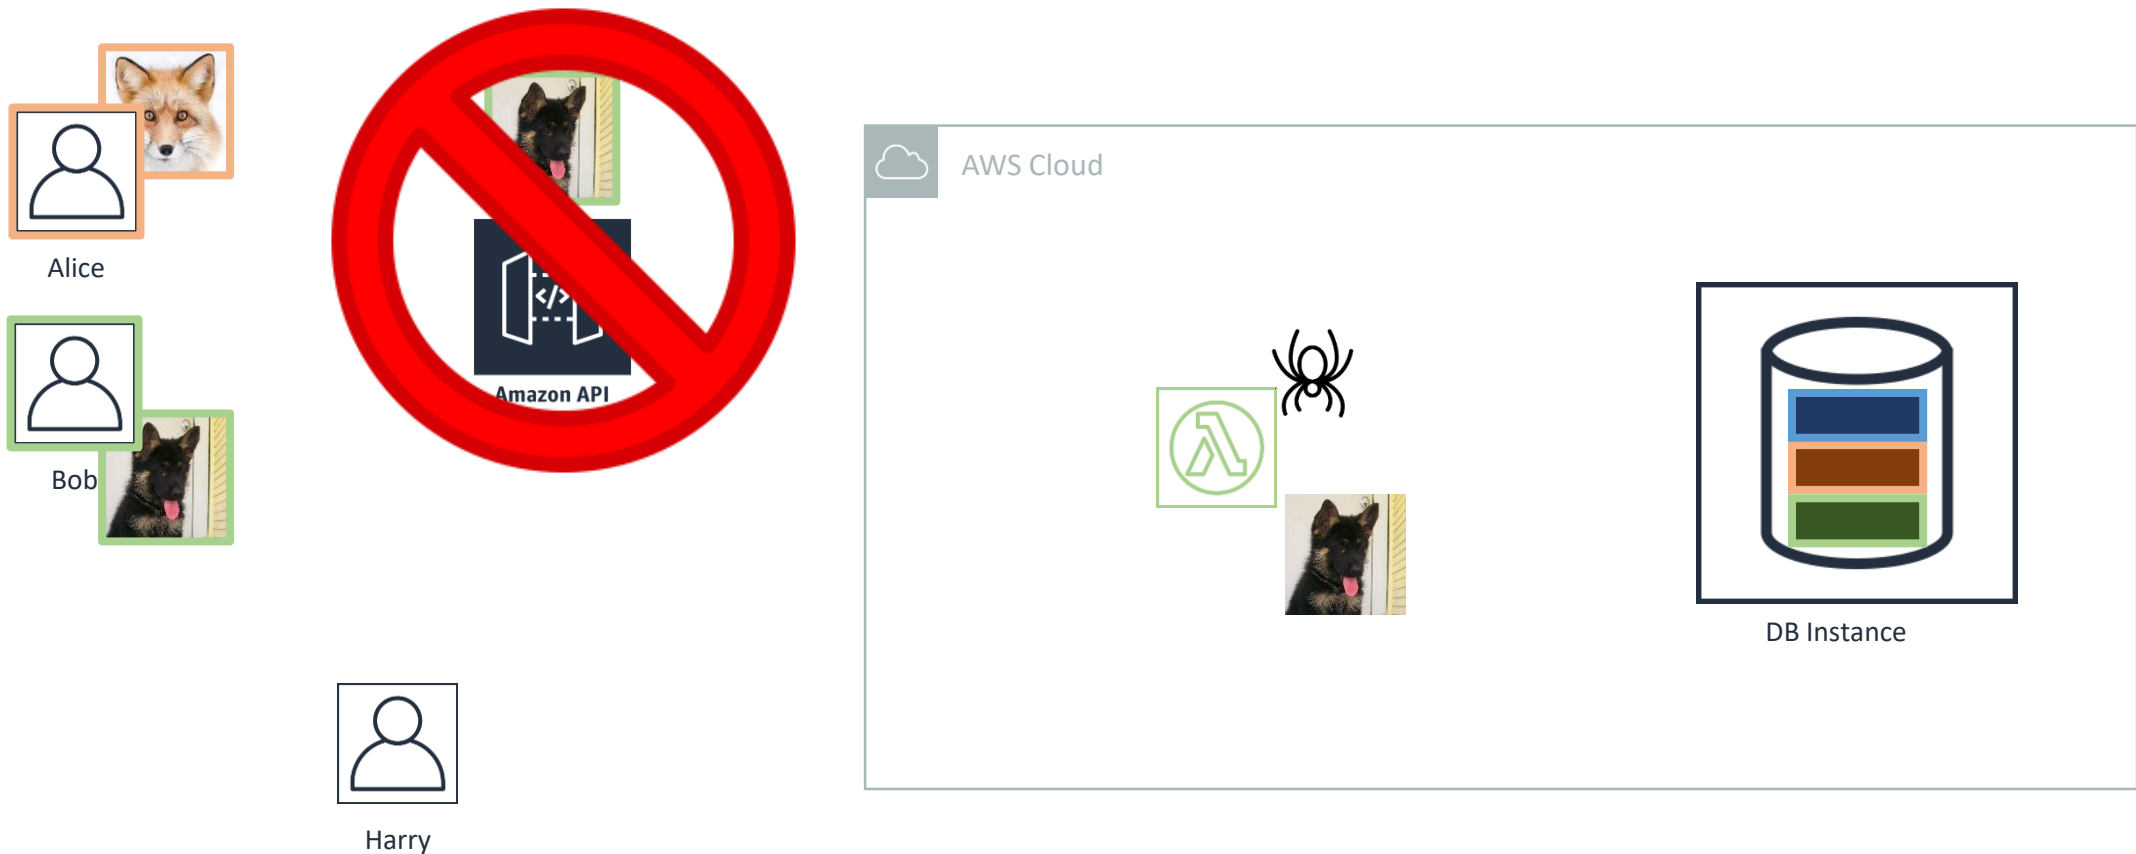
At which point we can check if the hacker is allowed to see the data





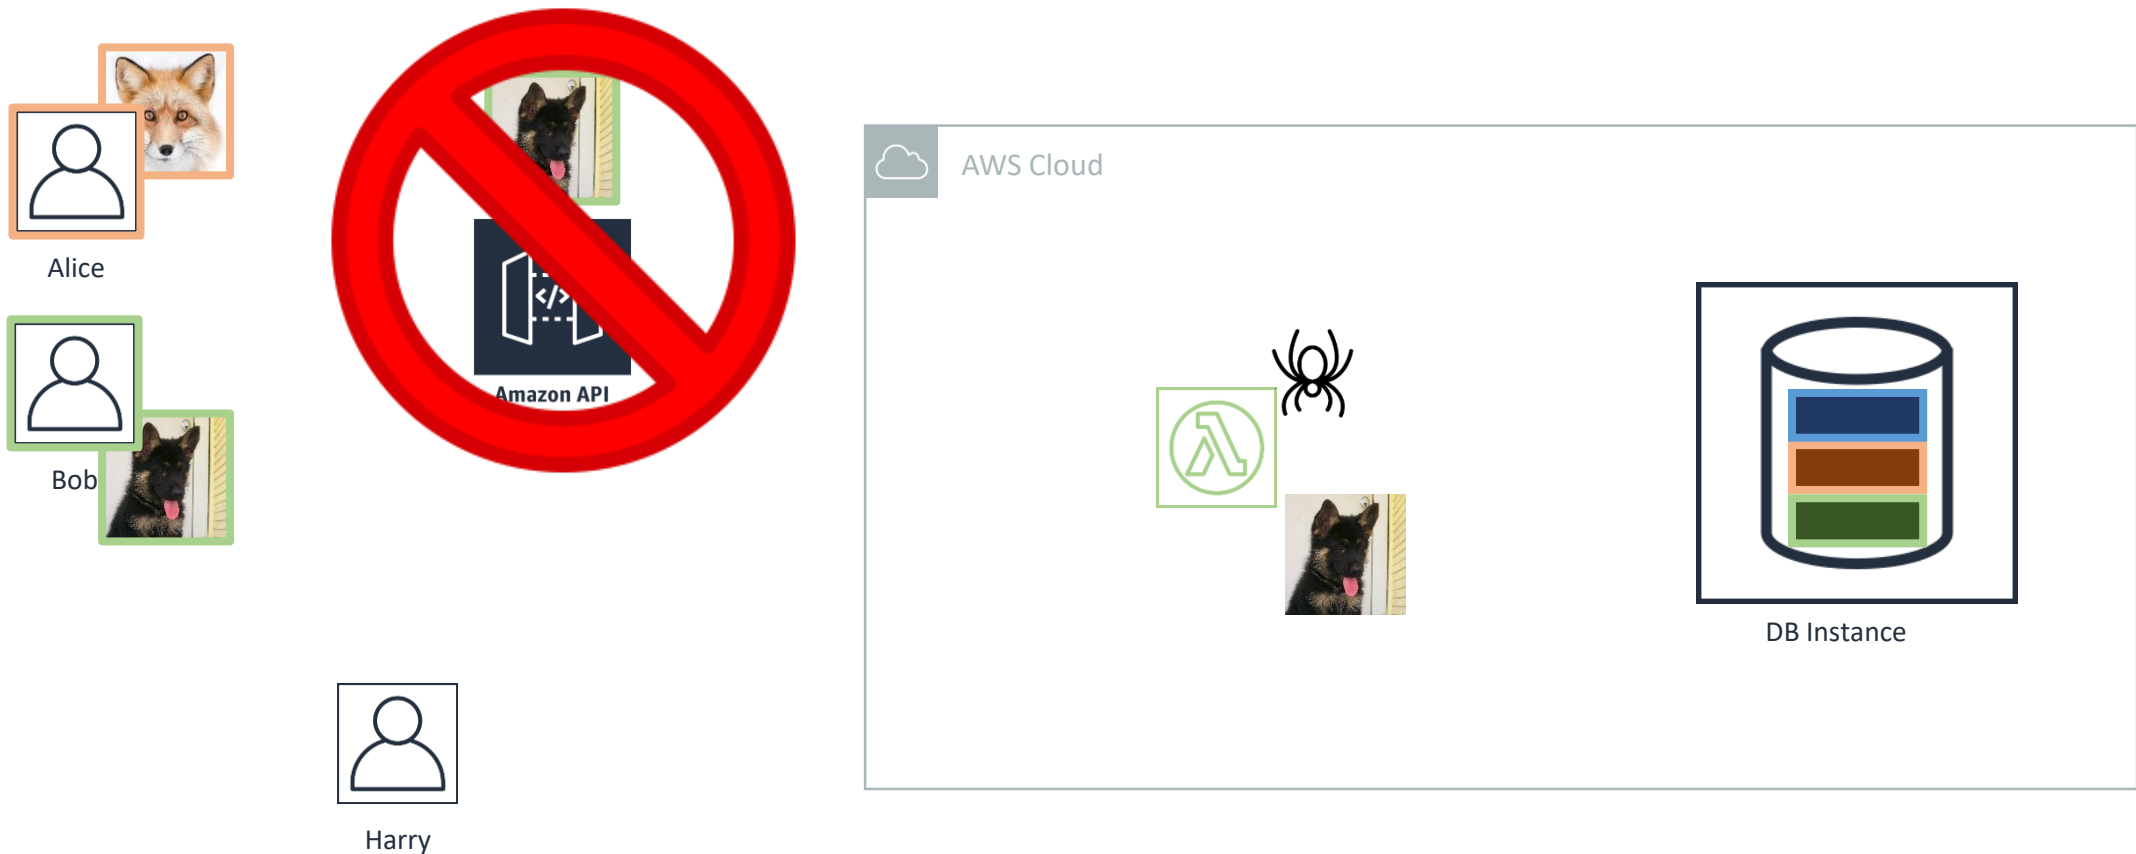
# Serverless+IFC

And block the request from being sent



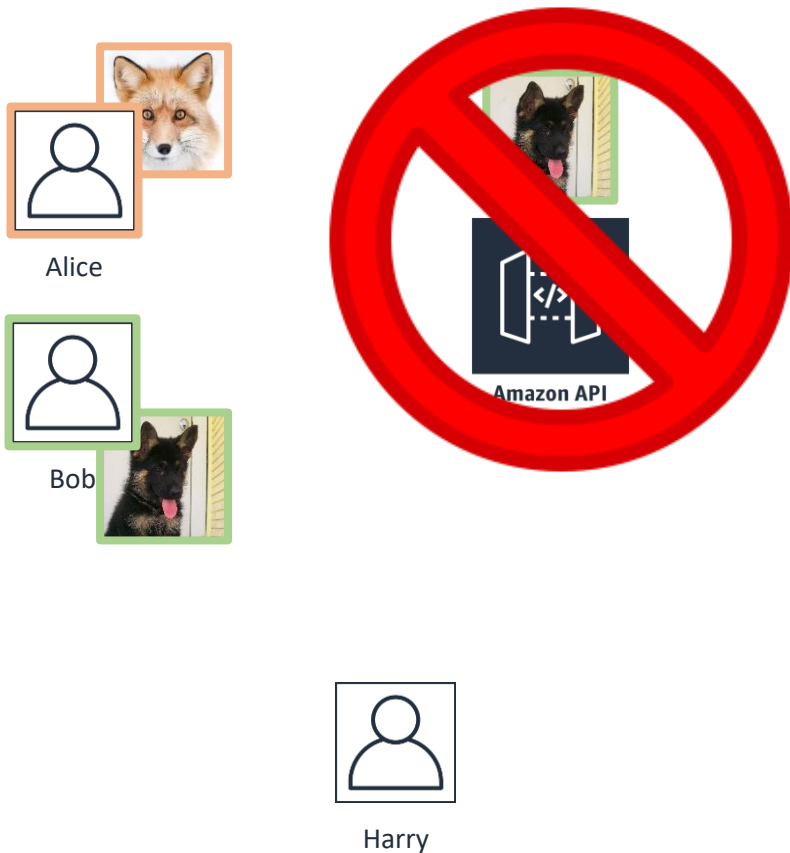
# Termination Channel

But – the attacker learned of the existence of secret data



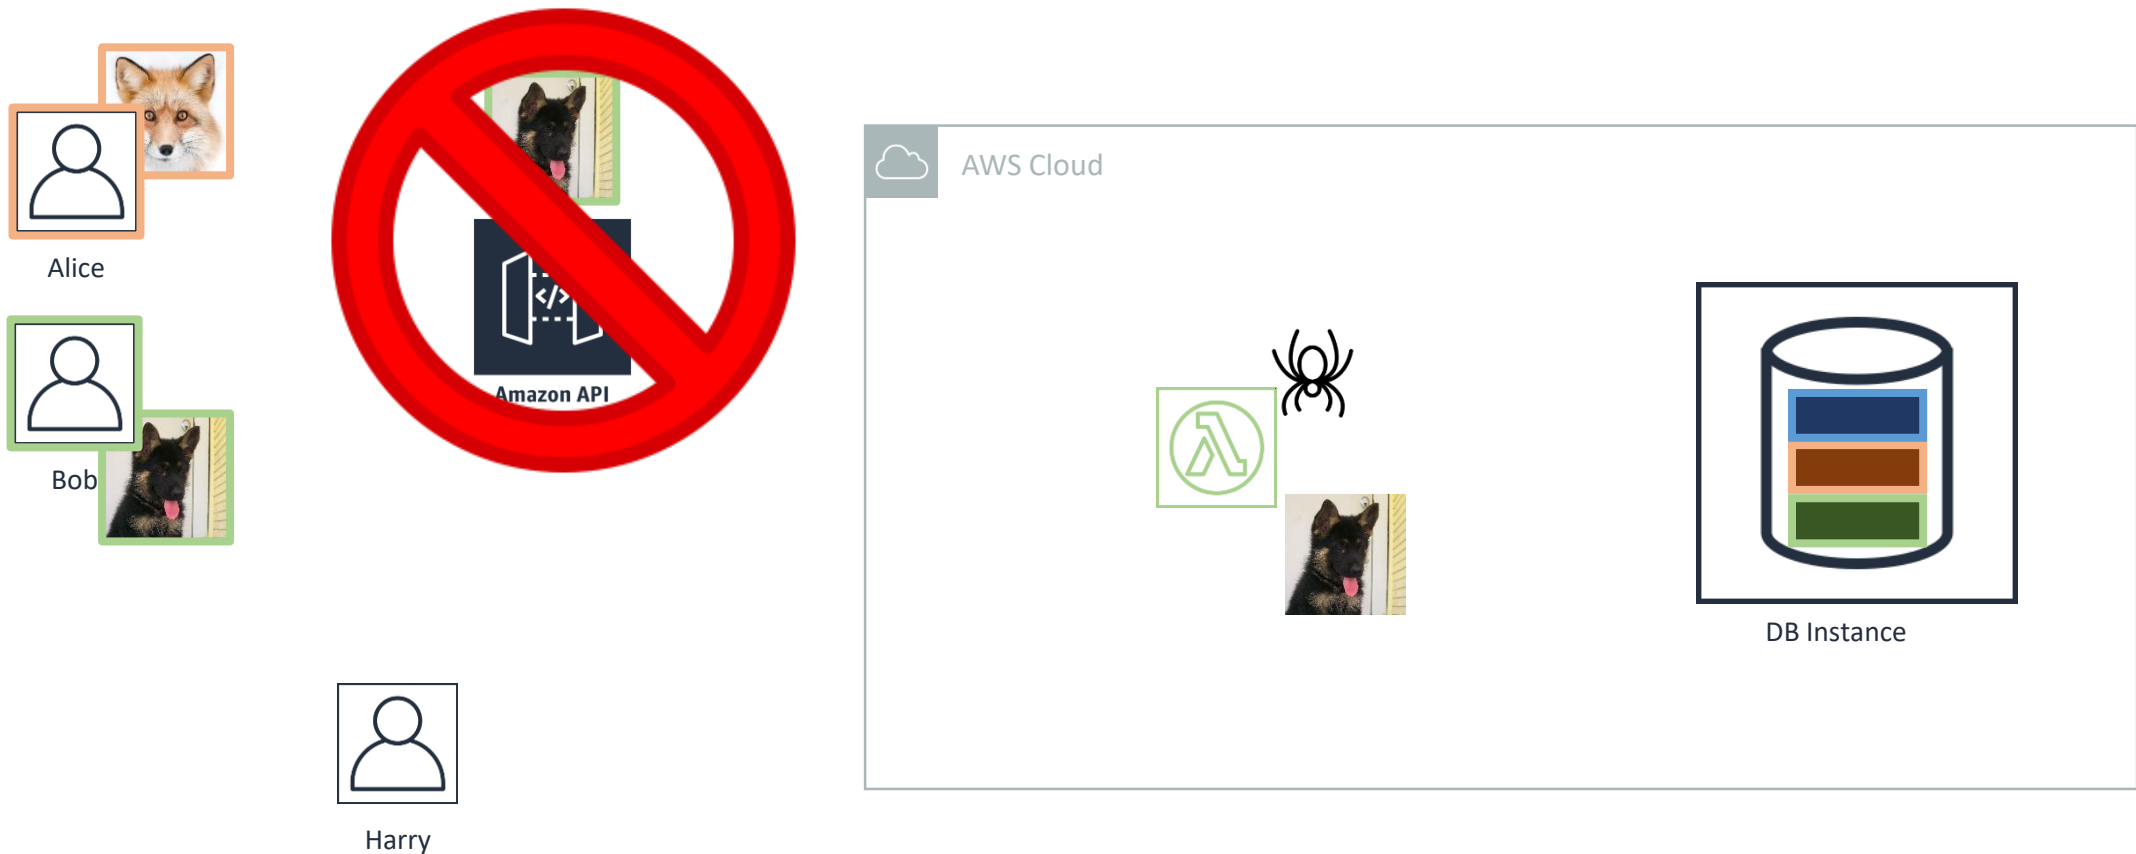
# Termination Channel

Which is bad if **the existence of a secret is itself a secret**



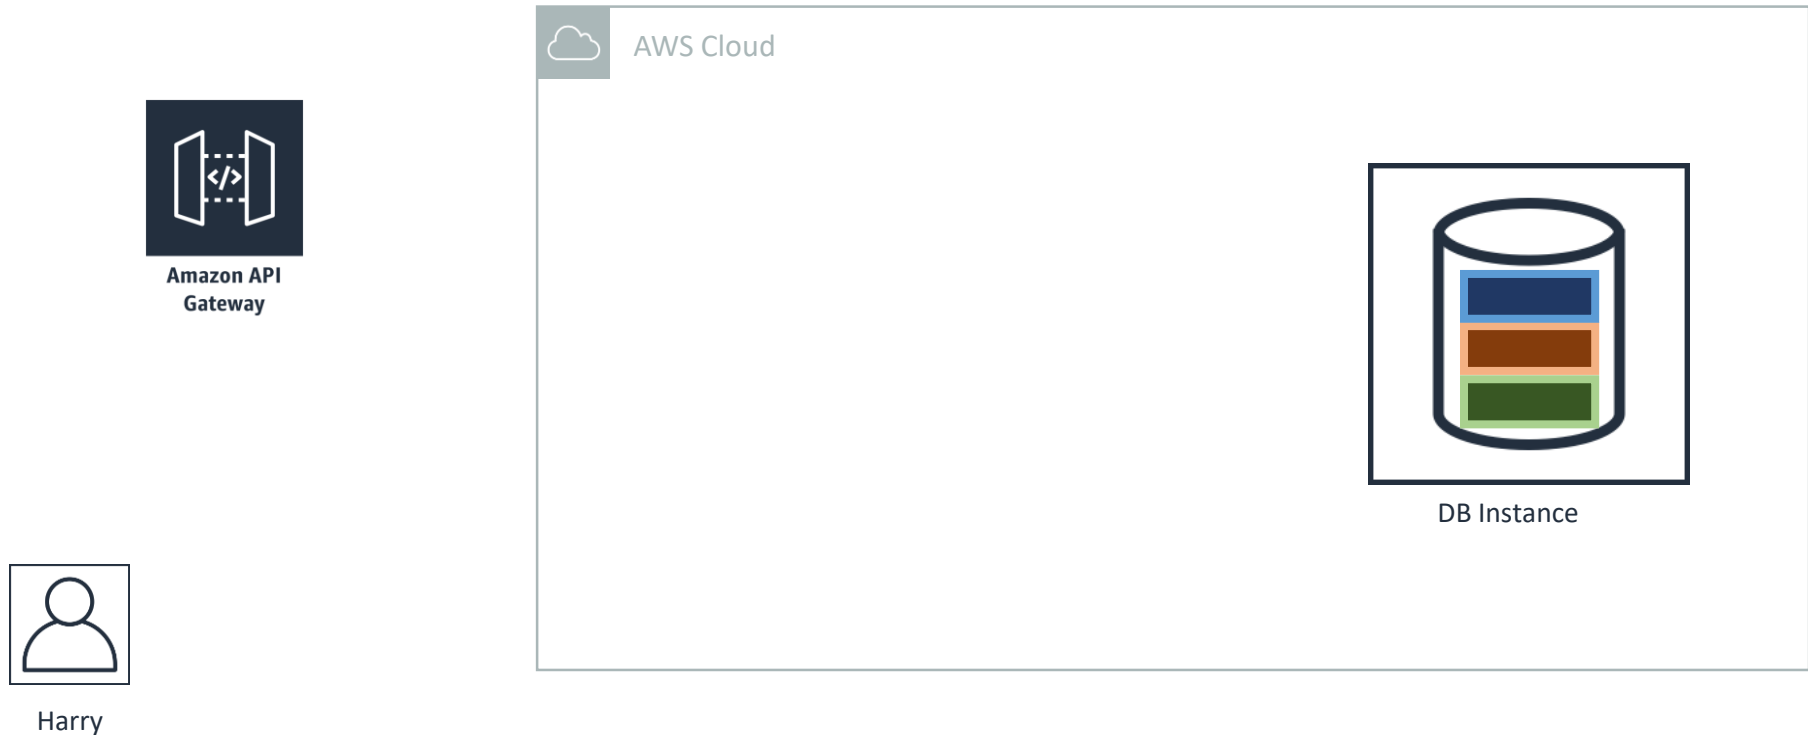
# Termination Channel

*Termination Insensitive Non-Interference* is insufficient!



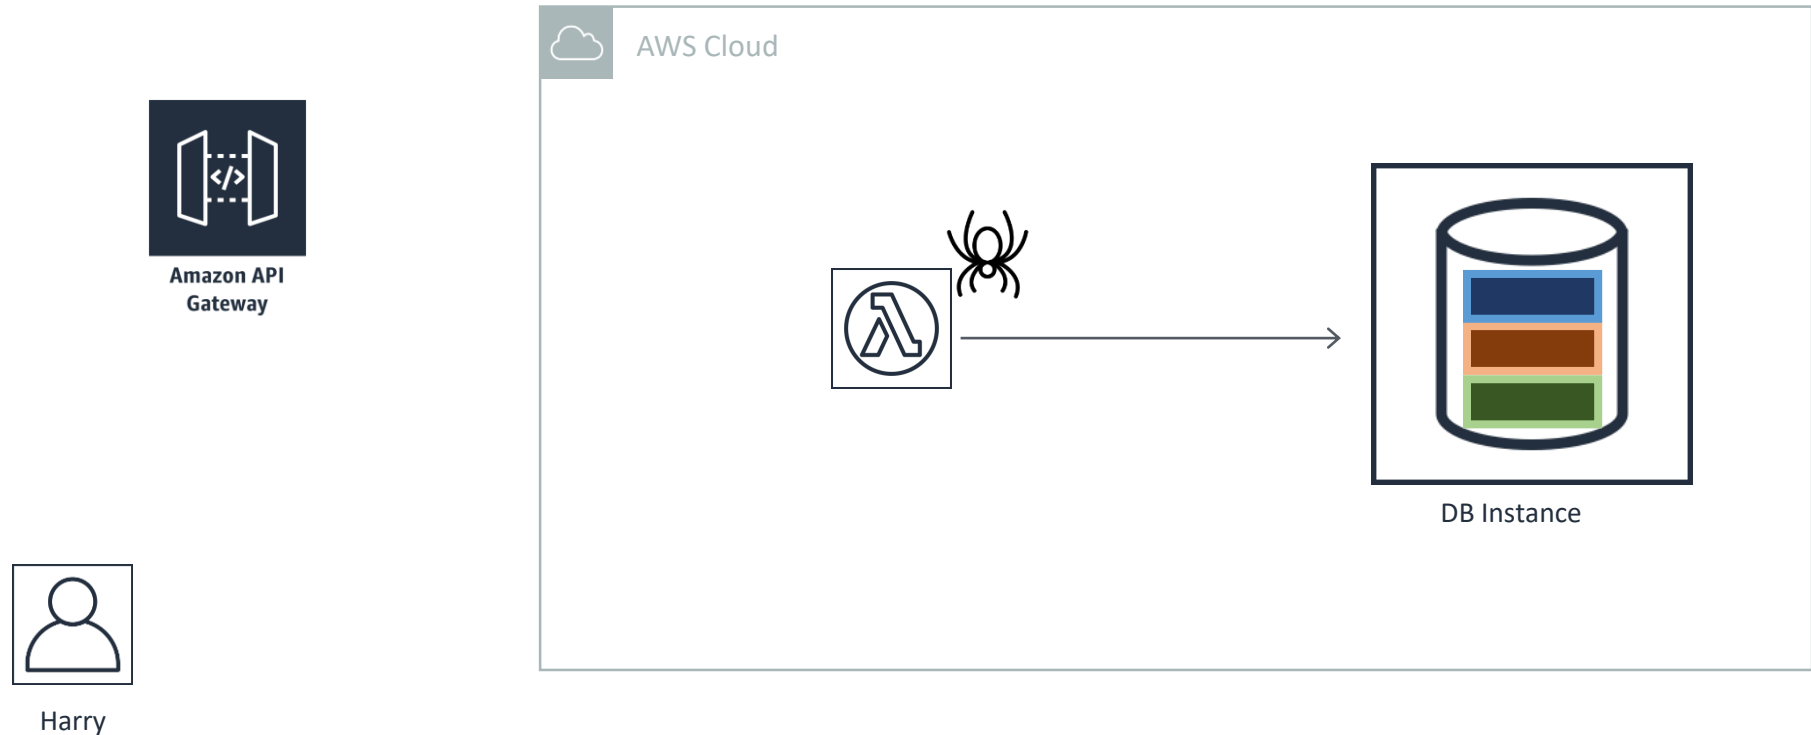
# Termination Channel

Attacker can encode a secret value by making certain values 'inaccessible'



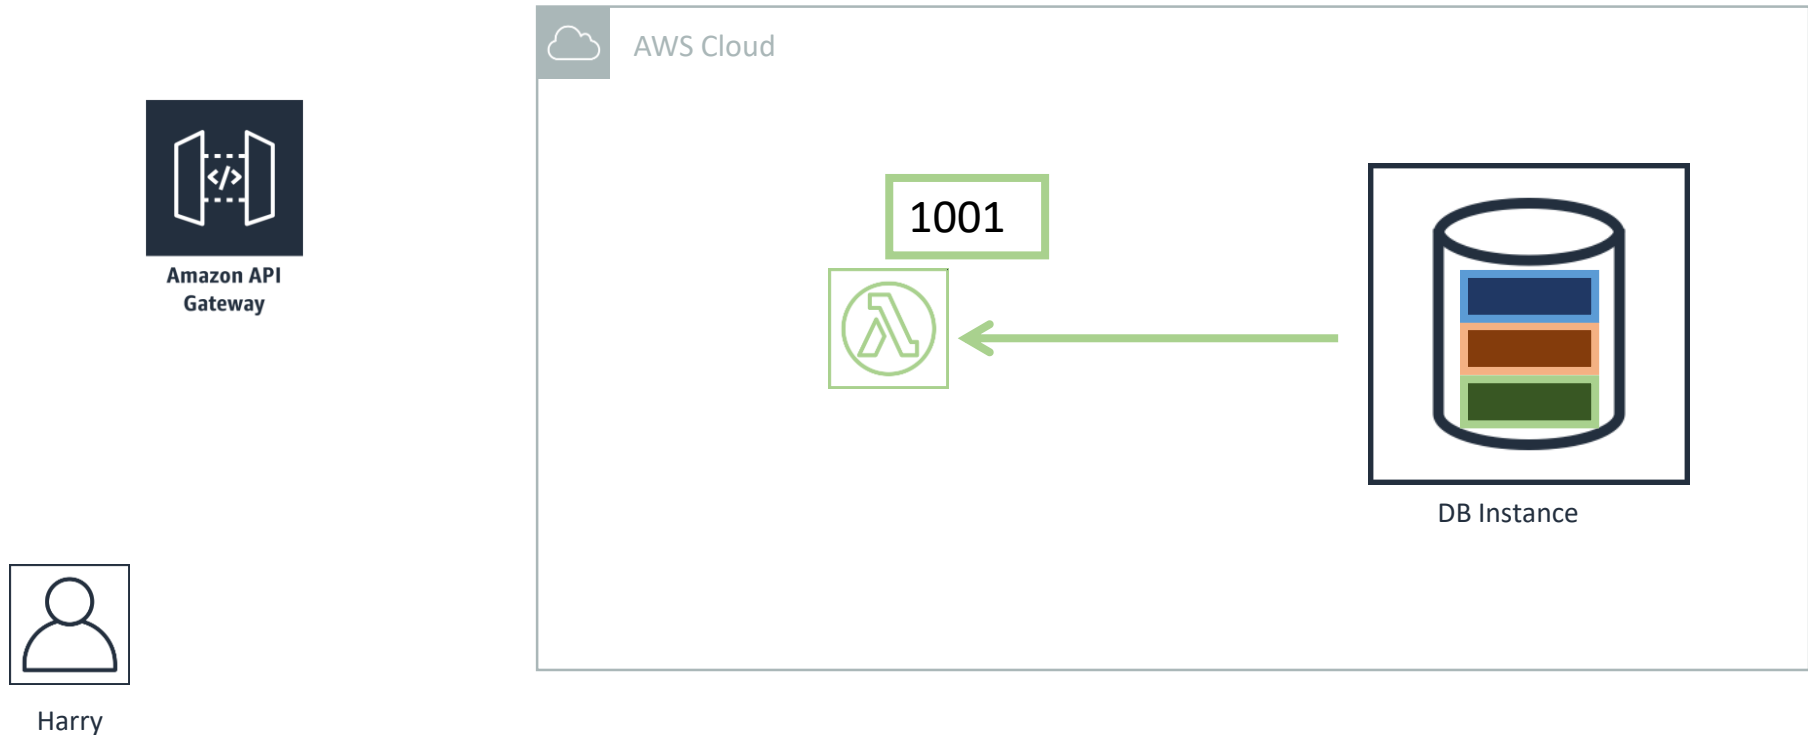
# Termination Channel

Attacker reads a secret value



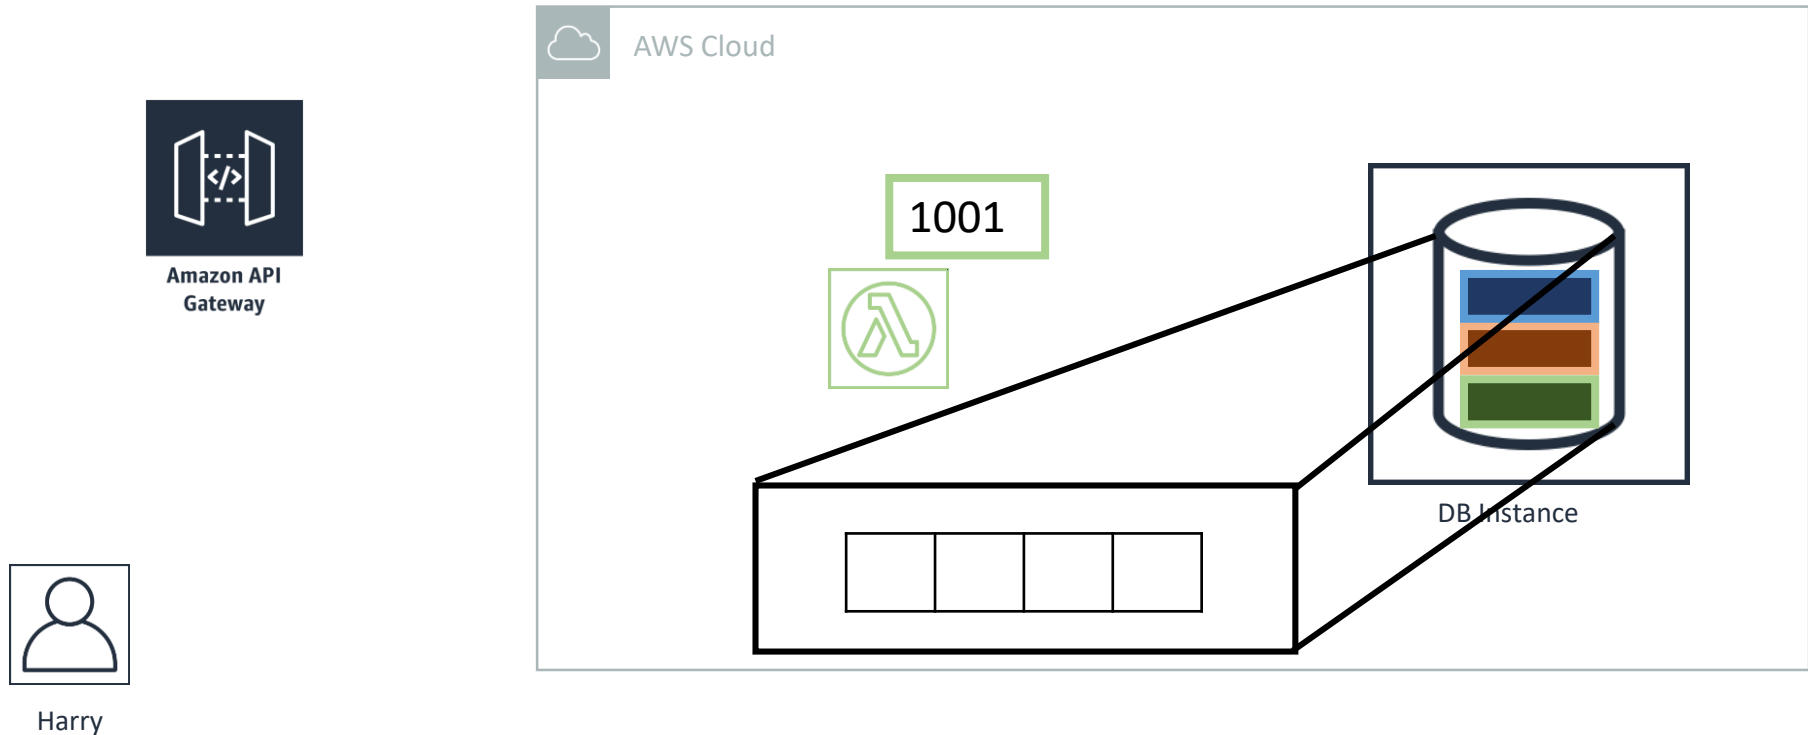
# Termination Channel

Attacker reads a secret value ('1001')



# Termination Channel

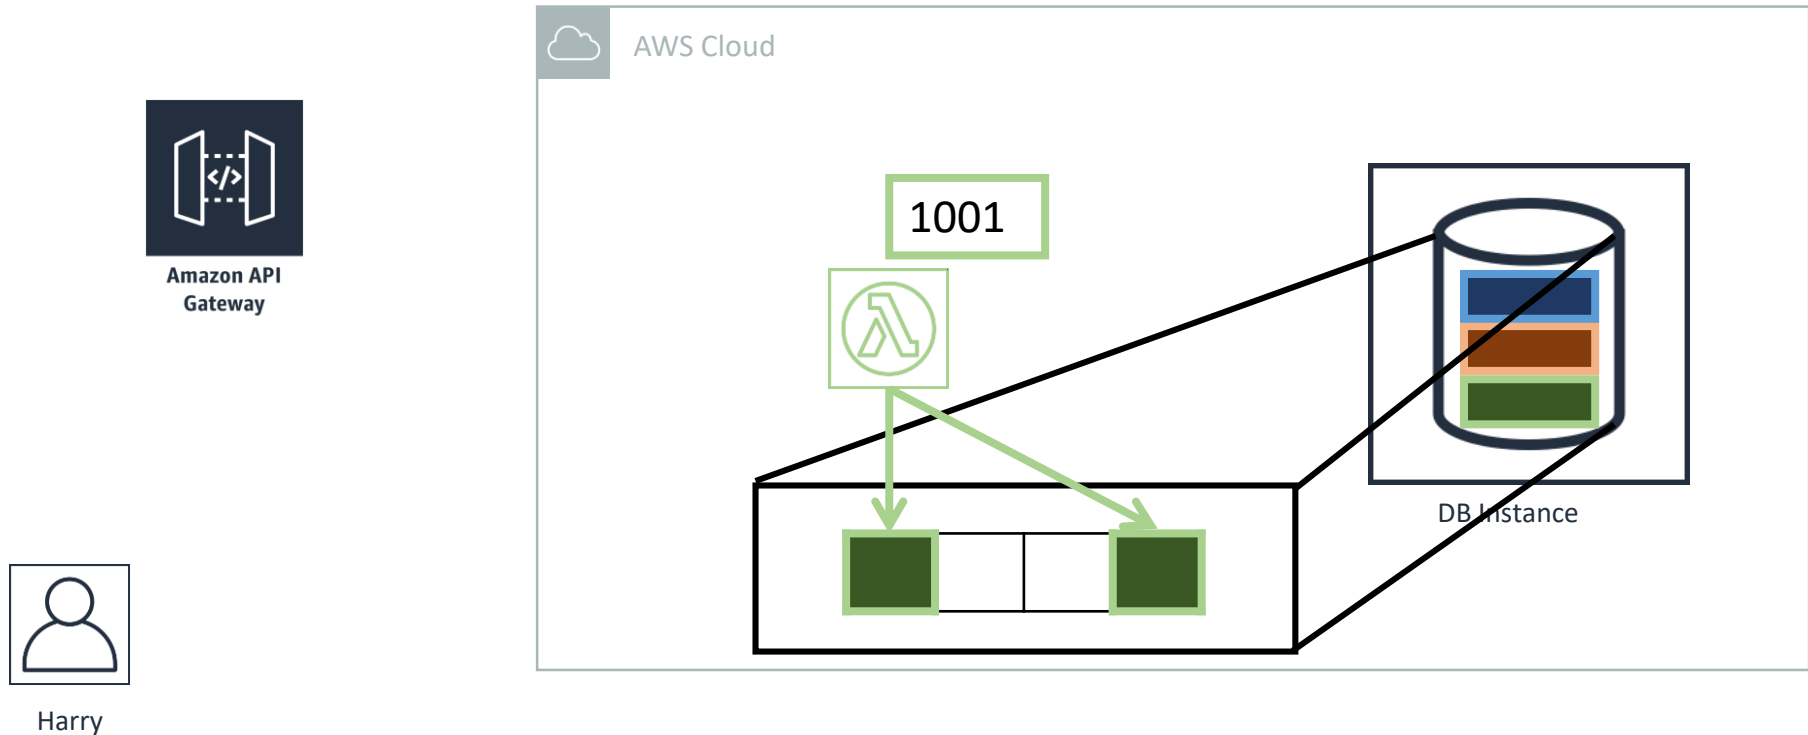
Write something to predefined locations





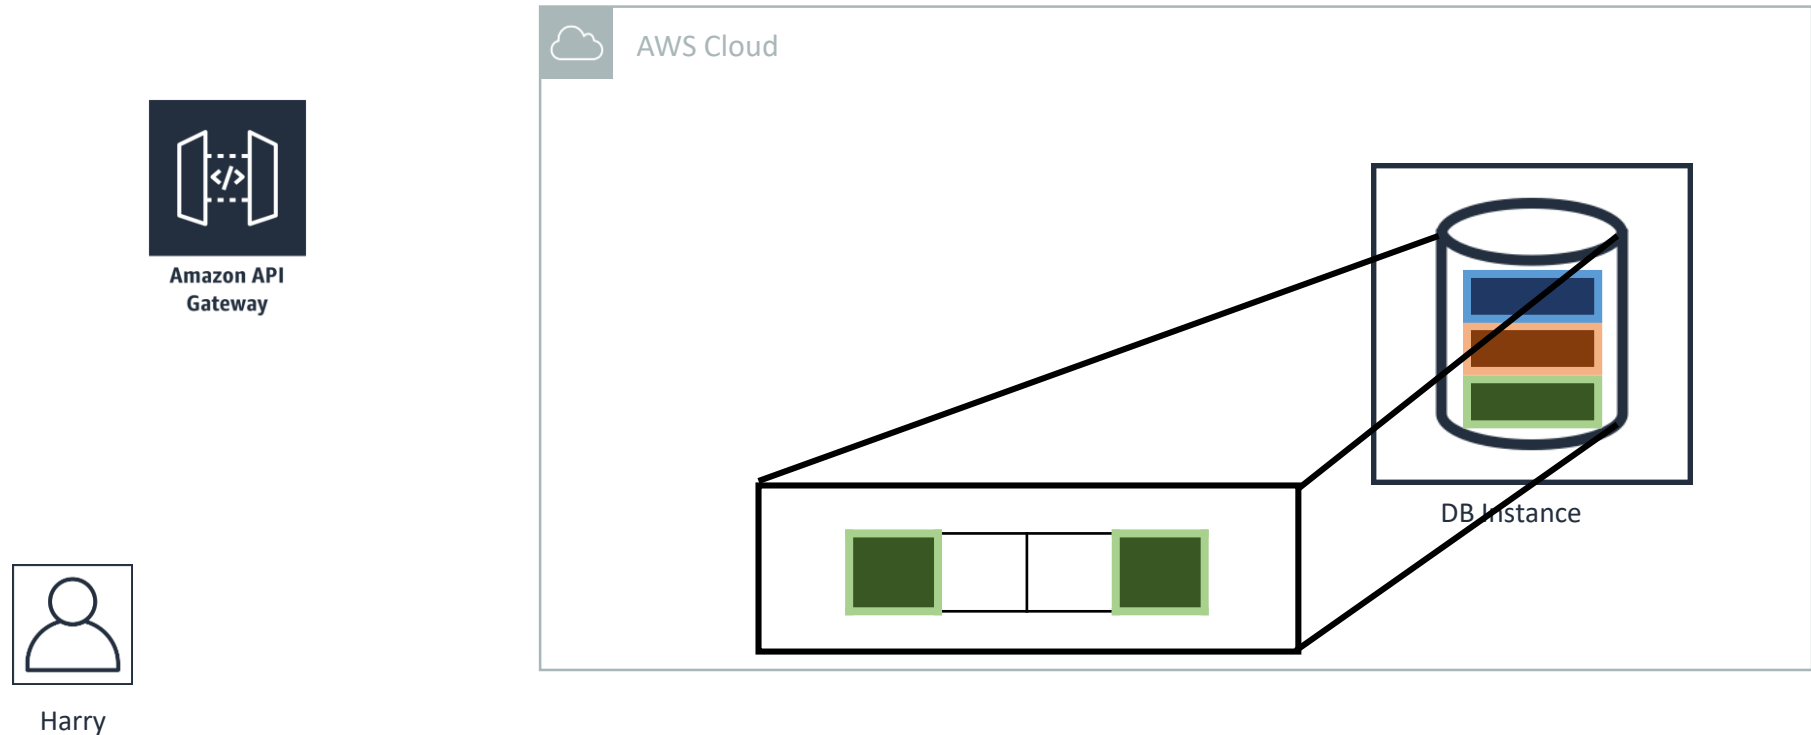
# Termination Channel

Encode secret as a memory storage pattern



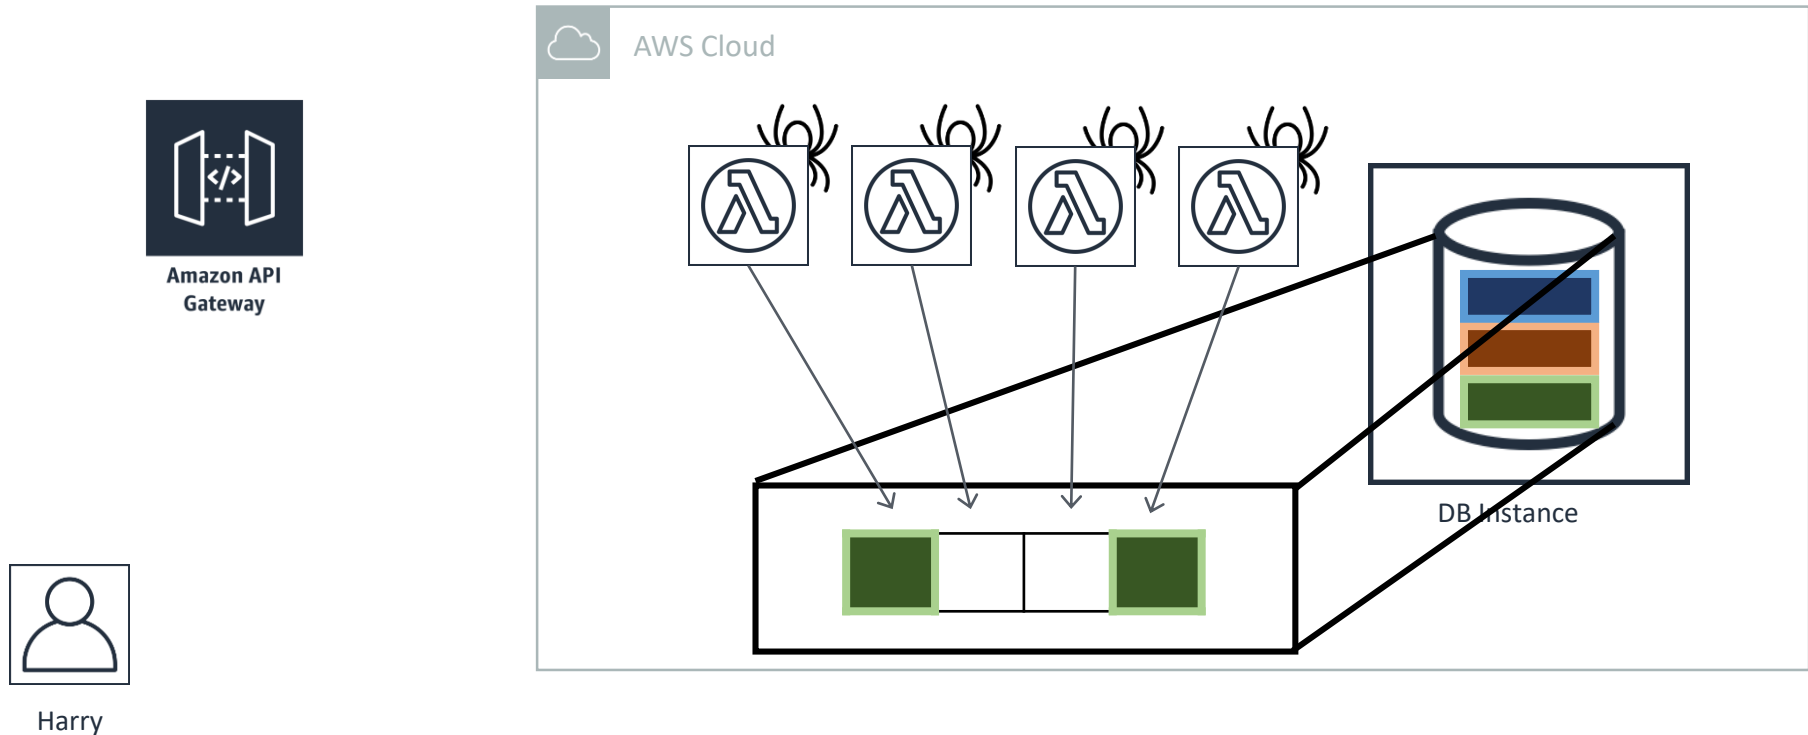
# Termination Channel

Attacker can leak every bit by accessing each cell separately



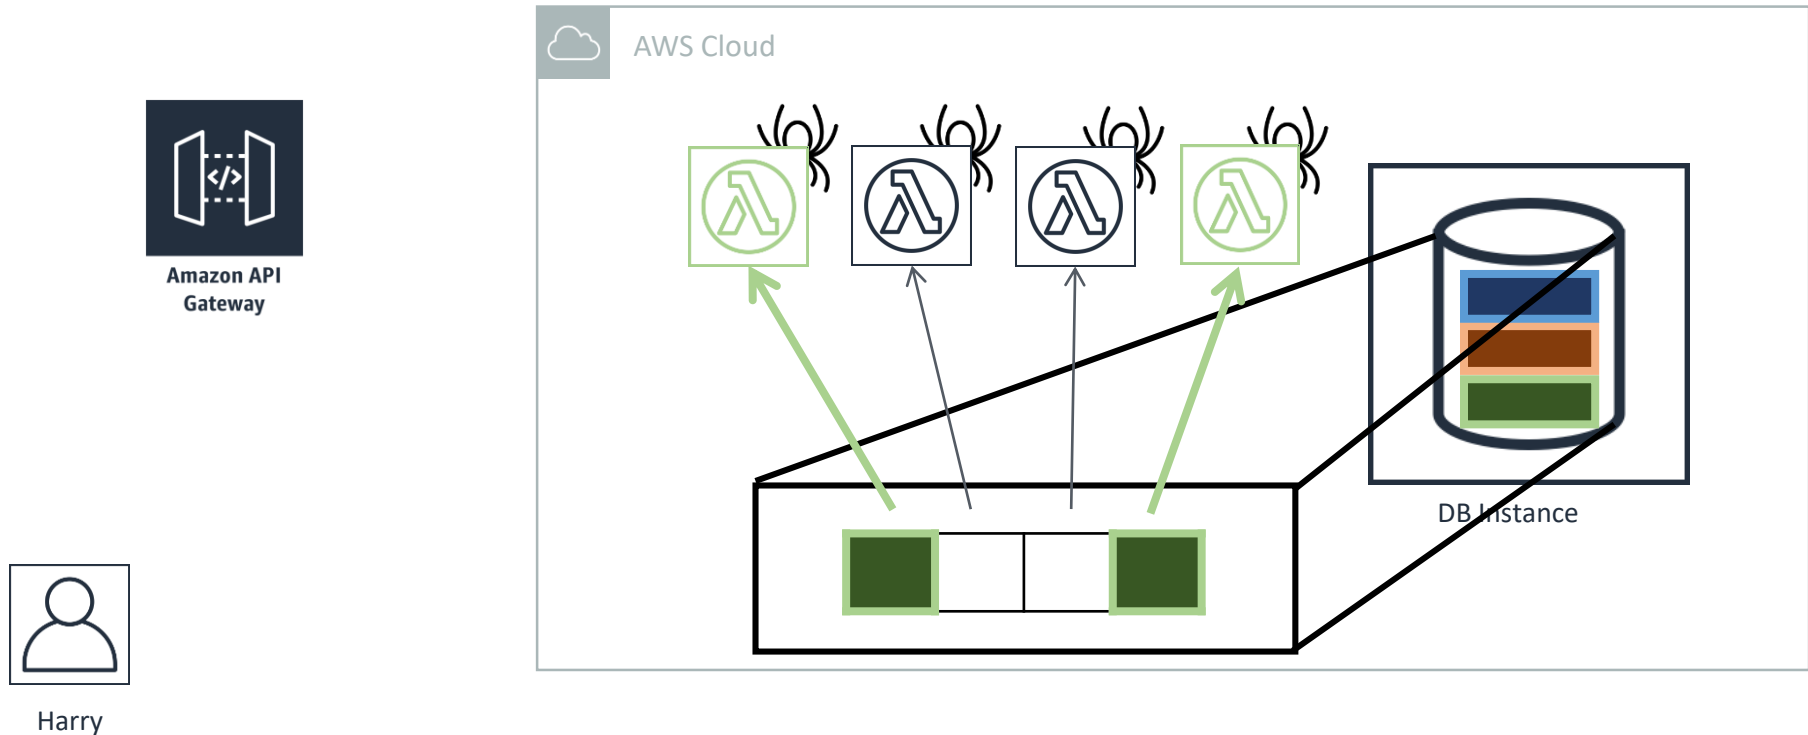
# Termination Channel

Attacker can leak every bit by accessing each cell separately



# Termination Channel

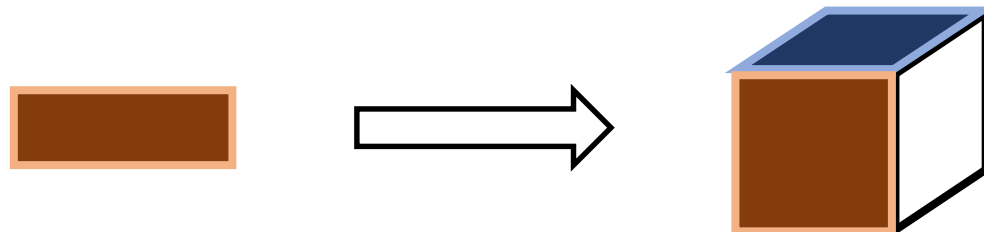
Attacker can leak every bit by accessing each cell separately



# Closing the Termination Channel

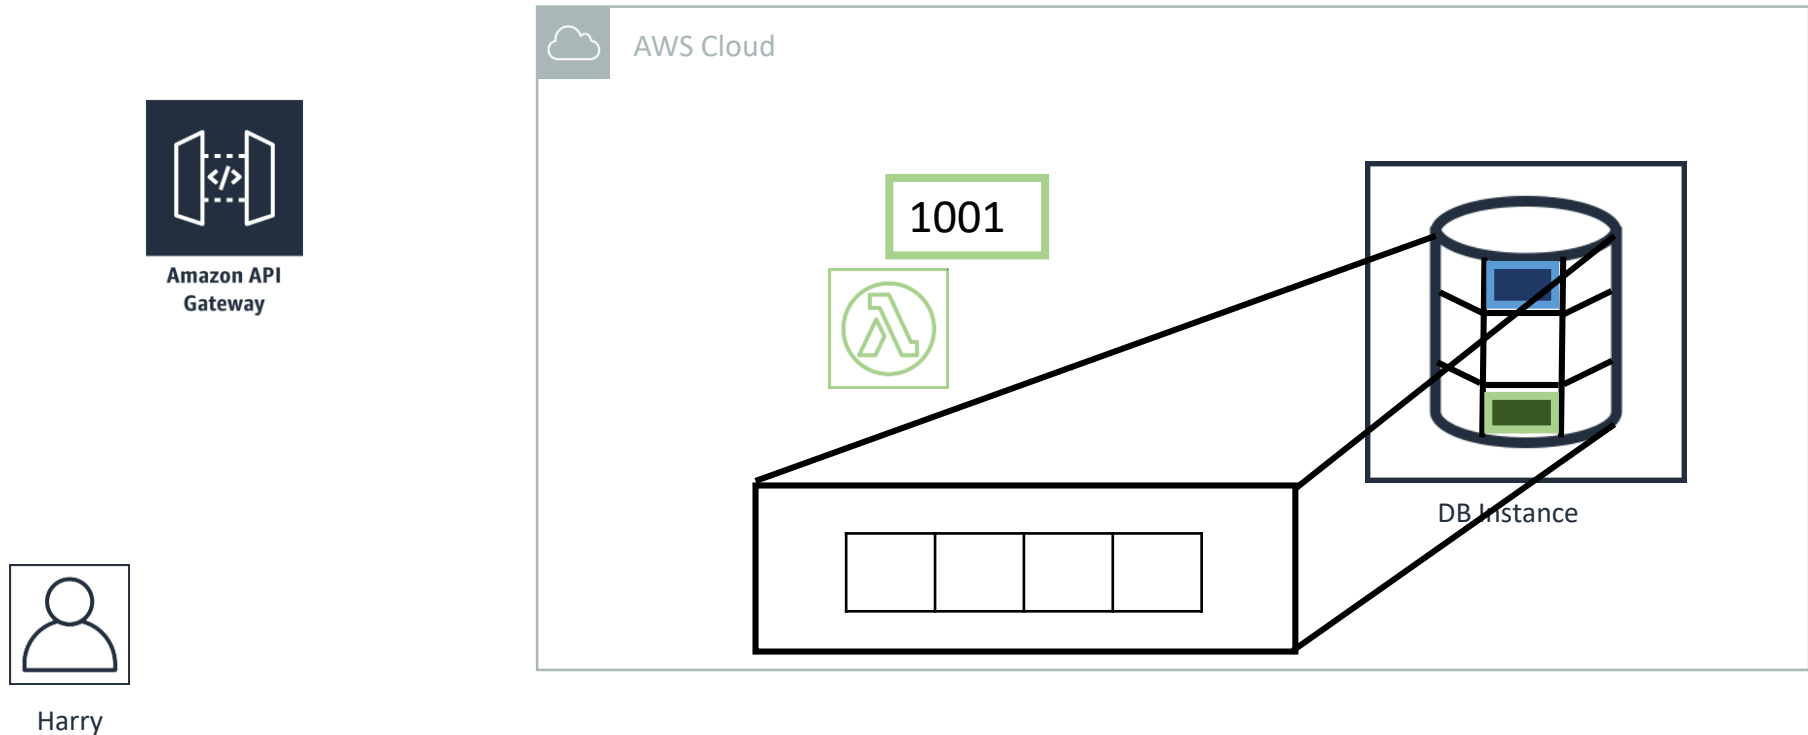
How do we get *Termination **Sensitive** Non-Interference*?

- Static labeling instead of floating labels
  - Label assigned at function invocation
  - Output behavior depends on invocation, not secret
- Faceted data store



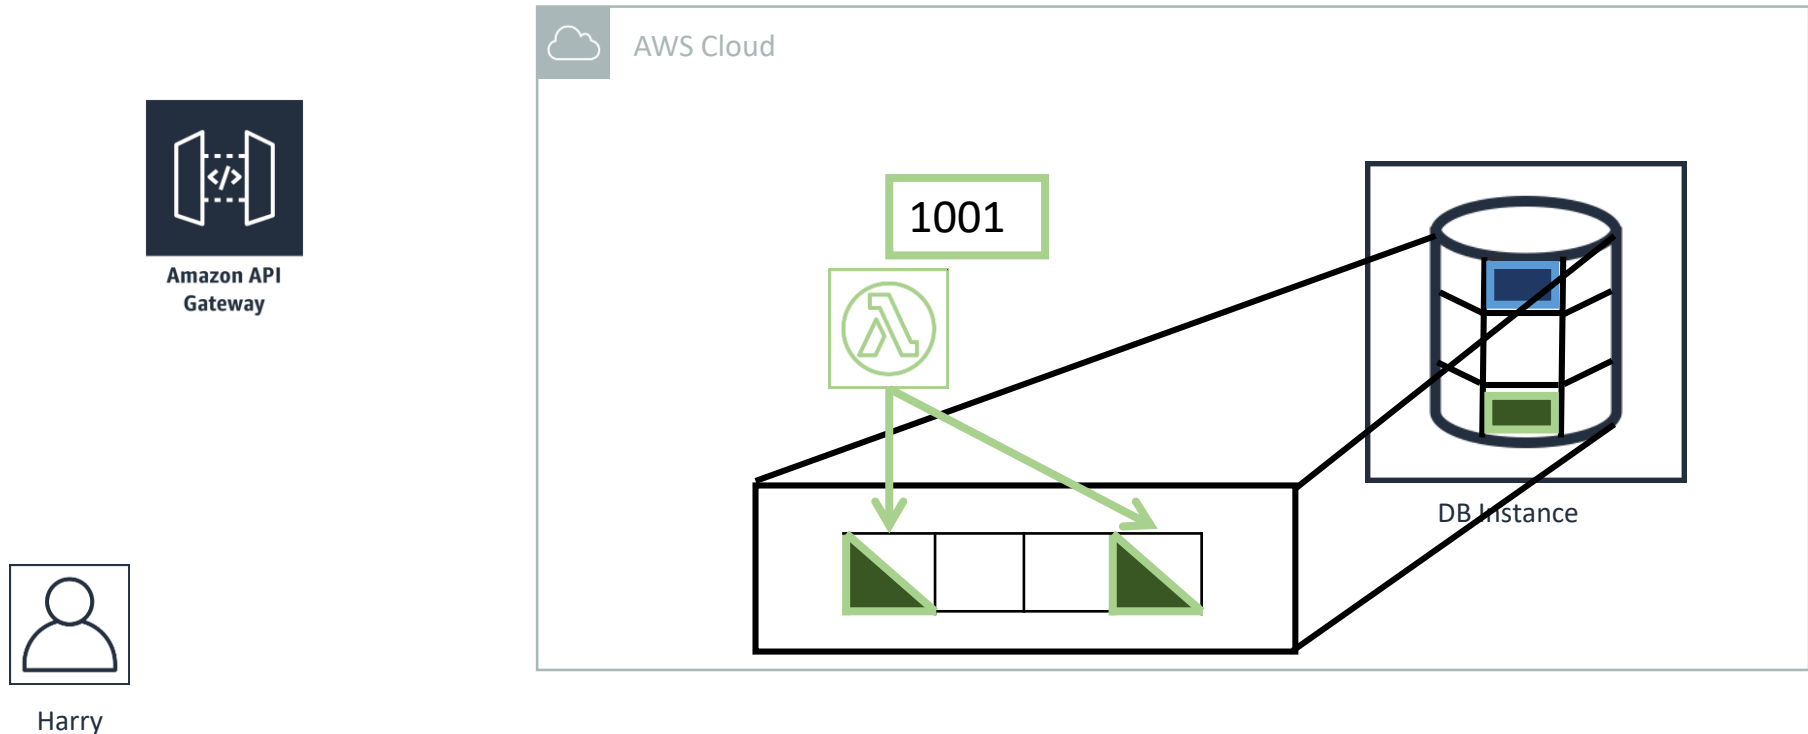
# Closing the Termination Channel

## Modified read and write semantics



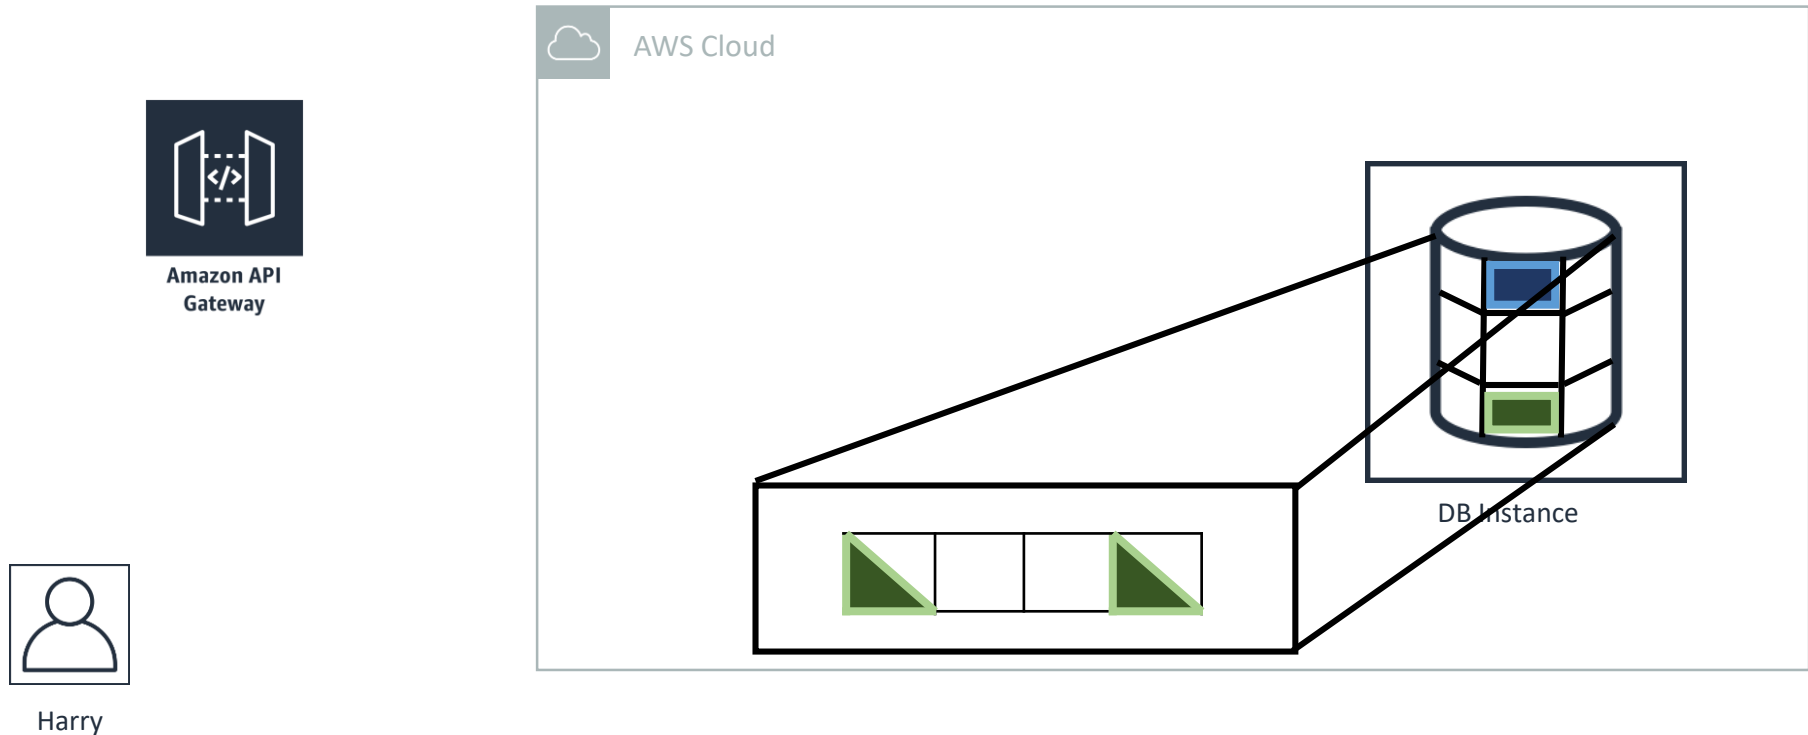
# Closing the Termination Channel

Writes add a facet to the location, instead of overwriting



# Closing the Termination Channel

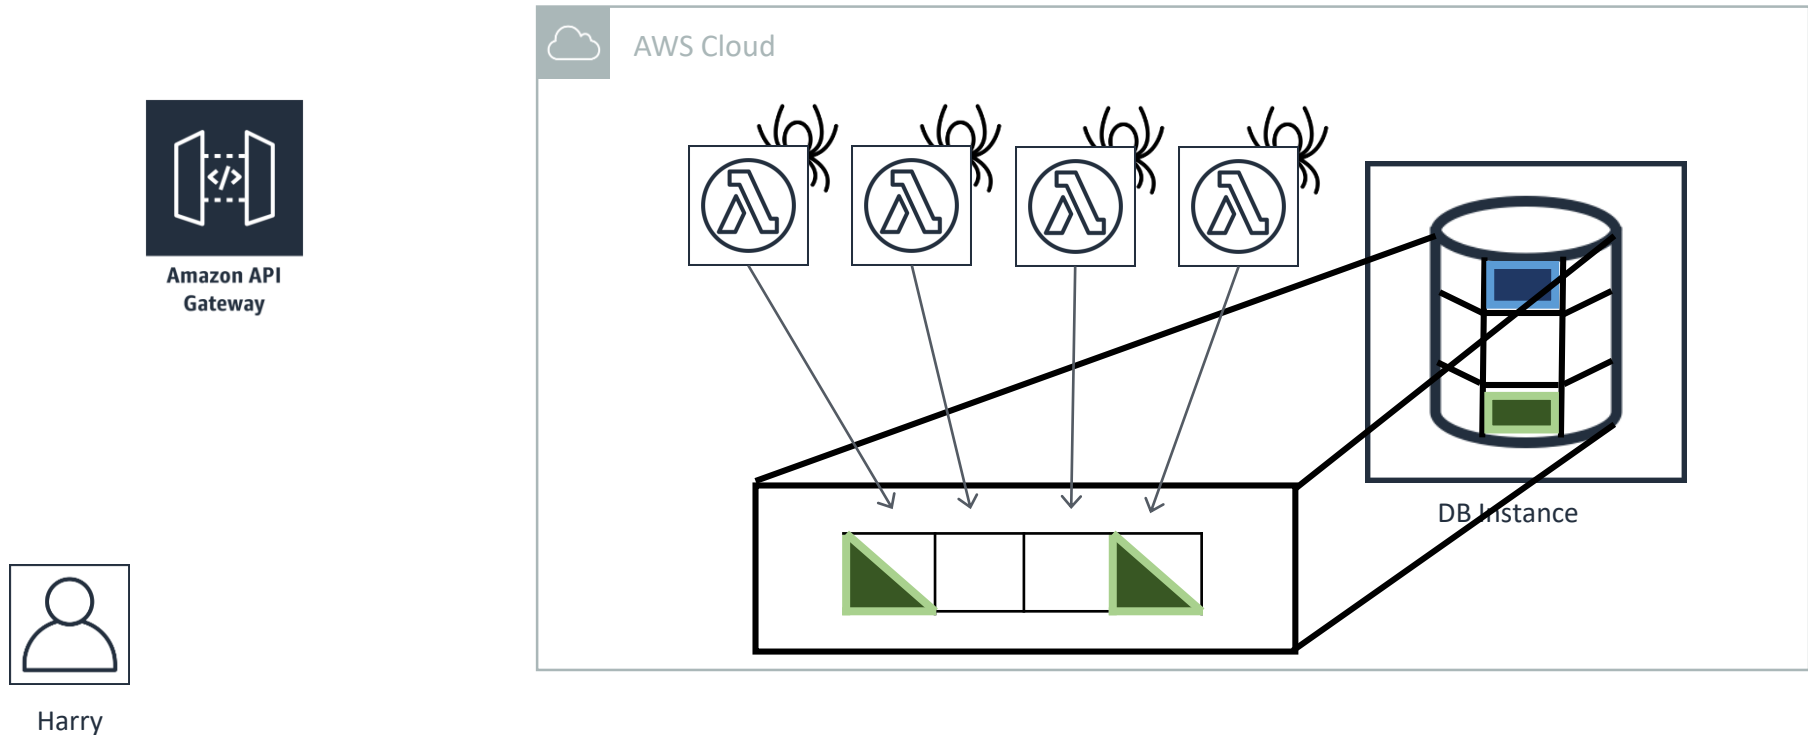
An execution only sees the facets that is allowed to see





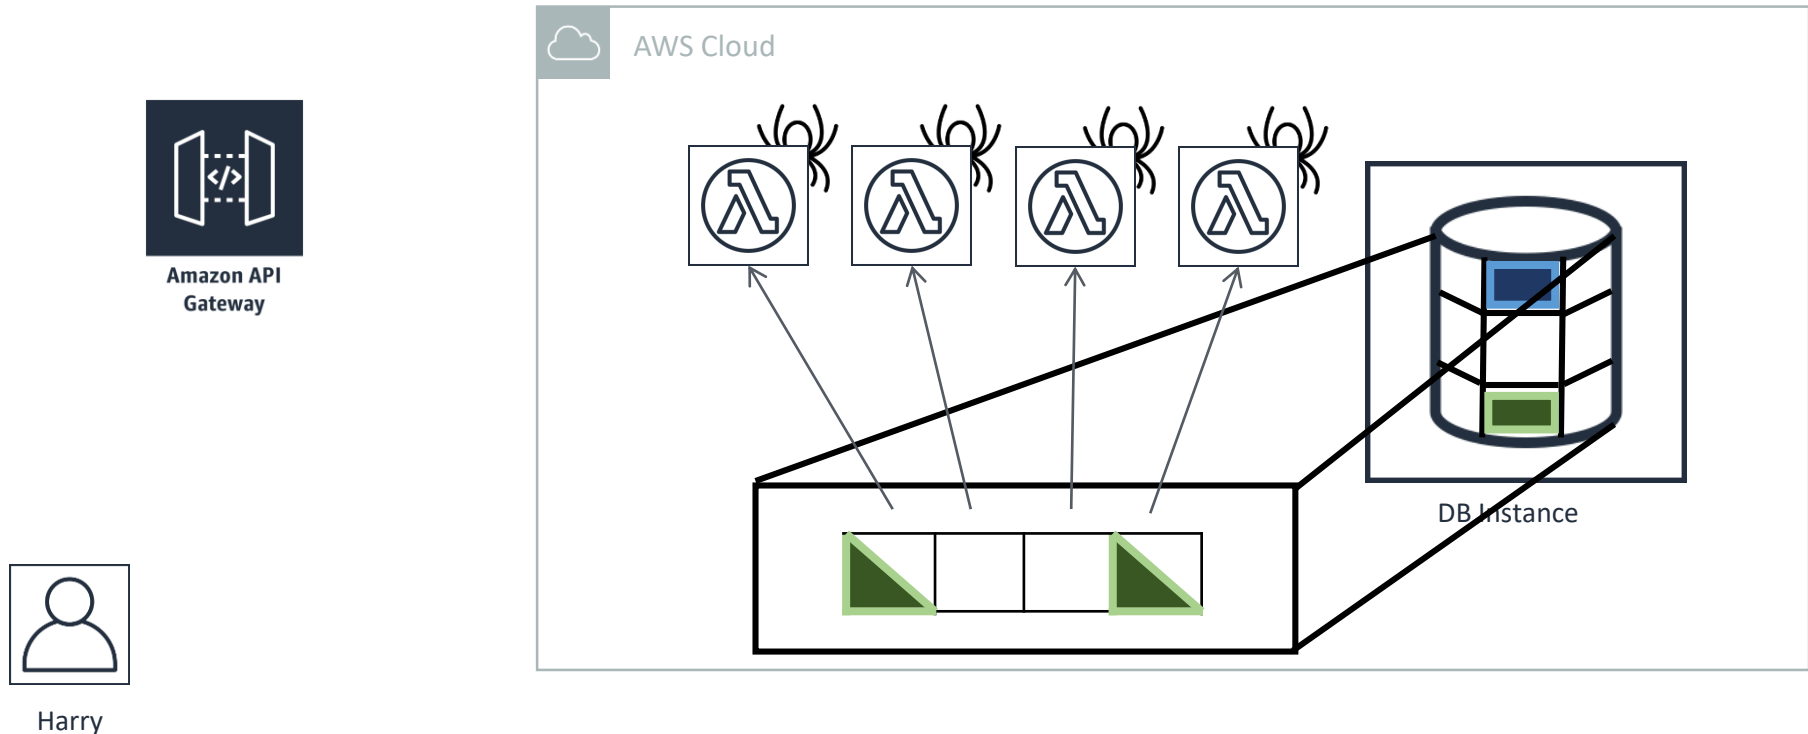
# Closing the Termination Channel

An execution only sees the facets that is allowed to see



# Closing the Termination Channel

An execution without the correct label can't distinguish between an empty cell, and a faceted cell



# Trapeze

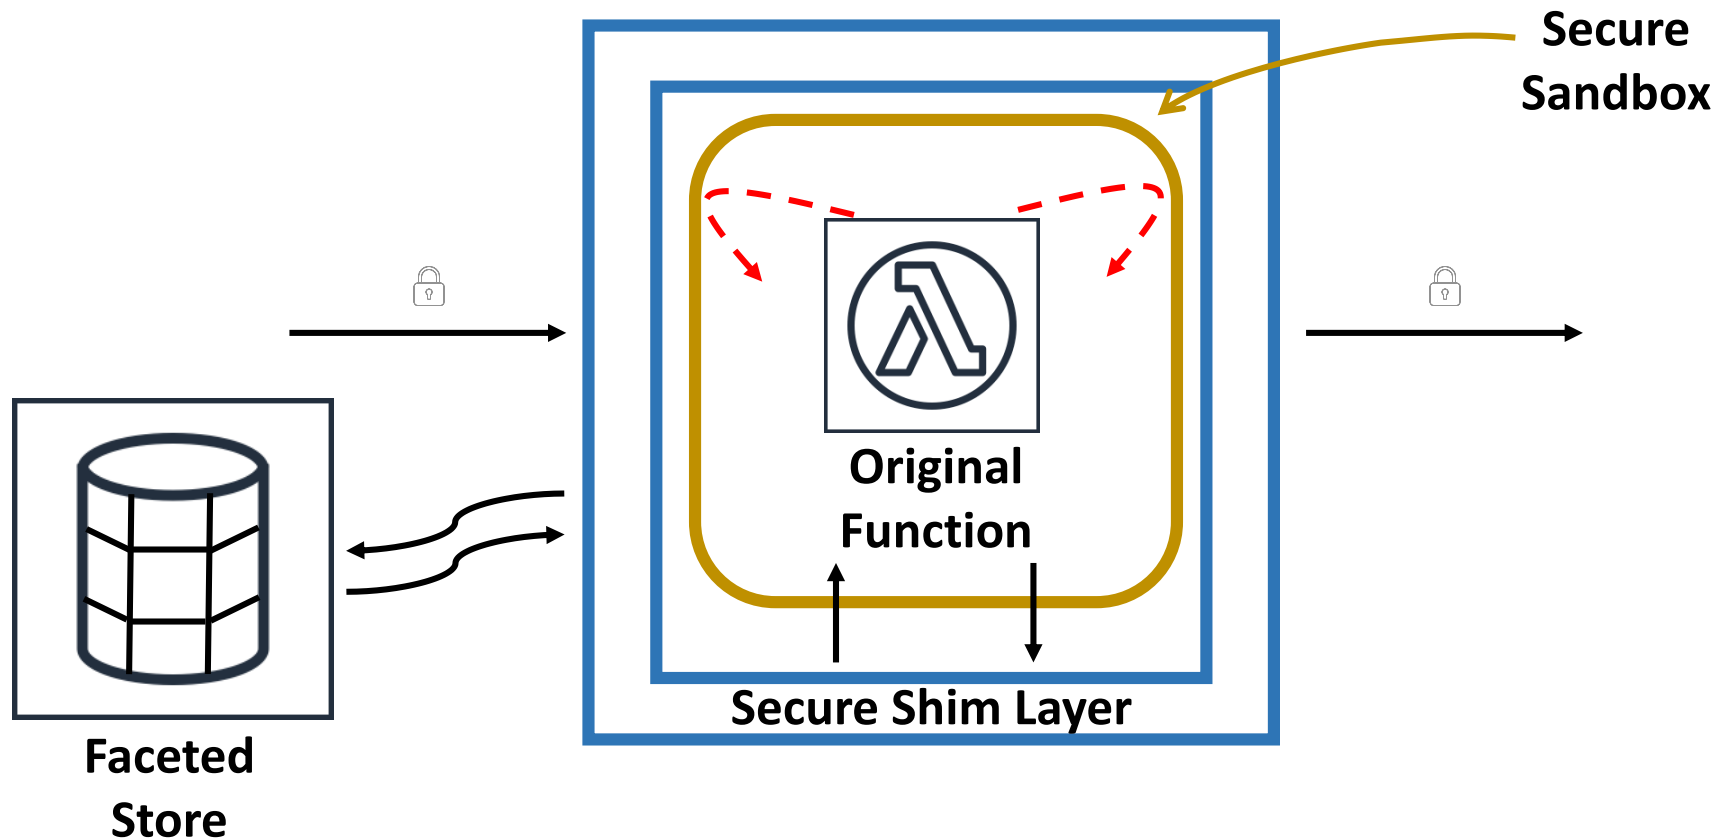
- Runtime enforcement of *Termination Sensitive Non-Interference*
  - Attacker cannot leak information by affecting execution termination
- Function executions get a security label when invoked
  - Label matches the security privilege of the invoker
- Faceted data store semantics
  - Items in the DB may have multiple values
  - Depending on the security label of the function that wrote the data

# Trapeze – Goals

- Practical IFC system for serverless applications
- Low overhead in runtime
- Allow running existing application securely with minimal modification
- Transparency – secure applications have the same semantics as when running without Trapeze

# Trapeze Architecture

Shim layer wraps functions, controlling communication and enforcing faceted store semantics



# Evaluation

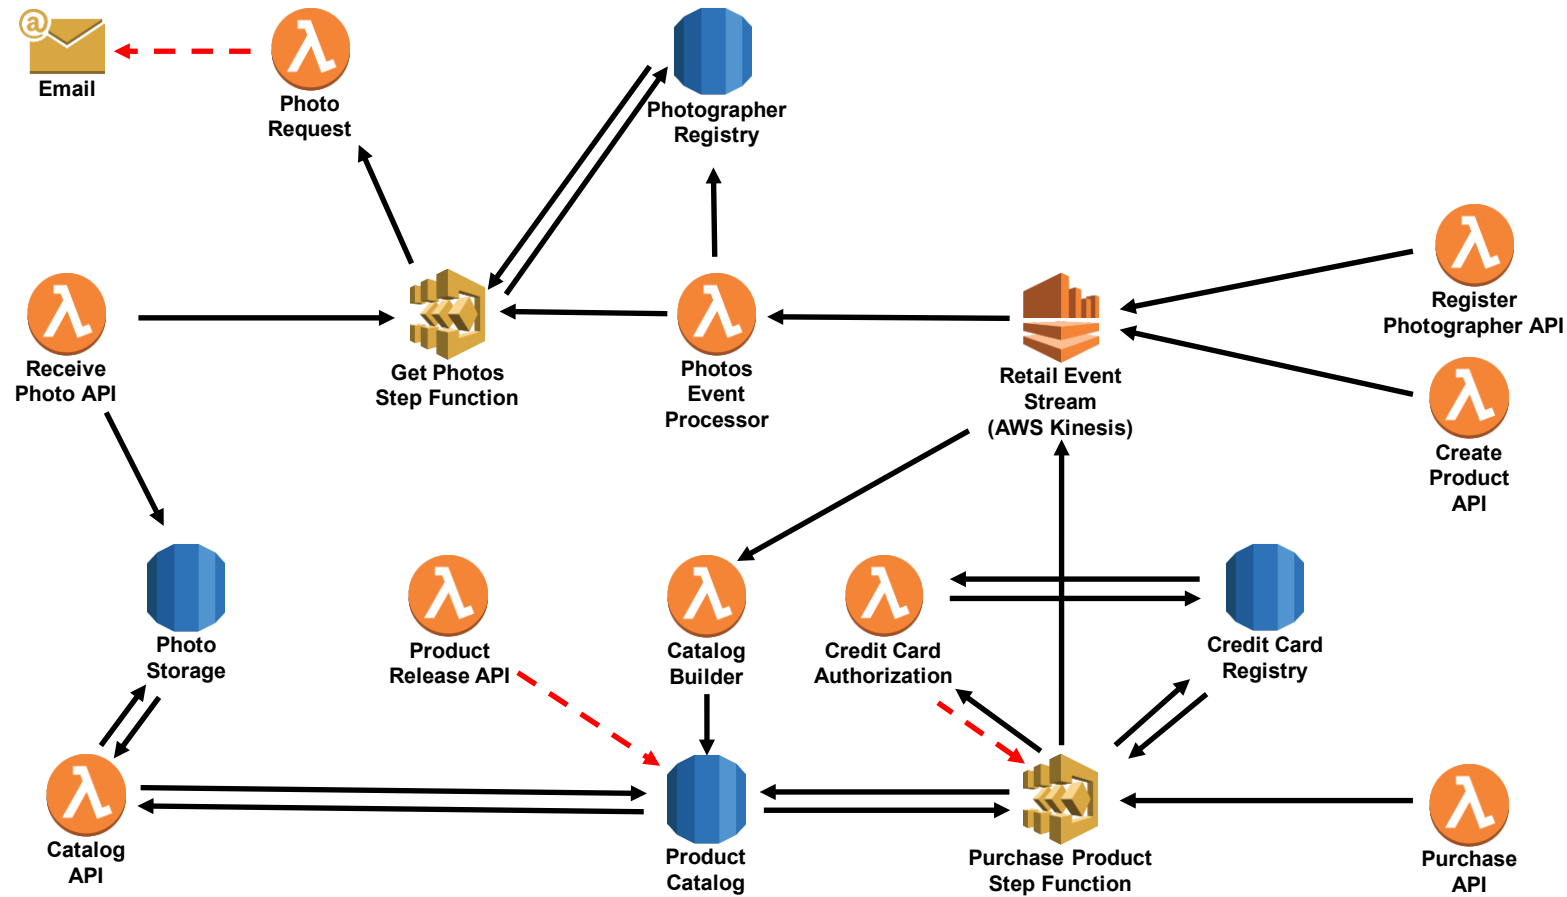
Test cases:

- Nordstrom Hello, Retail!
- gg
- Image feature extraction

Questions:

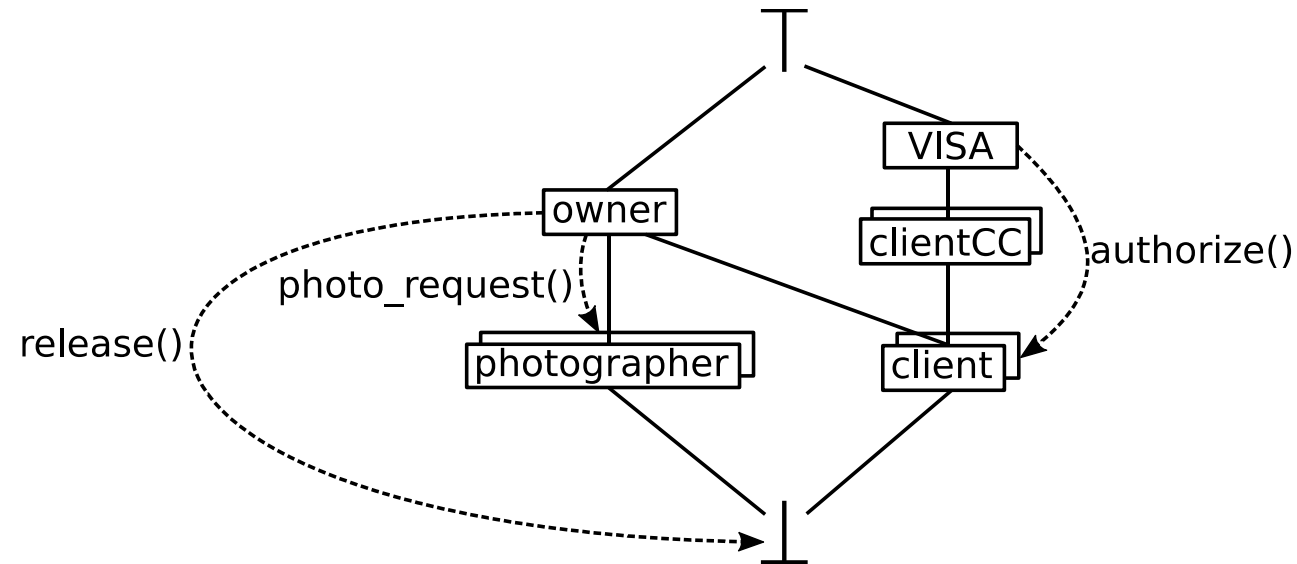
- Performance
- Ease of deployment
- Transparency

# Evaluation – Hello, Retail!



# Evaluation – Hello, Retail!

Security policy:





# Evaluation Results

Scenario	# $\lambda$	$\lambda$ runtime (ms)		$\Delta(\%)$	App. code	Modified
		baseline	Trapeze			
Hello, Retail!						
Update & Purchase	435	51,246	98,563	92.33	1300	3
Build & Browse	955	257,799	391,900	52.03		
gg						
mosh	111	681,173	654,448	-3.92	8800	2
git	917	2,602,500	2,660,721	2.24		
vim	234	1,242,873	1,338,128	7.66		
openssh	654	1,626,233	1,649,139	1.41		
Image Feature Extraction						
Upload	1	475	525	10.4	95	0
Extract	1	1,882	2,114	12.3		

# Evaluation Results

Scenario	# $\lambda$	$\lambda$ runtime (ms)		$\Delta(\%)$	App. code	Modified
		baseline	Trapeze			
Hello, Retail!						
Update & Purchase	435	51,246	98,563	92.33	1300	3
Build & Browse	955	257,799	391,900	52.03		
gg						
mosh	111	681,173	654,448	-3.92	8800	2
git	917	2,602,500	2,660,721	2.24		
vim	234	1,242,873	1,338,128	7.66		
openssh	654	1,626,233	1,649,139	1.41		
Image Feature Extraction						
Upload	1	475	525	10.4	95	0
Extract	1	1,882	2,114	12.3		

# Evaluation Results

Scenario	# $\lambda$	$\lambda$ runtime (ms)		$\Delta(\%)$	App. code	Modified
		baseline	Trapeze			
Hello, Retail!						
Update & Purchase	435	51,246	98,563	92.33	1300	3
Build & Browse	955	257,799	391,900	52.03		
gg						
mosh	111	681,173	654,448	-3.92	8800	2
git	917	2,602,500	2,660,721	2.24		
vim	234	1,242,873	1,338,128	7.66		
openssh	654	1,626,233	1,649,139	1.41		
Image Feature Extraction						
Upload	1	475	525	10.4	95	0
Extract	1	1,882	2,114	12.3		

# Summary

- Serverless computing
  - Security Risks
  - Security opportunities
- Trapeze
  - A practical IFC system for serverless applications
    - Guarantees TSNI
  - Low overhead (mostly) thanks to function granularity
  - High transparency

**GitHub** <https://github.com/kalevalp/trapeze>